

# P2 Documentation and Code Project

You are viewing draft content  
created with the use of Claude.AI



# P2 Assembly Programming

*A Human-Centered Approach to Parallel Processing*

June 2026

Version 3.0

## Tutorial Philosophy

Learn by doing, celebrate progress, have fun!

### Code Block Colors:

- **Green** – Spin2
- **Yellow** – PASM2
- **Purple** – CORDIC
- **Blue** – Multi-COG
- **Red** – Antipattern

### Teaching Elements:

- **Sidetracks** – deeper dives
- **Medicine Cabinet** – simpler alternatives
- **Your Turn** – hands-on exercises
- **Interludes** – stories & context

# Contents

<b>Copyright and License</b>	<b>8</b>
<b>Dedication</b>	<b>9</b>
<b>Acknowledgments</b>	<b>10</b>
<b>Preface: Welcome to the Journey</b>	<b>12</b>
<b>Chapter 1: Your First Spin</b>	<b>15</b>
Why P2? . . . . .	15
The Hook: Making Light . . . . .	15
What’s Really Happening . . . . .	16
Sidetrack: The Clock Preamble . . . . .	17
Let’s Make It Better . . . . .	18
Understanding COGs . . . . .	19
Your Turn: Experiments . . . . .	19
Sidetrack: Why Start at Address 0? . . . . .	21
Common Gotchas . . . . .	22
What We’ve Learned . . . . .	22
Coming Up Next . . . . .	22
<b>Chapter 2: Architecture Safari</b>	<b>23</b>
The Propeller Philosophy . . . . .	23
COG Anatomy 101 . . . . .	24
Meet the Hub: The Meeting Place . . . . .	25
Let’s See COGs in Action . . . . .	26
COG Communication: How They Talk . . . . .	26
The Timer: Everyone Gets One . . . . .	28
Why No Interrupts? (Usually) . . . . .	28
Real-World Example: Parallel Sensors . . . . .	29
Sidetrack: The Philosophy of Parallel . . . . .	30
Common Gotchas . . . . .	31
What We’ve Learned . . . . .	31
Your Turn: Experiments . . . . .	32
Coming Up Next . . . . .	33
<b>Chapter 3: Speaking PASM2</b>	<b>34</b>

---

The Hook: One Instruction, Many Powers . . . . .	34
Instruction Anatomy 101 . . . . .	34
The Basic Vocabulary . . . . .	35
Flow Control: Jump! . . . . .	37
Labels: Naming Your Places . . . . .	40
Data in DAT Blocks: Your Program's Pantry . . . . .	42
The Flags: C and Z (and Q!) . . . . .	48
Special Instructions That Will Blow Your Mind . . . . .	49
Real-World Example: Fast Memory Copy . . . . .	50
Common Gotchas . . . . .	51
Your Turn: Experiments . . . . .	52
Sidetrack: Why PASM2 Is Different . . . . .	53
What We've Learned . . . . .	54
Coming Up Next . . . . .	54
<b>Chapter 4: The Hub Connection</b> . . . . .	<b>55</b>
The Hook: Instant Communication . . . . .	55
Reading from Hub . . . . .	55
Writing to Hub . . . . .	55
The FIFO Pipeline . . . . .	56
Real-World Example: Video Buffer . . . . .	56
Your Turn: Experiments . . . . .	57
Common Gotchas . . . . .	57
What We've Learned . . . . .	58
Coming Up Next . . . . .	58
<b>Chapter 5: Mathematics Unleashed</b> . . . . .	<b>59</b>
The Hook: Hardware Multiply . . . . .	59
The Multiplication Revolution . . . . .	59
Division Without Tears . . . . .	59
64-Bit Operations . . . . .	60
Real-World Example: Fixed-Point Math . . . . .	60
Your Turn: Experiments . . . . .	61
Common Gotchas . . . . .	61
What We've Learned . . . . .	62
Coming Up Next . . . . .	62
<b>Chapter 6: Flags and Decisions</b> . . . . .	<b>63</b>
The Hook: Any Instruction Can Be Conditional . . . . .	63
The C and Z Flags . . . . .	63
Complex Conditions . . . . .	63
Skip Patterns - Conditional Blocks . . . . .	64
Your Turn: Experiments . . . . .	64

---

Common Gotchas . . . . .	65
What We've Learned . . . . .	65
Coming Up Next . . . . .	65
<b>Chapter 7: CORDIC Magic</b>	<b>66</b>
The Hook: Rotate a Point in 3 Lines . . . . .	66
What Just Happened? . . . . .	66
The CORDIC Pipeline - Your Mathematical Assembly Line . . . . .	67
Core CORDIC Operations . . . . .	67
Real-World Example: Spinning a Sprite . . . . .	69
Your Turn: CORDIC Experiments . . . . .	70
Advanced CORDIC: The Pipeline Dance . . . . .	72
CORDIC for Graphics . . . . .	73
CORDIC for Audio . . . . .	73
Common CORDIC Gotchas . . . . .	74
What About QLOG, QEXP? . . . . .	74
What We've Learned . . . . .	75
Coming Up Next . . . . .	76
<b>Chapter 8: Basic I/O</b>	<b>77</b>
The Hook: One Pin, Three Instructions, Infinite Possibilities . . . . .	77
Understanding P2 Pins . . . . .	77
Digital Output: Making Things Happen . . . . .	78
Digital Input: Reading the World . . . . .	79
Pin Timing: When Things Happen . . . . .	79
Real-World Example: Button Debouncing . . . . .	80
Bit-Banged Serial (The Basics) . . . . .	80
Your Turn: I/O Experiments . . . . .	81
Advanced Pin Control . . . . .	83
Common I/O Gotchas . . . . .	83
Timing Is Everything . . . . .	83
Real-World Example: Servo Control . . . . .	84
The Power of Simple . . . . .	84
What We've Learned . . . . .	84
A Note About Smart Pins . . . . .	85
Coming Up Next . . . . .	85
<b>Chapter 9: Streaming Data</b>	<b>86</b>
The Hook: 4KB in 4 Instructions . . . . .	86
Block Transfers: The Power Move . . . . .	86
The FIFO: Your Streaming Pipeline . . . . .	86
Writing Through the FIFO . . . . .	87
Real-World Example: Screen Buffer Clear . . . . .	87

---

Streaming with the Streamer . . . . .	88
FIFO and COG Execution . . . . .	88
Advanced Streaming Techniques . . . . .	89
Your Turn: Streaming Experiments . . . . .	90
Common Streaming Gotchas . . . . .	91
Performance Numbers . . . . .	91
Real-World Example: Audio Buffer . . . . .	92
What We've Learned . . . . .	93
Coming Up Next . . . . .	93
<b>Chapter 10: Hub Execution</b>	<b>94</b>
The Hook: Unlimited Code Space . . . . .	94
COG vs Hub Execution: The Trade-offs . . . . .	94
How Hub Execution Works . . . . .	95
Real-World Example: Menu System . . . . .	95
The Hub Execution FIFO . . . . .	97
Mixing COG and Hub Code . . . . .	97
Your Turn: Hub Execution Experiments . . . . .	98
Advanced Hub Execution . . . . .	99
Common Hub Execution Gotchas . . . . .	100
Real-World Example: Command Parser . . . . .	100
When to Use Hub Execution . . . . .	101
What We've Learned . . . . .	102
Coming Up Next . . . . .	102
<b>Chapter 11: Why No Interrupts?</b>	<b>103</b>
The Hook: Interrupts Without Interrupts . . . . .	103
The Interrupt Problem . . . . .	103
The Propeller Solution . . . . .	104
Real-World Example: Perfect Servo Control . . . . .	105
"But P2 HAS Interrupts!" . . . . .	105
The Event System: Better Than Interrupts . . . . .	107
Interrupt Horror Stories . . . . .	107
Your Turn: COG vs Interrupt Challenge . . . . .	108
The Philosophy Deep Dive . . . . .	109
When Interrupts Actually Make Sense . . . . .	110
Common "But What About..." Questions . . . . .	110
What We've Learned . . . . .	110
Coming Up Next . . . . .	111
<b>Chapter 12: Optimization Mastery</b>	<b>112</b>
The Hook: Double Your Speed with One Change . . . . .	112
Understanding the Pipeline . . . . .	112

---

Instruction Timing Basics . . . . .	113
REP: The Speed Loop . . . . .	113
SKIP: Conditional Execution on Steroids . . . . .	114
Hub Access Optimization . . . . .	114
The FIFO Fast Path . . . . .	115
Parallel Operations . . . . .	115
Real-World Example: Fast Memory Copy . . . . .	116
Your Turn: Optimization Challenges . . . . .	117
Advanced Techniques . . . . .	118
Common Optimization Gotchas . . . . .	119
Profiling and Measurement . . . . .	119
What We've Learned . . . . .	120
Coming Up Next . . . . .	120
<b>Chapter 13: LUT Memory - Your Private Lookup Table</b>	<b>122</b>
The Hook: A Lookup Table in 3 Cycles . . . . .	122
Why Another Memory? . . . . .	122
Reading and Writing the LUT . . . . .	123
LUT Sharing Between COGs . . . . .	124
Practical Examples . . . . .	125
LUT with the Streamer . . . . .	126
Common Gotchas . . . . .	127
Medicine Cabinet . . . . .	128
Your Turn . . . . .	128
What We've Learned . . . . .	129
Coming Up Next . . . . .	129
<b>Chapter 14: Smart Pins Orientation</b>	<b>130</b>
The Hook: A UART in 4 Lines . . . . .	130
What Are Smart Pins, Really? . . . . .	130
The Universal Smart Pin Pattern . . . . .	130
The Core Instructions . . . . .	131
Understanding the IN Flag . . . . .	132
Common Smart Pin Modes . . . . .	133
Configuration Values Demystified . . . . .	135
Common Gotchas . . . . .	135
Medicine Cabinet . . . . .	136
Your Turn . . . . .	137
What We've Learned . . . . .	138
Coming Up Next . . . . .	138
<b>Chapter 15: Event-Driven Programming</b>	<b>139</b>
The Hook: No More Polling Loops . . . . .	139

---

Why Events Matter . . . . .	139
The Four Selectable Events . . . . .	140
Configuring an Event . . . . .	140
Timer Events . . . . .	143
Waiting vs Polling . . . . .	143
Multiple Events . . . . .	144
Practical Examples . . . . .	144
ATN - Inter-COG Events . . . . .	146
Common Gotchas . . . . .	146
Medicine Cabinet . . . . .	147
Your Turn . . . . .	149
What We've Learned . . . . .	149
Coming Up Next . . . . .	150
<b>Chapter 16: Multi-COG Orchestration</b>	<b>151</b>
The Hook: A Complete System in 8 COGs . . . . .	151
Communication Patterns . . . . .	152
Synchronization Techniques . . . . .	154
Real-World Example: Robot Controller . . . . .	155
Your Turn: Multi-COG Project . . . . .	157
Design Principles for Multi-COG Systems . . . . .	158
Common Multi-COG Gotchas . . . . .	159
What We've Learned . . . . .	159
Your Journey Continues . . . . .	159
Chapter Summary . . . . .	160
Epilogue: The Journey Forward . . . . .	160
Further Reading . . . . .	161
<b>Appendix A: Platform Comparison</b>	<b>163</b>
The Landscape . . . . .	163
What Makes P2 Different . . . . .	163
Coming From ARM/STM32 . . . . .	165
Coming From ESP32 . . . . .	165
Coming From Arduino/AVR . . . . .	166
When P2 Is the Right Choice . . . . .	166
Platform Trade-offs . . . . .	166
Community Resources . . . . .	167
The P2 Hardware Ecosystem . . . . .	167
Summary . . . . .	168
<b>Index</b>	<b>169</b>

# Copyright and License

Copyright © 2025-2026 Iron Sheep Productions, LLC and Parallax Inc.

This work is licensed under the Creative Commons Attribution–NonCommercial–NoDerivatives 4.0 International License (CC BY-NC-ND 4.0).

You are free to:

- **Share** — copy and redistribute the material in any medium or format

Under the following terms:

- **Attribution** — You must give appropriate credit, provide a link to the license, and indicate if changes were made (for example, formatting or excerpting).
- **NonCommercial** — You may not use the material for commercial purposes.
- **NoDerivatives** — If you remix, transform, translate, or build upon the material, you may not distribute the modified material.

**Commercial use:** For uses that may be commercial (including paid courses, kits, or redistribution with products), please contact Iron Sheep Productions, LLC and Parallax Inc. (info@ironsheep.biz) for separate permission.

To view the full license, visit: <https://creativecommons.org/licenses/by-nc-nd/4.0/>

## Trademarks

Propeller, Propeller 2, P2, Spin, and the Parallax logo are trademarks of Parallax Inc.

## Disclaimer

The information in this manual is subject to change without notice. While every effort has been made to ensure accuracy, the authors and publishers assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

# Dedication

**To deSilva** — *Whose legendary P1 assembly tutorial taught a generation of programmers that assembly language could be approachable, enjoyable, and even fun. Your unique voice—combining technical precision with human warmth—showed us that great documentation teaches not just the mind, but speaks to the spirit of discovery.*

**To the Propeller Community** — *Who have spent countless hours exploring, documenting, and sharing their knowledge. From the early P1 pioneers to today’s P2 innovators, your collective wisdom makes this manual possible.*

**To Future Makers** — *May you find in these pages the same joy of discovery that we experienced. The Propeller 2 is more than a microcontroller—it’s an invitation to think differently about computing. Welcome to the journey.*

*“The best way to predict the future is to invent it.”* — Alan Kay

# Acknowledgments

This manual stands on the shoulders of giants. We gratefully acknowledge:

## Primary Contributors

**deSilva** - For creating the gold standard of microcontroller documentation with the P1 Assembly Tutorial. Your pedagogical approach, combining technical depth with human empathy, remains unmatched. This manual attempts to honor your legacy while adapting to the P2's capabilities.

**Iron Sheep Productions LLC (Stephen M Moraco)** - For extensive P2 documentation efforts, community tools, and the vision of creating an AI-optimized knowledge base. Your systematic approach to extracting and organizing P2 knowledge made this comprehensive manual possible.

**Chip Gracey** - Creator of the Propeller architecture. Thank you for giving us a microcontroller that thinks differently and challenges us to do the same.

## Community Contributors

**The Parallax Forums Community** - Your questions, answers, code examples, and endless experimentation have created a living knowledge base that no single author could match.

**Early P2 Adopters** - Who dealt with evolving documentation, changing specifications, and still produced amazing projects that showed us what was possible.

## Technical Reviewers

Special thanks to those who reviewed drafts, tested code examples, and provided invaluable feedback:

- The P2 Documentation Team at Parallax
- Community members who beta-tested examples
- Everyone who reported errors and suggested improvements

## Inspiration

**The MIT AI Lab** - For showing us that technical documentation can have personality

**Donald Knuth** - For proving that programming texts can be literature

**The Demoscene Community** - For pushing hardware beyond its limits and inspiring us to do the same

## Production Notes

This manual was created using:

- Knowledge extracted from official Parallax technical documentation and OBEX (Object Exchange) community contributions
- AI-assisted content generation trained on deSilva's writing style
- Community validation and real-world testing
- A commitment to making parallel processing accessible to everyone

*“If I have seen further, it is by standing on the shoulders of giants.”* — Isaac Newton

Any errors, omissions, or dad jokes that fell flat are entirely the responsibility of the authors, not our distinguished contributors.

# Preface: Welcome to the Journey

Well, here we are! You're about to embark on a journey into the heart of the Propeller 2, and I promise you, it's going to be quite different from what you might expect.

## A Different Kind of Processor

The Propeller 2 isn't just another microcontroller. Oh no, it's something far more interesting. Imagine, if you will, eight independent processors (we **CALL** them COGs) all working together in perfect harmony, sharing a common memory space, yet each running their own programs at full speed. No interrupts fighting for attention, no complex priority schemes, just eight brains working in parallel.

And if you think this sounds terribly complicated, you're probably right... but here's the secret: it's actually simpler than traditional architectures once you understand the philosophy.

## About This Manual

This manual follows in the footsteps of deSilva's legendary P1 tutorial. What does that mean? It means we're going to:

1. **Start with working code** - You'll be blinking LEDs before you know what hit you
2. **Learn by doing** - Theory is important, but nothing beats hands-on experience
3. **Have some fun** - Yes, assembly language can actually be enjoyable!
4. **Be honest about complexity** - When something is hard, we'll admit it and then show you how to handle it

## Who Is This For?

Are you a complete beginner to assembly language? Welcome! We'll take good care of you.

Are you a grizzled veteran who's been writing assembly since the 6502? Welcome! The P2 will still surprise you.

Are you somewhere in between? Perfect! This is exactly where you want to be.

The only requirement is curiosity and a willingness to think a bit differently about how computers can work.

## How to Use This Manual

**The Fast Track** — If you're impatient (and who isn't?), jump straight to Chapter 1. Get that LED blinking. Feel the satisfaction. Then come back here when you're ready for more.

**The Scenic Route** — Read the chapters in order. Each builds on the previous one, and I've hidden little gems of knowledge throughout that will make later chapters easier.

**The Reference Approach** — Already know what you're looking for? The table of contents and index are your friends. The appendices contain every instruction, every Smart Pin mode, every CORDIC operation.

## What Makes the P2 Special?

Let me count the ways:

- **8 symmetric COGs** - No master/slave relationships, all COGs are equal
- **64 Smart Pins** - Each pin has its own processor for I/O operations
- **CORDIC engine** - Hardware trigonometry and coordinate transformations
- **Hardware multiply/divide** - Finally! Real math at hardware speed
- **512KB of RAM** - Shared by all COGs with deterministic access timing
- **No interrupts** - Well, actually there are interrupts, but we'll talk about why you probably don't want them

## A Note on Our Approach

The best technical documentation remembers you're human. You'll get frustrated. You'll make mistakes. Your code won't work the first time (or the second, or sometimes even the third).

That's normal. That's learning. And that's why this manual provides plenty of "medicine" along the way - simpler alternatives, working examples, and the occasional moment of levity to keep your spirits up.

## The deSilva Spirit

Throughout this manual, you'll encounter the teaching spirit of deSilva. When you see phrases like:

- "Well, ..." - We're about to correct a common assumption
- "Uff!" - We just got through something complex
- "Have Fun!" - We mean it, this stuff is actually enjoyable

These aren't just quirks; they're signals that we remember you're human and we're on this journey together.

## Ready?

Take a deep breath. Pour yourself your favorite beverage. Open your development environment.

Let's make some magic happen with the Propeller 2!

*“The Propeller architecture is based on the simple idea that the best way to avoid the complexity of interrupts is to have enough processors that you don't need them.”*

— Chip Gracey, creator of the Propeller

**Turn the page, and let's blink that LED! →**

# Chapter 1: Your First Spin

*Let's blink an LED and change your life*

## Why P2?

Before we dive into code, let me tell you why you're in for something different.

If you've fought with interrupt priority conflicts on an ARM, watched your timing jitter because of cache misses, or discovered that the UART you need is only available on pins you're already using... well, the P2 was designed by someone who got tired of those problems too.

Here's the P2 philosophy in a nutshell:

**Instead of one processor fighting with interrupts**, you get eight complete, identical processors (COGs) that run truly in parallel. Your serial handler never delays your motor control. Your sensor sampling never misses a deadline. Each task owns its own processor.

**Instead of fixed peripherals**, every one of the 64 pins contains its own programmable state machine. Any pin can become a UART, PWM output, quadrature encoder, ADC - whatever you need, wherever you need it.

**Instead of timing that depends on cache luck**, the hub memory has deterministic access. Your timing loops work the same way every time.

**Instead of calling math libraries**, there's a hardware CORDIC that gives you sine, cosine, and arctangent (via vector rotate and convert) with results ready in exactly 55 clocks. Every time.

Does this mean P2 is perfect for everything? Of course not. But if your projects involve multiple real-time tasks, precise timing, video or audio generation, or just running out of peripheral pins - you're in the right place.

For a full comparison to ARM, ESP32, Arduino, and PIC platforms, see Appendix A. But you probably want to blink that LED first, don't you?

## The Hook: Making Light

I know you're absolutely crazy to have your first instruction executed, so let's not waste any time. Here's a complete PASM2 program that blinks an LED on pin 56 (that's the built-in LED on the P2 Eval board):

```
1 CON
2   _clkfreq = 200_000_000      ' 200 MHz system clock
3
```

*continues on next page →*

```

4 DAT
5 ' LED Blinker - Your first PASM2 program!
6     ORG     0           ' Start at COG address 0
7
8     DRVH    #56         ' Drive pin 56 high (LED on)
9     WAITX   ##50_000_000 ' Wait 0.25 seconds at 200MHz
10    DRVL    #56         ' Drive pin 56 low (LED off)
11    WAITX   ##50_000_000 ' Wait 0.25 seconds
12    JMP     #$-6        ' Jump back 6 longs (## adds hidden AUGD)

```

That’s it! Five lines of code and you have a blinking LED. Load this into any COG and watch the magic happen.

## What’s Really Happening

Well, now that you’ve seen it work (you did try it, right?), let’s talk about what’s actually going on here.

### The Instructions Decoded

**ORG 0** - This tells the assembler to start placing code at COG address 0. Every COG has its own private 512 longs (2KB) of memory, and execution always starts at address 0 when a COG is loaded.

**DRVH #56** - This drives pin 56 high (3.3V). The ‘h’ means high. The ‘#’ means we’re using an immediate value (the actual number 56) rather than the contents of register 56. One instruction, and your LED is on!

**WAITX ##50\_000\_000** - This waits for 50 million clock cycles. At 200 MHz, that’s 0.25 seconds. Notice the double ‘##’? That means this is a 32-bit immediate value. Single ‘#’ only gives us 9 bits.

**DRVL #56** - Drive low. LED off. You get the pattern.

**JMP #\$-6** - Jump back 6 longs. The ‘\$’ means “current address”, so ‘\$-6’ means “6 addresses back from here”. Why 6? Each ## immediate generates a hidden **AUGD** instruction (it augments the **WAITX** destination operand), so we have 6 longs total: **DRVH**, **AUGD**, **WAITX**, **DRVL**, **AUGD**, **WAITX**. Infinite loop achieved!

### But Wait, There’s More!

“Hold on,” you might say, “how does this even get into the COG?”

Ah, excellent question! In the real world, you’d typically launch this from Spin2 (the high-level language) like this:

```

1 CON
2   _clkfreq = 200_000_000   ' 200 MHz system clock
3
4 PUB main()
5   COGINIT(COGEXEC_NEW, @blink_code, 0)   ' Start PASM2 in new COG
6   repeat   ' Keep the main COG alive
7
8 DAT
9   ORG     0
10 blink_code
11   DRVH   #56
12   WAITX  ##50_000_000
13   DRVL   #56
14   WAITX  ##50_000_000
15   JMP    #-$-6

```

The **COGINIT** instruction loads your PASM2 code from hub memory into a fresh COG and starts it running. Meanwhile, your Spin2 code keeps running in its own COG. You now have parallel processing!

## Sidetrack

### The Clock Preamble

Notice the **CON** section at the top of that example? Every P2 program needs to configure its system clock:

```

1 CON
2   _clkfreq = 200_000_000   ' 200 MHz system clock

```

This tells the P2 to run at 200 MHz using your board's crystal oscillator. Without it, the chip runs at a sluggish ~20 MHz on its internal RC oscillator—and timing-dependent code (including **DEBUG** output) won't behave as expected.

At 200 MHz with most instructions taking 2 clocks, each COG executes approximately 100 million instructions per second (100 MIPS). With 8 COGs running in parallel, that's 800 MIPS of total processing power—and that's before Smart Pins start handling I/O autonomously.

**From here on, we'll omit this preamble from examples to keep them focused on the concept being taught.** When you create your own files, always include it at the top before your **PUB** or **DAT** sections.

## Let's Make It Better

The blinker works, but it's a bit rigid, isn't it? What if we want to change the blink rate? Let's use a register:

```

1          ORG      0
2
3          MOV      delay, ##50_000_000    ' Set delay to 0.25 sec
4                                          ' (at 200 MHz)
5
6 .blink  DRVH      #56                    ' LED on
7          WAITX    delay                  ' Wait
8          DRVL      #56                    ' LED off
9          WAITX    delay                  ' Wait
10         JMP      #.blink                 ' Repeat forever
11
12 delay   LONG     0                      ' Storage for delay value

```

Uff! Look at that - we're using a register now! The **MOV** instruction copies our delay value into a register (which we cleverly named 'delay'). Now we can change the blink rate by modifying just one value.

*A note on terminology: P2 documentation often uses “register” to refer to any long in COG RAM. Unlike ARM or x86 where registers are a small, special set (R0-R15, EAX, etc.), every COG RAM location can be used as a general-purpose register. However, the last 16 locations (496-511) have special functions: addresses 496-503 are dual-purpose (usable as RAM if interrupts aren't used), and 504-511 are special-purpose registers (PTRA, PTRB, DIRA, DIRB, OUTA, OUTB, INA, INB). When you see “register” in P2 context, think “COG RAM location.”*

## Understanding COGs

Here's something important: each COG is a complete processor with its own memory. When we loaded our blink program, it was copied from hub memory into COG memory. The COG then executes it independently, without any further connection to hub memory (unless we explicitly read or write to it).

Think of it like this:

- **Hub memory** is the meeting place (512KB shared by all)
- **COG memory** is private workspace (2KB per COG)
- Loading a COG is like making a photocopy - the COG gets its own copy to run

This is why our blinker keeps running even after the Spin2 code that launched it goes into an infinite repeat loop. The COG is independent!

## Your Turn: Experiments

Now for the fun part. Try these modifications:

### Experiment 1: Different Patterns

Make the LED blink in a pattern: short-short-long (like SOS):

```

1      ORG      0
2
3      MOV      short, ##20_000_000    ' 0.1 second (at 200 MHz)
4      MOV      long_d, ##60_000_000   ' 0.3 seconds (at 200 MHz)
5
6 .pattern DRVH    #56                  ' Short pulse 1
7      WAITX    short
8      DRVL     #56
9      WAITX    short
10
11     DRVH     #56                      ' Short pulse 2
12     WAITX    short
13     DRVL     #56
14     WAITX    short
15
16     DRVH     #56                      ' Long pulse
17     WAITX    long_d
18     DRVL     #56

```

*continues on next page →*

```

19      WAITX   long_d
20
21      JMP     #.pattern
22
23 short LONG   0
24 long_d LONG  0

```

## Experiment 2: Multiple LEDs

Blink LEDs on pins 56 and 57 alternately:

```

1      ORG     0
2
3 .loop DRVH   #56           ' LED 56 on
4      DRVL   #57           ' LED 57 off
5      WAITX  ##50_000_000 ' 0.25 sec at 200 MHz
6
7      DRVL   #56           ' LED 56 off
8      DRVH   #57           ' LED 57 on
9      WAITX  ##50_000_000 ' 0.25 sec at 200 MHz
10
11     JMP     #.loop

```

## Experiment 3: Fading (Advanced)

This one's a bit tricky - we'll use PWM to fade the LED:

```

1      ORG     0
2
3      DIRL   #56           ' Reset the pin before configuring
4      WRPIN  ##P_PWM_TRIANGLE, #56 ' Configure pin 56 for PWM
5      WXPIN  ##$0100_0001, #56   ' Frame = 256 base periods
6      DIRH   #56           ' Enable the pin
7
8 .fade WYPIN  level, #56       ' Set duty cycle
9      WAITX  ##100_000        ' Small delay
10     ADD    level, #1         ' Increment brightness
11     AND    level, #$FF      ' Wrap at 256
12     JMP    #.fade

```

```

13
14 level    LONG    0

```

Don't worry if the PWM example seems complex - we'll cover Smart Pins in detail in Chapter 8!

### The Medicine Cabinet

Feeling overwhelmed? Here's the simplified prescription:

**Minimum viable blinker** - Just 3 instructions:

```

1  .loop    DRVNOT  #56          ' Toggle pin 56
2          WAITX   ##50_000_000 ' 0.25s at 200MHz
3          JMP     #.loop        ' Repeat

```

The **DRVNOT** instruction toggles a pin - if it's high, make it low; if it's low, make it high. Sometimes simpler is better!

### Sidetrack

#### Why Start at Address 0?

You might wonder why COG code always starts at address 0. It's actually quite elegant:

When a COG is started with `COGINIT`, the hardware:

1. Stops the COG (if it was running)
2. Copies 504 longs from hub to COG memory (addresses 0-503); the top 8 (504-511) are special hardware registers, not loaded
3. Starts execution at COG address 0

This means every COG program starts fresh, with a clean slate. No residual state, no confusion. It's like each COG gets a fresh brain transplant every time it starts!

The last 16 longs (addresses 496-511) have special functions: 496-503 are dual-purpose (usable as RAM if interrupts not used), and 504-511 are special-purpose registers (PTRA, PTRB, DIRA, DIRB, OUTA, OUTB, INA, INB). We'll explore these later.

## Common Gotchas

Before we move on, let me save you some debugging time:

1. **Forgetting the ##** - Using `WAITX #25_000_000` will NOT wait for 0.25 seconds! Single `#` only allows values up to 511.
2. **Wrong pin number** - The P2 Eval board's LEDs are on pins 56-63. The P2 Edge module might have different assignments.
3. **Clock setup required** - P2 boots at ~20MHz (internal RC oscillator). Most programs configure 200MHz with a crystal. Our examples assume 200MHz - adjust **WAITX** values if your clock differs.
4. **COG already running** - If you `COGINIT` to a specific COG that's already running something else, it will be stopped and replaced. Use `COGEXEC_NEW` to automatically find a free COG.

## What We've Learned

Let's celebrate what you've accomplished:

- Written your first PASM2 program
- Controlled hardware (LED) directly
- Used immediate values (`#` and `##`)
- Created loops with **JMP**
- Understood COG independence
- Modified code for different patterns

That's quite a lot for Chapter 1!

## Coming Up Next

In Chapter 2, we'll take our "Architecture Safari" and explore:

- How 8 COGs really work together
- The hub memory system and the "egg beater"
- Why the P2 doesn't need interrupts
- How to make COGs talk to each other

But for now, enjoy your blinking LED. You've just taken your first step into parallel processing!

**Have Fun!** And remember, every expert was once a beginner who kept their LED blinking when everyone else gave up.

# Chapter 2: Architecture Safari

*Eight brains are better than one*

## The Propeller Philosophy

Before we dive into the technical details, let's talk philosophy. Why would anyone design a micro-controller with eight processors?

The answer is beautifully simple: to avoid complexity.

“Wait,” you might say, “eight processors sounds MORE complex, not less!”

Well, consider the traditional approach:

- One processor trying to do everything
- Interrupts constantly breaking your flow
- Priority levels to juggle
- Context switching overhead
- Race conditions and timing nightmares

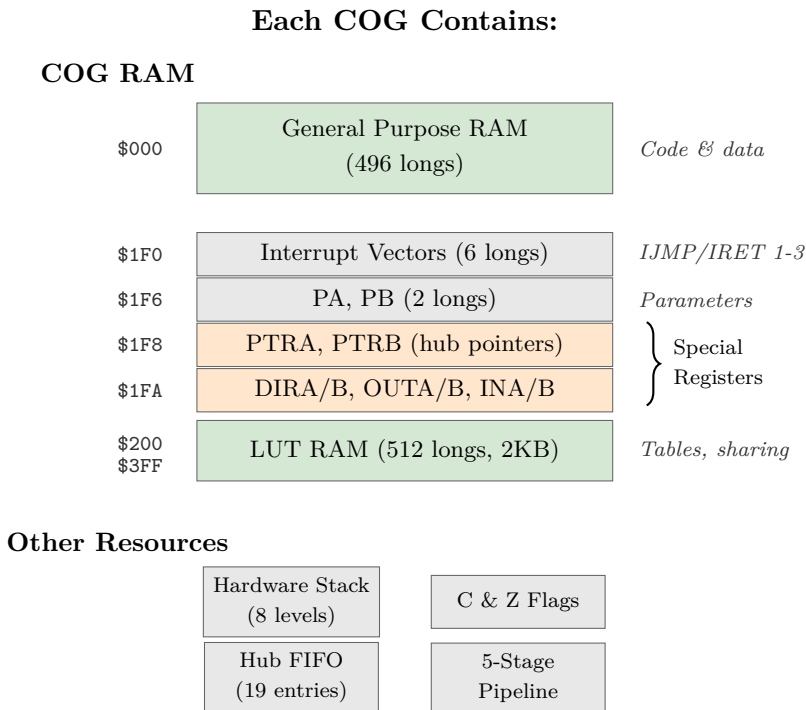
Now consider the Propeller way:

- Eight processors, each doing one thing well
- No interrupts needed (why interrupt when you have a dedicated processor?)
- No priorities (all COGs are equal)
- Deterministic timing (you know EXACTLY when things happen)
- True parallel processing (not time-slicing)

It's like the difference between one stressed-out juggler trying to keep eight balls in the air versus eight relaxed people each tossing one ball. Which seems simpler?

## COG Anatomy 101

Let's dissect a COG and see what makes it tick:



But here's the beautiful part: COGs are identical. There's no "master" COG or "special" COG. Any COG can do anything any other COG can do. Democracy in silicon!

### The 512-Long Limit

Each COG has exactly 512 longs (2048 bytes) of memory. The first 496 longs are yours to use for code and data. The last 16 are special registers (but not like P1 - we'll get to that).

"Only 496 instructions?" you might cry. "That's tiny!"

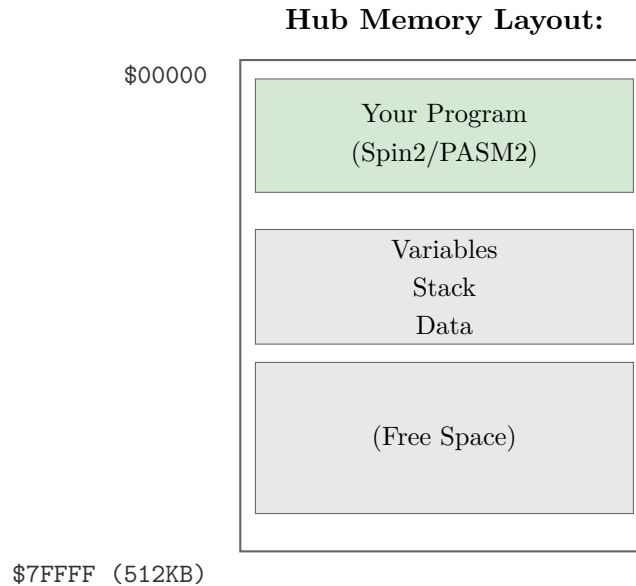
Well, yes and no. Remember:

1. PASM2 instructions are powerful - one instruction often does what takes several in other processors
2. You have EIGHT of these COGs
3. There's hub execution mode for larger programs (Chapter 10)
4. Most real-time tasks **FIT** easily in 496 instructions

Think of it like haiku - the constraint forces elegance.

## Meet the Hub: The Meeting Place

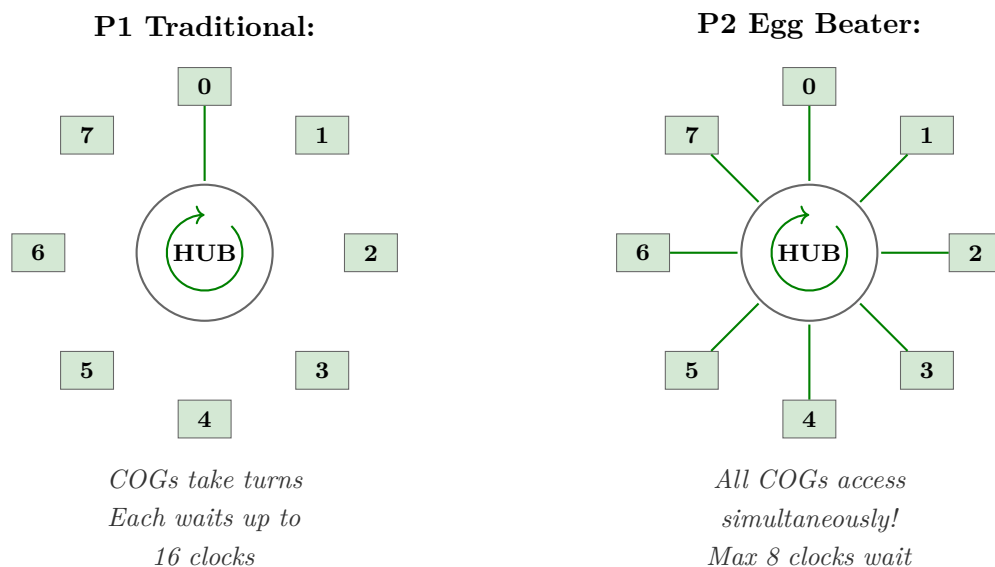
The hub is where COGs come together. It's 512KB of RAM shared by all COGs, and it's where the magic of cooperation happens.



## The Egg Beater Revolution

Now here's where P2 gets clever. In P1, COGs took turns accessing the hub in a round-robin fashion. If you missed your slot, you waited for the wheel to come around again.

P2 uses what we call the “egg beater” model. Imagine eight beaters (COGs) all whipping through the same bowl (hub) simultaneously, but their paths are cleverly arranged so they never collide:



The practical result? Hub access is MUCH faster and more predictable. Instead of waiting up to 16 clocks (P1), you wait at most 8 clocks (P2), and often less if you align your accesses properly.

## Let's See COGs in Action

Here's a simple demonstration of multiple COGs working together:

```

1  ' Multi-COG LED Pattern Demo
2  PUB main() | i
3      repeat i from 0 to 3
4          COGINIT(COGEXEC_NEW, @cog_code, 56 + i) ' Start 4 COGs
5      repeat ' Main COG just watches
6
7  DAT
8      ORG      0
9  cog_code
10     MOV      pin_num, ptra          ' Pin number was passed in via PTRA
11
12     .loop    DRVNOT pin_num         ' Toggle our LED
13         SHL      pin_num, #24      ' Pin number to bits 24-31
14         OR       pin_num, ##10_000_000 ' Combine with delay
15         WAITX    pin_num           ' Wait (varies per COG!)
16         SHR      pin_num, #24      ' Restore pin number
17         JMP      #.loop
18
19     pin_num LONG  0

```

What's happening here:

1. The main Spin2 code starts 4 COGs
2. Each COG gets a different pin number (56, 57, 58, 59)
3. Each COG blinks its LED at a slightly different rate
4. All four LEDs blink independently and simultaneously!

## COG Communication: How They Talk

COGs are independent, but they're not isolated. They can communicate through hub memory:

## Method 1: Simple Variables

```

1 ' COG 1: Writer
2     MOV     value, #42
3     WRLONG  value, ##$1000 ' Write to hub address $1000
4
5 ' COG 2: Reader
6     RDLONG  result, ##$1000 ' Read from hub address $1000

```

## Method 2: Locks (When It Matters)

When multiple COGs might write to the same location, we need locks:

```

1 ' Get a lock
2 .try_lock
3     LOCKTRY lock_id wc      ' Try to get lock (C=1 if we got it)
4     if_nc JMP     #.try_lock ' Keep trying until we get it
5
6     ' Critical section - we have the lock!
7     RDLONG  value, ##shared_addr
8     ADD     value, #1
9     WRLONG  value, ##shared_addr
10
11     LOCKREL lock_id        ' Release the lock
12
13 lock_id LONG    0          ' Lock 0-15

```

## Method 3: Mailboxes (Elegant)

A mailbox is just a hub location where COGs leave messages:

```

1 ' COG A: Leave a message
2     WRLONG  message, ##mailbox
3
4 ' COG B: Check for messages
5 .check RDLONG  data, ##mailbox wz
6     if_z JMP     #.check      ' Keep checking if empty
7     WRLONG  #0, ##mailbox    ' Clear mailbox
8     ' Process the message in 'data'

```

## The Timer: Everyone Gets One

Each COG has its own 64-bit timer, always counting system clocks. This is incredibly useful:

```

1 ' Method 1: Simple delay
2     GETCT    start_time
3     ADDCT1   start_time, ##1_000_000
4     WAITCT1           ' Wait exactly 1,000,000 clocks
5
6 ' Method 2: Periodic events
7     GETCT    time
8 .loop  ADDCT1   time, ##10_000_000
9     WAITCT1           ' Wait for next 10M clock interval
10    DRVNOT   #56      ' Toggle LED
11    JMP      #.loop   ' Perfectly periodic!

```

The beauty? Each COG's timer is independent. No shared resource conflicts!

## Why No Interrupts? (Usually)

Here's a controversial P2 feature: it HAS interrupts, but you probably shouldn't use them. Why?

Because with 8 COGs, you don't need interrupts! Instead of interrupting important work, just dedicate a COG to monitoring whatever would have triggered the interrupt:

```

1 ' Traditional (with interrupts):
2 ' Main code runs, gets interrupted, handles event, returns
3
4 ' Propeller way:
5 ' COG 1: Main code runs uninterrupted
6 ' COG 2: Watches for event continuously
7 pin_watcher
8     TESTP    #BUTTON_PIN wc
9     if_c JMP    #button_pressed
10    JMP      #pin_watcher
11
12 button_pressed
13    WRLONG   ##1, ##button_flag ' Signal other COGs
14    JMP      #pin_watcher

```

No interrupt latency, no context switching, no priority inversion. Just dedicated, deterministic monitoring.

## Real-World Example: Parallel Sensors

Let's read four different sensors simultaneously:

```
1 ' Each COG runs this with different parameters
2 sensor_reader
3     RDLONG  sensor_pin, ptra[0]    ' Get pin assignment
4     RDLONG  hub_addr, ptra[1]     ' Where to store results
5
6 read_loop
7     ' Read sensor (simplified - real sensors need protocols)
8     TESTP   sensor_pin wc
9     RCL     value, #1              ' Accumulate bits
10
11     ' Every 32 reads, store to hub
12     INCMOD counter, #31 wc
13     if_c WRLONG value, hub_addr
14     if_c ADD  hub_addr, #4
15
16     JMP     #read_loop
17
18 sensor_pin LONG 0
19 hub_addr   LONG 0
20 value      LONG 0
21 counter    LONG 0
```

Four COGs running this code = four sensors being read truly simultaneously. Try doing that with a single processor and interrupts!

## The Medicine Cabinet

Feeling overwhelmed by all this parallel processing? Here's your prescription:

**Start simple:** Use just one or two COGs at first

```

1 ' Just two COGs - main program and one helper
2 PUB main()
3     COGINIT(COGEXEC_NEW, @helper, 0)
4     ' Your main code here

```

**Debug one COG at a time:** Get each COG working alone before combining

```

1 ' Test COG in isolation first
2 debug_cog
3     DRVH     #MY_DEBUG_LED ' Visual confirmation it's running
4     ' Your actual code here

```

**Use Spin2 for coordination:** Let the high-level language handle the complex stuff

```

1 ' Spin2 manages COGs, PASM2 does the real-time work
2 PUB orchestrator()
3     startSensorCog(0)
4     startMotorCog(1)
5     startCommsCog(2)
6     ' Spin2 coordinates, PASM2 executes

```

## Sidetrack

### The Philosophy of Parallel

The Propeller's design philosophy comes from a simple observation: in the real world, things happen in parallel, not in sequence.

Consider your car:

- The engine runs continuously
- The radio plays independently
- The climate control maintains temperature
- The dashboard updates displays
- The ABS monitors wheel speed

*continues on next page →*

These aren't taking turns - they're all happening simultaneously. The Propeller models this reality directly. Instead of one processor frantically time-slicing between tasks, you have eight processors each focused on their job.

It's not just different - it's more natural.

## Common Gotchas

Save yourself some debugging time:

1. **COG RAM is copied, not shared** - Changes in COG RAM don't affect hub RAM unless you explicitly write them back
2. **COGs start at 0** - Always! Your code better be there.
3. **Hub addresses are byte addresses** - COG addresses are **LONG** addresses. Don't mix them up!

```
1  RDLONG  value, ##$1000  ' Reads from hub byte address $1000
2  MOV     value, $100     ' Moves from COG long address $100 (256)
3  ' Note: COG RAM is only 512 longs ($000-$1FF)!
```

4. **PTRA/PTRB are your friends** - These special registers make hub access much easier
5. **COGs are truly independent** - Stopping one COG doesn't affect others (unless they're waiting for it)

## What We've Learned

Look at what you now understand:

- Why eight processors is simpler than one with interrupts
- How COGs are structured and limited
- The hub memory system and egg beater access
- Multiple ways for COGs to communicate
- Why interrupts are usually unnecessary
- How to think in parallel

## Your Turn: Experiments

### Experiment 1: COG Counter

Start COGs to increment different hub locations. With **COGEXEC\_NEW**, the loop will start up to 7 new COGs (since COG 0 runs Spin2):

```

1 PUB main() | i
2     repeat i from 0 to 7
3         COGINIT(COGEXEC_NEW, @counter, $1000 + (i * 4))
4     repeat
5         ' Monitor the counters in hub RAM
6
7 DAT
8     ORG    0
9 counter MOV    hub_ptr, ptra           ' Our hub address arrived in PTRa
10 .loop  RDLONG  value, hub_ptr
11        ADD    value, #1
12        WRLONG  value, hub_ptr
13        WAITX  ##1_000_000
14        JMP    #.loop
15
16 hub_ptr LONG   0
17 value   LONG   0

```

### Experiment 2: Parallel Pattern

Make 8 LEDs display a moving pattern, with each COG controlling one LED:

```

1 ' Each COG gets LED pin in ptra
2     ORG    0
3     RDLONG pin, ptra
4     RDLONG delay, ptrb           ' Different delay per COG
5
6 .flash DRVH   pin
7        WAITX  delay
8        DRVL   pin
9        WAITX  delay
10       SHL    delay, #1           ' Double the delay
11       CMP    delay, ##100_000_000 wcz

```

*continues on next page →*

```
12    if_a MOV    delay, ##1_000_000  ' Reset if too long
13          JMP    #.flash
```

## Coming Up Next

In Chapter 3, “Speaking PASM2”, we’ll dive deep into the instruction set:

- The anatomy of an instruction
- Conditional execution that will blow your mind
- Math operations that actually make sense
- Why PASM2 is unlike any assembly you’ve used

But for now, appreciate what you’ve learned: you understand the Propeller’s parallel philosophy. That’s not just technical knowledge - it’s a new way of thinking about computing.

**Have Fun!** Remember, parallel processing isn’t harder - it’s different. And different can be wonderful.

# Chapter 3: Speaking PASM2

*Learning the native tongue*

## The Hook: One Instruction, Many Powers

Look at this single PASM2 instruction:

```
1      ADD      value, #1 wc
```

This one line:


- **ADDS** 1 to ‘value’
- Optionally sets the carry flag
- Executes in exactly 2 clock cycles
- Can be conditional
- Can even modify itself!

In most processors, that would take multiple instructions. In PASM2, it’s just one. Let’s learn to speak this powerful language.

## Instruction Anatomy 101

Every PASM2 instruction follows the same basic pattern:

### PASM2 Instruction Format:



Let’s dissect a real instruction:

### Example: if\_z add total, #10 wc



## The Basic Vocabulary

### Moving Data Around

The **MOV** family - your bread and butter:

```

1  ' Basic move
2      MOV    dest, source    ' dest = source
3      MOV    x, #42         ' x = 42 (immediate)
4      MOV    x, ##70000     ' x = 70000 (32-bit immediate)
5
6  ' But wait, there's more!
7      mvn   dest, source    ' dest = NOT source (inverted)
8      ABS   dest, source    ' dest = |source|
9      NEG   dest, source    ' dest = -source
10
11 ' And the mind-blowing ones
12     ALTD  dest, source    ' Modify NEXT inst's dest field!
13     ALTS  dest, source    ' Modify NEXT inst's source field!

```

Well, that escalated quickly! Don't worry about **ALTD**/**ALTS** yet - just know they exist and they're amazing.

### Math Without Tears

P2 has hardware multiply and divide. Let that sink in. Hardware. Multiply. And. Divide.

```

1  ' Addition and subtraction
2      ADD   x, y            ' x = x + y
3      SUB   x, y            ' x = x - y
4      ADDS  x, y            ' Signed add
5      SUBS  x, y            ' Signed subtract
6
7  ' The revolution: hardware multiply!
8      MUL   x, y            ' x = x * y (low 32 bits)
9      MULS  x, y            ' Signed multiply
10
11 ' And hardware divide!
12     QDIV  x, y            ' Start division x/y
13     GETQX result         ' Get quotient
14     GETQY remainder      ' Get remainder

```

Here's a complete multiply example:

```

1  ' Simple 16x16->32 multiply (2 clocks)
2      MOV    x, #123
3      MOV    y, #456
4      MUL    x, y          ' Result: 123 * 456 = 56088 in x
5      ' Uses lower 16 bits of each operand!
6
7  ' For full 32x32->64 multiply, use CORDIC:
8      QMUL   x, y          ' Start multiply (uses full 32 bits)
9      ' ... 55 clocks of other work ...
10     GETQX  low           ' Lower 32 bits of result
11     GETQY  high          ' Upper 32 bits of result

```

Uff! In the old days, we'd write loops for this. Now hardware does it!

## Logic Operations

Your Boolean friends:

```

1      AND    x, mask      ' x = x AND mask
2      OR     x, bits      ' x = x OR bits
3      XOR    x, toggle    ' x = x XOR toggle
4      NOT    x            ' x = NOT x (XOR with $FFFFFFFF)
5
6  ' Bit manipulation
7      BITL   x, #5        ' Clear bit 5 of x
8      BITH   x, #5        ' Set bit 5 of x
9      BITNOT x, #5        ' Toggle bit 5 of x
10     TESTB  x, #5 wc     ' Test bit 5, result in C flag

```

## Shifting and Rotating

Moving bits around:

```

1      SHL    x, #3        ' Shift left 3 bits
2      SHR    x, #3        ' Shift right 3 bits
3      SAR    x, #3        ' Arithmetic shift right (signed)
4      ROL    x, #3        ' Rotate left 3 bits
5      ROR    x, #3        ' Rotate right 3 bits

```

*continues on next page →*

```

6
7 ' Variable shifts (amount in register)
8     SHL    x, y        ' Shift x left by y bits
9
10 ' Fancy ones
11     REV    x          ' Reverse bit order (!! )
12     MERGEB x         ' Merge bits, not bytes (inverse of SPLITB)

```

## Flow Control: Jump!

### Unconditional Jumps

```

1     JMP    #target    ' Jump to target
2     JMP    target    ' Jump to address in target register
3
4 ' Relative jumps
5     JMP    #$$-4     ' Jump back 4 longs (addresses)
6     JMP    #$$+8     ' Jump forward 8 instructions

```

### Conditional Execution (The Magic)

Here's where PASM2 gets beautiful. ANY instruction can be conditional:

```

1 if_z  ADD    x, #1    ' Only add if Z flag set
2 if_nz ADD    x, #1    ' Only add if Z flag clear
3 if_c  ADD    x, #1    ' Only add if C flag set
4 if_nc ADD    x, #1    ' Only add if C flag clear

```

The basic conditions:

Condition	Meaning
<code>if_z</code>	If Z flag set (result was zero)
<code>if_nz</code>	If Z flag clear (result not zero)
<code>if_c</code>	If C flag set (carry/borrow occurred)
<code>if_nc</code>	If C flag clear
<code>if_c_and_z</code>	If both C and Z set
<code>if_c_or_z</code>	If either C or Z set
<code>if_c_eq_z</code>	If C equals Z
<code>if_c_ne_z</code>	If C not equal to Z

And the comparison conditions (use after **CMP**):

Condition	Meaning
<code>if_a</code>	If above (unsigned greater)
<code>if_b</code>	If below (unsigned less)
<code>if_ae</code>	If above or equal
<code>if_be</code>	If below or equal

## The Call/Return Dance

```

1      CALL    #subroutine    ' Call subroutine
2      RET                                ' Return from subroutine
3
4  ' But here's the P2 twist - CALL uses internal stack
5  subroutine
6      ' Do something useful
7      RET                                ' Returns to instruction after CALL
8
9  ' You get 8 levels of hardware stack!
```

## The `_RET_` Prefix: Return With Benefits

Here's a clever trick the P2 offers. What if you could execute an instruction *and* return from a subroutine in one go? That's exactly what the `_RET_` prefix does.

```

1 ' Normal way: Two instructions
2 add_and_return
3     ADD     x, y           ' Do the add
4     RET                               ' Then return (RET is a ~4-cycle branch)
5
6 ' _RET_ way: One instruction!
7 add_and_return
8     _ret_   ADD     x, y   ' Add AND return (saves 2 cycles)

```

The `_RET_` prefix says: “Execute this instruction, then return.” It’s like getting a free return ticket with your instruction. The add happens, flags get set normally, and then—pop!—you’re back at the caller.

### When does `_RET_` NOT return?

Here’s the catch: if the instruction itself branches, no return happens. The branch wins:

```

1     _ret_   JMP     #somewhere   ' JMP wins - no return
2     _ret_   CALL    #helper     ' CALL wins - no return
3     _ret_   DJNZ   count, #loop  ' Branch? No return. Zero? Return!

```

That last one is interesting! If `count` isn’t zero, `DJNZ` branches and no return. But when `count` hits zero, no branch occurs, so you get your return. Clever, right?

### One-Instruction Subroutines

This is where `_RET_` really shines:

```

1 ' Toggle pin 0 - entire subroutine is ONE instruction!
2 toggle_led
3     _ret_   DRVNOT #0           ' Toggle and return
4
5 ' Read all inputs - also just one instruction
6 read_inputs
7     _ret_   MOV     result, ina  ' Copy INA and return
8
9 ' Usage:
10    CALL    #toggle_led         ' Blink!
11    CALL    #read_inputs        ' result now has INA

```

A normal subroutine needs at least two instructions (the work + `RET`). With `_RET_`, you can have genuinely single-instruction subroutines. Your code gets smaller and faster.

## The Medicine: `_RET_` Quick Reference

Pattern	What Happens
<code>_ret_ add x, y</code>	ADD executes, then return
<code>_ret_ jmp #label</code>	JMP executes, NO return (branch wins)
<code>_ret_ djnz n, #loop</code>	If $n > 0$ : branch, no return. If $n = 0$ : no branch, return

**Important:** Unlike `RET wcz`, the `_RET_` prefix does NOT restore C and Z flags from the stack. If you need flag restoration, use the regular `RET wcz` instruction.

## Labels: Naming Your Places

You've been using labels throughout this chapter without us properly introducing them. How rude of me! Let's fix that.

### Global Labels: The Big Signposts

A global label is just a name at the start of a line:

```

1 DAT          ORG
2
3 send_byte    RDBYTE x, ptr      ' Global label
4              WYPIN  x, tx_pin
5              RET
6
7 receive_byte TESTP  rx_pin  wc  ' Another global label
8              RDPIN  x, rx_pin
9              RET

```

Global labels are visible everywhere in your DAT block. You can jump to them, **CALL** them, reference them from Spin2 - they're your main signposts.

### Local Labels: The Little Helpers

But here's a problem. What if every routine needs a loop? You can't have two labels called `loop` - the assembler would be terribly confused.

Enter local labels. Prefix a name with a dot (`.`) and it becomes local:

```

1 DAT          ORG
2

```

*continues on next page →*

```

3 send_byte      RDBYTE x, ptr
4 .loop         TESTP  tx_pin  wc   ' Local to send_byte
5             if_nc  JMP    #.loop
6             WYPIN  x, tx_pin
7             RET
8
9 receive_byte   TESTP  rx_pin  wc   ' New scope begins here
10            if_nc  JMP    #.wait
11 .wait        TESTP  rx_pin  wc   ' Local to receive_byte
12            if_nc  JMP    #.wait
13            RDPIN  x, rx_pin
14 .loop        SHR    x, #24      ' Different .loop, OK!
15            RET

```

Each global label starts a new “scope”. The `.loop` under `send_byte` is completely separate from the `.loop` under `receive_byte`. You can reuse `.loop`, `.done`, `.retry`, `.exit` to your heart’s content.

## The Colon Alternative

You might also see local labels with a colon prefix:

```

1 :loop         DJNZ   count, #:loop ' Same as .loop

```

Both `:` and `.` work identically. I prefer the dot - it’s what modern convention has settled on - but you’ll see both in the wild.

## Reference Operators: Finding Your Labels

When you reference a label, you need to tell the assembler what you want:

```

1 ' In COG code (after ORG):
2     JMP    #my_routine    ' # = immediate COG address
3     CALL  #.helper       ' # works for local labels too
4     MOV   x, #data_table  ' Get COG address of data
5
6 ' For hub addresses (used with Spin2):
7     MOV   ptr, @hub_data  ' @ = hub address of label

```

The `#` means “immediate value” - use this for jumps and calls within COG code. The `@` means “hub address” - use this when passing addresses to Spin2 or for hub memory operations.

## Scope Boundaries: When Local Labels Reset

Here's the rule: **every global label or data definition starts a new local scope.**

```

1 func_a      MOV    x, #1      ' Scope #1 begins
2 .loop      DJNZ   x, #.loop   ' .loop in scope #1
3
4 data_block  LONG   0, 0, 0, 0 ' Scope #2 begins (data!)
5
6 func_b      MOV    y, #2      ' Scope #3 begins
7 .loop      DJNZ   y, #.loop   ' .loop in scope #3, OK!
8 .done      RET

```

This is wonderfully useful - your utility routines can all use `.loop` and `.done` without stepping on each other's toes.

## The Medicine: Quick Reference

What	Syntax	Example
Global label	name	my_routine
Local label	.name or :name	.loop, :done
Jump to label	#label	jmp #.loop
Hub address	@label	mov ptr, @data

## Common Gotchas

1. **Forgetting the dot:** `loop` is global, `.loop` is local. If you accidentally create a global `loop`, you'll get conflicts.
2. **Scope surprise:** Data definitions (`LONG`, `WORD`, `BYTE`) also start new scopes. If you put data between two parts of a routine, your local labels won't work!
3. **The 30-character limit:** For compatibility with all tools, keep label names under 30 characters. `this_is_a_really_long_label_name` might cause trouble.

## Data in DAT Blocks: Your Program's Pantry

Speaking of data definitions starting new scopes... we should probably talk about how to actually declare data! You've seen snippets like `counter LONG 0` scattered through our examples, but there's a whole world of data declaration waiting for you.

## The Three Sizes

Just like Goldilocks, you have three choices:

```

1 DAT
2 my_byte      BYTE    $FF          ' 1 byte (8 bits)
3 my_word      WORD    $1234        ' 2 bytes (16 bits)
4 my_long      LONG    $DEADBEEF    ' 4 bytes (32 bits)

```

When do you use each? Well:

- **BYTE** for characters, small counters, flags, or when every byte of memory counts
- **WORD** for medium values, 16-bit peripherals, or when **BYTE** is too small but **LONG** is wasteful
- **LONG** for everything else - addresses, large numbers, and “I don’t want to think about it”

If you’re unsure, use **LONG**. Memory is cheap, debugging overflow errors is not.

## Multiple Values on One Line

Here’s a convenience - you can list multiple values after a single type:

```

1 DAT
2 primes      LONG    2, 3, 5, 7, 11, 13, 17, 19
3 gpio_pins   BYTE    16, 17, 18, 19, 20, 21
4 message     BYTE    "Hello, P2!", 0      ' String with null term

```

The assembler just lays them out consecutively in memory. That string? It’s just bytes - each character followed by a zero at the end.

## Arrays: The Repetition Trick

Need 100 bytes of zeros? Don’t type them all out:

```

1 DAT
2 buffer      BYTE    0[100]          ' 100 zero bytes
3 lookup_table WORD    $FFFF[64]      ' 64 words, all $FFFF
4 scratch_pad LONG    0[32]          ' 32 longs of zeros

```

The [count] syntax repeats the value. This is your friend for buffers, tables, and anywhere you need initialized storage.

You can even combine values and repetition:

```

1 DAT
2 mixed_init      BYTE    $AA, $BB, 0[10], $CC, $DD
3                '      ^      ^      ^      ^
4                '      Two vals, then 10 zeros, then two more

```

## BYTEFIT and WORDFIT: The Safety Net

Here's a subtle trap. What if you accidentally write:

```

1 DAT
2 oops            BYTE    1000          ' 1000 won't fit in a byte!

```

The assembler will silently truncate 1000 to 232 (the low 8 bits). Your program will run, but with wrong values. Debugging that is no fun at all.

Enter the safety net:

```

1 DAT
2 safe_byte      bytefit 100, 200, 255  ' OK - all fit in a byte
3 danger_byte    bytefit 100, 200, 300  ' ERROR! 300 > 255, error!
4
5 safe_word      wordfit 1000, 50000    ' OK - all fit in a word
6 danger_word    wordfit 1000, 70000    ' ERROR! 70000 > 65535

```

Use **BYTEFIT** and **WORDFIT** when you want the assembler to check that your constants actually **FIT**. It's like having a compiler catch your mistakes before they become 3 AM debugging sessions.

## Alignment: Keeping Things Tidy

The P2 is quite forgiving about alignment - it can read a **LONG** from any address. But aligned access is faster and cleaner. Sometimes you need to force alignment:

```

1 DAT
2 some_bytes     BYTE    1, 2, 3        ' 3 bytes
3               ALIGNW                    ' Align to word boundary
4 next_word      WORD    $1234          ' Now properly aligned
5
6 more_bytes     BYTE    "ABC"          ' 3 bytes
7               ALIGNL                    ' Align to long boundary
8 next_long      LONG    $DEADBEEF      ' Now on a 4-byte boundary

```

When does alignment matter? Mostly when you're:

- Mixing data sizes in the same DAT block
- Creating structures that Spin2 code will access
- Optimizing for maximum hub access speed

If you're just starting out, don't worry about it. Add alignment when you hit problems.

## A Complete Example

Let's put it all together:

```

1  DAT          ORG
2
3  ' Your code goes here
4  entry        MOV    ptra, ##buffer_addr
5                MOV    count, #BUFFER_SIZE
6  .fill        WRBYTE fill_value, ptra++
7                DJNZ   count, #.fill
8                JMP    #$
9
10 ' Constants (read-only, effectively)
11 fill_value   BYTE   $55
12 BUFFER_SIZE  LONG   256
13
14 ' Lookup tables
15                ALIGNL
16 sin_table    WORD   0, 1608, 3212, 4808, 6393   ' Sine table
17                WORD   7962, 9512, 11039, 12540   ' ... continues
18
19 ' Working storage
20                ALIGNL
21 buffer_addr  LONG   0                          ' Set by Spin2 at startup
22 temp        LONG   0
23 result      LONG   0
24
25 ' Reserve uninitialized space
26                ALIGNL
27 scratch     RES    16                          ' Reserve 16 longs

```

Notice the pattern: code first, then constants, then working storage, then reserved space. This keeps things organized and makes your DAT block readable.

## The Medicine: Quick Reference

Declaration	Meaning	Example
<code>byte val</code>	8-bit value	<code>byte \$FF</code>
<code>word val</code>	16-bit value	<code>word \$1234</code>
<code>long val</code>	32-bit value	<code>long \$DEADBEEF</code>
<code>val[n]</code>	Repeat n times	<code>byte 0[100]</code>
<code>bytefit</code>	Byte with range check	<code>bytefit 100, 200</code>
<code>wordfit</code>	Word with range check	<code>wordfit 1000, 50000</code>
<code>alignw</code>	Align to word	(no value)
<code>alignl</code>	Align to long	(no value)
<code>res n</code>	Reserve n longs	<code>res 16</code>

## Common Gotchas

1. **Forgetting the label:** Every piece of data needs a label if you want to reference it. Anonymous data just wastes space.
2. **String termination:** `BYTE "Hello"` doesn't include a null terminator. Add `, 0` if you need one!
3. **RES is in longs:** `RES 10` reserves 10 *longs* (40 bytes), not 10 bytes. This trips up everyone at least once.
4. **Alignment after RES:** `RES` doesn't affect alignment. If you need alignment after reserved space, add an explicit `ALIGNL` or `ALIGNW`.

## Including External Files: FILE

Sometimes you have binary data sitting in a file - a font bitmap, a sound sample, a pre-computed lookup table. Rather than manually converting it to hex values (ugh!), you can pull it straight in:

```

1 DAT
2 font_data      file    "myfont.bin"      ' Import entire file
3 sound_sample  file    "beep.raw"        ' Raw audio data
4 lut_table     file    "precalc.dat"      ' Pre-computed values

```

The assembler reads the file at compile time and drops its raw bytes right into your DAT block. The label gives you a way to reference where the data starts.

## Where does it search?

1. Same folder as your source file
2. Library paths (if configured)

### Practical example - embedded bitmap:

```

1 DAT          ORG
2 entry        MOV    ptr, ##@splash_screen
3              ' ... display routine using ptr
4
5 ' The image data lives right here in your code
6 splash_screen file  "logo_128x64.bin"  ' 1024 bytes of pixel data

```

No conversion scripts, no copy-paste errors. Just reference the file and it's part of your program. The icing on the cake? If you update the file, recompile, and your program has the new data.

### String and Data Generation Methods

Beyond manually typing values, you have some helpers for creating data:

#### @“text” - Inline String Address

Need a string address without declaring a separate label?

```

1      MOV    ptr, @"Hello!"      ' ptr points to "Hello!" in hub
2      CALL  #print_string
3
4      MOV    ptr, @"Error: "     ' Another string, inline
5      CALL  #print_string

```

The @“text” syntax creates the string in hub memory and gives you its address. It's like an anonymous label for a string constant. Each unique string gets stored once, even if you reference it multiple times.

#### STRING(“text”) and LSTRING(“text”)

These work similarly but in different contexts:

```

1 ' In Spin2 code (not PASM):
2 DEBUG(STRING("Temperature: "))  ' Zero-terminated string address
3 DEBUG(LSTRING("Status"))       ' Length byte first, then string

```

`STRING()` returns the hub address of a zero-terminated string - same as what C programmers expect. `LSTRING()` puts a length **BYTE** at the front, which is handy when you need to know the string length without scanning for null.

### **BYTE[], WORD[], LONG[] - Data Arrays**

In Spin2, you can create inline data arrays:

```

1 ' Spin2 examples:
2 lookup := BYTE[10, 20, 30, 40, 50]      ' Returns address of byte array
3 config := LONG[$DEAD_BEEF, $CAFE_BABE]
```

These are primarily Spin2 features, but they generate hub data that your PASM code can access if you know the addresses.

### **The Pattern:**

Method	Result	Use Case
<code>file "name"</code>	Raw binary data	Images, audio, lookup tables
<code>@"text"</code>	String address	Quick inline strings in PASM
<code>STRING("text")</code>	Zero-terminated string address	Spin2 string constants
<code>LSTRING("text")</code>	Length-prefixed string	When you need length upfront

### **The Flags: C and Z (and Q!)**

Flags are your friends. They remember things:

```

1 ' Z flag - was the result zero?
2     SUB    x, y wz      ' Set Z if x-y equals zero
3 if_z     JMP    #equal   ' Jump if they were equal
4
5 ' C flag - did we carry/borrow?
6     ADD    x, y wc      ' Set C if addition overflowed
7 if_c     JMP    #overflow ' Handle overflow
8
9 ' Both at once!
10     CMP    x, y wcz     ' Compare and set both flags
11 if_a     JMP    #x_greater ' Jump if x > y (unsigned)
```

The Q flag is special - it's used by CORDIC operations (Chapter 7).

## Special Instructions That Will Blow Your Mind

### SKIP - The Instruction Skipper

```

1      SKIP    ##%00001010    ' LSB first: 1=skip, 0=execute
2      ADD     x, #1          ' Executed (bit 0 = 0)
3      ADD     y, #1          ' Skipped! (bit 1 = 1)
4      ADD     z, #1          ' Executed (bit 2 = 0)
5      SUB     a, #1          ' Skipped! (bit 3 = 1)
6      SUB     b, #1          ' Executed (bit 4 = 0)
7      ' ... bit N controls the Nth instruction after SKIP

```

This is like having conditional execution on steroids!

### REP - Hardware Loops

```

1      REP     #4, #5          ' Repeat next 4 instructions 5 times
2      ADD     x, #1
3      SUB     y, #1
4      ROL     z, #1
5      ROR     w, #1
6      ' These 4 instructions execute 5 times total
7      ' No loop overhead!

```

### ALTD/ALTS - Instruction Modification

```

1      ' Modify the next instruction's destination
2      MOV     index, #10
3      ALTD    index, #array  ' Next instruction's dest = array+10
4      MOV     0-0, value     ' Actually moves to array[10]!

```

This replaces self-modifying code from P1. Much cleaner!

#### Sidetrack

**Hub-Exec Note:** All ALTx instructions — **ALTI**, **ALTS**, **ALTD**, **ALTR**, **ALTB**, **ALTSN**, **ALTSB**, **ALTSW**, **ALTGN**, **ALTGB**, **ALTGW** — work identically in cog-exec and hub-exec modes. They act on the next pipelined instruction regardless of whether

*continues on next page →*

it came from cog/LUT memory or the hub-prefetch FIFO. So when you graduate to hub execution in Chapter 10, every ALTx trick you learn here still works.

## Real-World Example: Fast Memory Copy

Let's combine what we've learned:

```
1  ' Fast block copy using REP
2  fast_copy
3      MOV    ptra, ##source_addr    ' Source pointer
4      MOV    ptrb, ##dest_addr     ' Destination pointer
5
6      REP    #2, ##256              ' Repeat 256 times
7      RDLONG temp, ptra++          ' Read and increment
8      WRLONG temp, ptrb++          ' Write and increment
9      ' 1024 bytes (256 longs) copied with no loop overhead!
10
11 temp    LONG    0
```

### The Medicine Cabinet

Feeling overwhelmed? Here's your simplified prescription:

#### Minimum Instructions to Know

```

1  ' Moving data
2      MOV      dest, source  ' Copy data
3
4  ' Math
5      ADD      dest, source  ' Addition
6      SUB      dest, source  ' Subtraction
7
8  ' Logic
9      AND      dest, source  ' AND operation
10     OR       dest, source  ' OR operation
11
12  ' Flow
13     JMP      #label       ' Jump
14     CALL     #label       ' Call subroutine
15     RET
16
17  ' Flags
18     CMP      x, y wcz     ' Compare and set flags
19     if_z    JMP      #label       ' Conditional jump

```

Master these 10 instructions and you can write real programs!

## Common Gotchas

### 1. Immediate values:

- #value for 9-bit immediates (0-511)
- ##value for 32-bit immediates
- Forgetting # uses the register at that address!

### 2. Flag confusion:

- wz sets Z flag, wc sets C flag, wcz sets both
- No flag update means flags unchanged

### 3. PTR A/PTR B are special:

```

1   RDLONG  x, ptra++      ' Read and auto-increment
2   RDLONG  x, ++ptra     ' Pre-increment then read
3   RDLONG  x, ptra--     ' Read and auto-decrement
4   RDLONG  x, ptra[5]    ' Read from ptra + 5*4

```

#### 4. Address confusion:

- COG addresses are in longs (0-511)
- Hub addresses are in bytes (0-524287)

## Your Turn: Experiments

### Experiment 1: Conditional Counter

Count up if button pressed, down if not:

```

1       ORG      0
2
3  .loop  TESTP   #BUTTON_PIN wz ' Test button
4  if_c   ADD     counter, #1    ' Increment if pressed
5  if_nc  SUB     counter, #1    ' Decrement if not
6
7       WRLONG   counter, ##HUB_ADDR ' Display count
8       WAITX   ##1_000_000
9       JMP     #.loop
10
11 counter LONG  0

```

### Experiment 2: Pattern Matcher

Find a pattern in data:

```

1       ORG      0
2
3       MOV     pattern, ##$DEADBEEF
4       MOV     ptra, ##data_start
5
6  .search RDLONG value, ptra++
7       CMP     value, pattern wz
8  if_z   JMP     #.found

```

*continues on next page →*

```

 9          CMP    ptra, ##data_end wcz
10 if_b     JMP    #.search
11          JMP    #.not_found
12
13 .found   ' Pattern found!
14          DRVH   #SUCCESS_LED
15          JMP    #$
16
17 .not_found
18          DRVH   #FAIL_LED
19          JMP    #$

```

### Experiment 3: Speed Test

Compare multiply methods:

```

 1 ' Method 1: Hardware multiply
 2     GETCT    start_time
 3     MUL     x, y
 4     GETCT    end_time
 5     SUB     end_time, start_time
 6     ' Result: 2 clocks!
 7
 8 ' Method 2: Shift and add (old school)
 9     GETCT    start_time
10     ' ... shift/add loop here
11     GETCT    end_time
12     ' Result: Many more clocks!

```

#### Sidetrack

##### Why PASM2 Is Different

Most assembly languages are thin wrappers over hardware. PASM2 is different - it's designed for humans:

1. **Symmetry:** Every instruction can use every addressing mode
2. **Orthogonality:** Features combine predictably
3. **Conditional everything:** Not just jumps, ANY instruction
4. **No special cases:** General-purpose registers, no accumulator

This isn't accident - it's philosophy. The P2 was designed to make assembly programming pleasant.

## What We've Learned

- Instruction anatomy and structure
- Basic data movement and math
- Hardware multiply and divide (!)
- Conditional execution on any instruction
- Special instructions (**SKIP**, **REP**, **ALT\***)
- Flag operations and testing
- Why PASM2 is human-friendly

## Coming Up Next

Chapter 4, “The Hub Connection”, explores:

- Reading and writing hub memory
- The FIFO and fast block transfers
- Hub execution mode
- Sharing data between COGs

You now speak basic PASM2. Time to learn how COGs communicate!

**Have Fun!** Remember, PASM2 isn't like other assembly languages - it's actually enjoyable!

# Chapter 4: The Hub Connection

*How COGs share and care*

## The Hook: Instant Communication

```

1 ' COG 1: Leave a message
2     WRLONG  ##$DEADBEEF, ##$1000
3
4 ' COG 2: Get the message
5     RDLONG  message, ##$1000
6     ' message now contains $DEADBEEF!
```

That's it - COGs talking through hub memory. But there's so much more...

## Reading from Hub

The basics are simple:

```

1     RDBYTE  value, hubaddr    ' Read 1 byte
2     RDWORD  value, hubaddr    ' Read 2 bytes (word)
3     RDLONG  value, hubaddr    ' Read 4 bytes (long)
4
5 ' With PTR A/PTR B magic
6     RDLONG  value, ptr++      ' Read and increment pointer
7     RDLONG  value, ++ptr     ' Increment then read
8     RDLONG  value, ptr[5]    ' Read from ptr + 5*4
```

## Writing to Hub

Just as easy:

```

1     WRBYTE  value, hubaddr    ' Write 1 byte
2     WRWORD  value, hubaddr    ' Write 2 bytes
3     WRLONG  value, hubaddr    ' Write 4 bytes
4
5 ' The mighty block transfer
```

*continues on next page →*

*↔ continued from previous page*

```

6         SETQ    #16-1           ' Transfer 16 longs
7         RDLONG  buffer, hubaddr  ' Reads 16 longs in one go!

```

## The FIFO Pipeline

Here's where P2 gets serious about speed:

```

1 ' Start the FIFO
2     RDFAST  #0, ##data_start  ' Start fast read
3
4 ' Now read at maximum speed
5 .loop  RFLONG  value           ' Read from FIFO
6     ' Process value
7     DJNZ   count, #.loop      ' Decrement and jump if not zero
8
9 ' No hub timing worries - FIFO handles it all!

```

## Real-World Example: Video Buffer

```

1 ' Fast screen clear using block transfer
2 ' Note: SETQ/SETQ2 maximum is 511 (for 512 longs)
3 ' For larger fills, we loop in 512-long chunks
4 clear_screen
5     MOV     hub_ptr, ##screen_buffer
6     MOV     chunks, ##640*480/512  ' Number of 512-long chunks
7     MOV     color, ##$00_00_00_00  ' Black
8
9 .loop  SETQ    #512-1           ' Transfer 512 longs (max)
10     WRLONG  color, hub_ptr      ' Fill this chunk
11     ADD     hub_ptr, ##512*4     ' Advance 512 longs (2KB)
12     DJNZ   chunks, #.loop
13     ' Full screen cleared with minimal loop overhead!

```

### The Medicine Cabinet

#### Simple hub access pattern:

```

1  ' Just use PTRA for everything
2      MOV    ptra, ##hub_address
3      RDLONG value, ptra++
4      ' That's all you really need!
```

## Your Turn: Experiments

Now you know the moves. Try these to make them stick:

1. **Round-trip:** Write a long to hub at address \$1000, then read it back into a different register. Verify the values match using a comparison and a flag.
2. **Block fill:** Use **SETQ** + **WRLONG** to fill 64 consecutive longs in hub memory with the value \$DEAD\_BEEF. Then read them back one by one and confirm.
3. **FIFO race:** Use **RDFAST** to stream a sequence of longs from hub. Compare the throughput to a loop of individual **RDLONG** instructions. The difference should astonish you.
4. **Pointer arithmetic:** Set **PTRA** to a hub address, then use **RDLONG** `val, ptra[3]` to read the third **LONG** ahead — without moving the pointer. Experiment with `++`, `--`, and indexed forms.

## Common Gotchas

Before you pull your hair out debugging hub access:

1. **Forgetting the ##** — Hub addresses are 20-bit. A bare `#address` only encodes 9 bits, so you'll hit the wrong memory. Always use `##` for hub addresses outside the 0–511 range.
2. **Long-aligned access** — **RDLONG** and **WRLONG** work on **LONG** boundaries. If your address isn't a multiple of 4, the P2 silently masks the low bits. For byte-precise access, use **RDBYTE** / **WRBYTE**.
3. **SETQ block size** — **SETQ** `#N-1` transfers `N` longs (not `N-1`). The `-1` is because the encoded field is “count minus one.” Off-by-one bugs love this one.
4. **PTRA vs. PTRB** — Both are pointer registers, but they're independent. Don't assume **PTRA** holds anything if you've been working with **PTRB**.
5. **Hub-exec vs. cog-exec timing** — In cog-exec, **RDLONG** is 9–16 cycles. In hub-exec, it's 9–26 cycles. Inner loops that hammer hub should run from COG memory when possible.

## What We've Learned

Look at all the ground we've covered:

- How to read and write hub memory from COG code
- Why pointer registers (PTRA, PTRB) exist and when to use them
- How the FIFO pipeline accelerates sequential hub reads
- Block transfers with **SETQ** for moving big chunks at once
- The pattern for clearing or filling buffer-sized regions

You now have the entire COG Hub vocabulary at your disposal.

## Coming Up Next

In Chapter 5, we'll unleash the P2's mathematical muscle:

- Hardware multiply and divide that don't require library calls
- 64-bit arithmetic without sweating
- A preview of the CORDIC coprocessor (it gets a full chapter later)
- Fixed-point math for fast trigonometry

But first, take a break. Hub access is one of those topics where letting it sit a day helps it lock in. Try one of the experiments above tomorrow.

**Have Fun!** And remember — the FIFO is doing real work even when your code looks idle. Trust it.

# Chapter 5: Mathematics Unleashed

*Hardware multiply and divide - finally!*

## The Hook: Hardware Multiply

```

1      MUL      x, y          ' 16x16->32 bit unsigned multiply
2      ' Result in x (lower 16 bits of each operand used)
3
4      ' For full 32x32->64 bit multiply, use CORDIC:
5      QMUL     x, y          ' Start 32x32->64 multiply
6      ' ... other work (55 clocks) ...
7      GETQX    low          ' Get lower 32 bits
8      GETQY    high         ' Get upper 32 bits

```

Remember doing this with shifts and **ADDS**? Those days are over!

## The Multiplication Revolution

```

1  ' Unsigned multiply
2      MUL      result, value  ' result = low 32 bits
3
4  ' Signed multiply
5      Muls     result, value  ' Signed version
6
7  ' Scale by a power of two
8      SHR      result, #1     ' Scale by 0.5 (divide by 2)

```

## Division Without Tears

```

1  ' Start division
2      QDIV     dividend, divisor ' Start the operation
3
4  ' Get results (takes 55 clocks)
5      GETQX    quotient        ' Get quotient
6      GETQY    remainder       ' Get remainder
7

```

*continues on next page →*

```

8 ' Fractional division
9     QFRAC    numerator, denominator
10    GETQX    fraction           ' 32-bit fraction

```

## 64-Bit Operations

```

1 ' 64-bit add
2     ADD     low1, low2 wc
3     ADDX    high1, high2
4
5 ' 64-bit multiply (uses CORDIC)
6     QMUL    x, y                ' Start 32x32->64 multiply
7     ' ... 55 clocks ...
8     GETQX   low                 ' Lower 32 bits
9     GETQY   high                ' Upper 32 bits

```

## Real-World Example: Fixed-Point Math

```

1 ' 16.16 fixed point multiply (uses CORDIC for full precision)
2 fixed_mul
3     QMUL    a, b                ' Start 32x32->64 unsigned multiply
4     ' ... 55 clocks (do other work) ...
5     GETQX   low                 ' Lower 32 bits
6     GETQY   high                ' Upper 32 bits
7     ' Extract middle 32 bits for 16.16 result:
8     SHL     high, #16           ' Upper 16 bits of result
9     SHR     low, #16            ' Lower 16 bits of result
10    OR      a, low              ' Combine for 16.16 format
11    OR      a, high

```

### The Medicine Cabinet

#### Quick math reference:

- **MUL** D, S — unsigned  $16 \times 16 \rightarrow 32$
- **MULS** D, S — signed  $16 \times 16 \rightarrow 32$
- **QMUL** D, S — full  $32 \times 32 \rightarrow 64$  (read via **GETQX**/**GETQY** after 55 clocks)
- **QDIV** D, S — full 32-bit divide (read via **GETQX** quotient / **GETQY** remainder after 55 clocks)
- **QFRAC** D, S — fractional divide (returns 32-bit fraction in **GETQX**)
- 64-bit add: **ADD** + **ADDX** chained with **WC**

For everyday integer work, **MUL**/**MULS** are 2 clocks and you're done. For precision (full 64-bit results, fixed-point math, signed division), **QMUL**/**QDIV** route through the CORDIC and pay 55 clocks — but they don't block the COG, so you can interleave other work.

## Your Turn: Experiments

Stretch your math muscles:

1. **Compute the average:** Read 8 longs from hub, sum them with **ADD**, then divide by 8 using a **SHR** (or **QDIV** if you want exact). Compare both approaches.
2. **Fractional reciprocal:** Use **QFRAC** to compute  $2^{-32} / x$  for various  $x$  values. You've just built a hardware reciprocal table.
3. **Pipeline overlap:** Start a **QMUL**, do 30+ clocks of other work (compute something else, update a counter), then **GETQX**/**GETQY**. Measure total cycles vs. doing the multiply blocking-style.
4. **64-bit counter:** Build a 64-bit increment loop using **ADD** + **ADDX**. Add to it in a tight loop and watch the high **WORD** advance when the low one wraps.

## Common Gotchas

The math instructions hide a few traps:

1. **MUL is unsigned, MULS is signed** — Using the wrong one on negative values gives spectacular nonsense. When in doubt, **MULS**.
2. **MUL gives only low 32 bits** — For the full 64-bit result, you must use **QMUL** + **GETQX**/**GETQY**. The high **WORD** is silently discarded by **MUL**.
3. **GETQX/GETQY block until ready** — They wait for the CORDIC. If you **CALL** them too early, your COG stalls. If you **CALL** them later than necessary, you've wasted cycles. The sweet spot is starting the CORDIC, doing exactly 55 clocks of other work, then reading.

4. **Don't queue a new CORDIC op while one is pending** — A second **QMUL/QDIV/QFRAC** before reading results overwrites the queue. Use **GETQX/GETQY** first.

## What We've Learned

A whole arithmetic library, in your back pocket:

- Hardware multiply (**MUL, MULS**) in just 2 clocks
- Division and modulo via **QDIV + GETQX/GETQY**
- Fractional division via **QFRAC** for fast reciprocals
- 64-bit chains using **ADD/ADDX, SUB/SUBX** with **WC**
- Fixed-point patterns for 16.16 math

You no longer have to write multiply-by-shift-and-add. Those days really are over.

## Coming Up Next

In Chapter 6 we'll meet the P2's most quietly powerful feature: **conditional execution**. Every instruction can be conditional — no branches needed, deterministic timing preserved. We'll cover:

- The **C** and **Z** flags and how to update them
- Building complex conditions with **WC, WZ, WCZ** combined with **IF\_x** prefixes
- **SKIP** and **SKIPF** for skipping patterns of instructions

If conditional execution doesn't change how you think about flow control, nothing will.

**Have Fun!** And remember — the CORDIC is a coprocessor with infinite patience. Use it.

# Chapter 6: Flags and Decisions

*Making choices at machine speed*

## The Hook: Any Instruction Can Be Conditional

```

1          CMP    x, y wcz          ' Compare x and y
2 if_a     MOV    result, x         ' If x > y, result = x
3 if_be    MOV    result, y         ' If x <= y, result = y
4          ' Max function in 3 instructions!
```

## The C and Z Flags

```

1 ' Z Flag - Zero detection
2          SUB    x, y wz           ' Z=1 if x equals y
3 if_z     JMP    #equal            ' Jump if equal
4
5 ' C Flag - Carry/Borrow
6          ADD    x, y wc           ' C=1 if overflow
7 if_c     JMP    #overflow         ' Handle overflow
```

## Complex Conditions

```

1 ' Combining flags
2          CMP    x, y wcz
3 if_a     JMP    #greater          ' x > y (unsigned)
4 if_b     JMP    #less             ' x < y (unsigned)
5 if_z     JMP    #equal           ' x == y
6
7 ' Signed comparisons
8          CMPS   x, y wcz         ' Signed compare
9 if_gt    JMP    #greater          ' x > y (signed)
10 if_lt   JMP    #less             ' x < y (signed)
```

## Skip Patterns - Conditional Blocks

```

1      SKIPF  pattern          ' Set skip pattern
2      ADD    x, #1            ' Maybe executed
3      SUB    y, #1            ' Maybe executed
4      MOV    z, #0            ' Maybe executed
5      ' Pattern determines what runs!
```

### The Medicine Cabinet

#### Conditional execution quick reference:

- **WC** — write the C flag with the instruction's carry-out
- **WZ** — write the Z flag based on result == 0
- **WCZ** — write both flags
- **IF\_C**, **IF\_NC** — execute next instruction only if C set / clear
- **IF\_Z**, **IF\_NZ** — same for Z
- **IF\_A**, **IF\_B**, **IF\_AE**, **IF\_BE** — unsigned comparisons after **CMP**
- **IF\_GT**, **IF\_LT**, **IF\_GE**, **IF\_LE** — signed comparisons after **CMPS**
- **IF\_E**, **IF\_NE** — equal / not equal (same as **IF\_Z/IF\_NZ** after compare)

Any instruction can carry a condition prefix and still take exactly 2 clocks. No branches, no pipeline flush, no surprise.

## Your Turn: Experiments

Practice the conditional mindset:

1. **Branchless absolute value:** Compute `result = ABS(x)` using **ABS** — then prove to yourself the same thing works branchless using **CMPS** + a conditional **NEG**.
2. **Saturating add:** Add two unsigned values; if the result overflows (**WC** sets C), clamp to `$FFFF_FFFF`. No jumps.
3. **Min/max:** Build a min function using **CMP** + conditional **MOV**. Then build max. Three instructions each.
4. **Skip-pattern selector:** Use **SKIPF** with a runtime bit pattern to selectively execute 4 different instructions based on a state **BYTE**. This is the building block for hand-coded jump tables.

## Common Gotchas

The conditional mindset has its own traps:

1. **Forgetting the effect** — `CMP x, y` with no flag effect doesn't update C or Z. You need **WC**, **WZ**, or **WCZ**. The default for **CMP** is no update.
2. **Signed vs. unsigned comparison** — **CMP** sets flags for unsigned semantics (**IF\_A/IF\_B**). **CMPS** sets them for signed (**IF\_GT/IF\_LT**). Mixing them produces baffling bugs near zero and at the high bit.
3. **One instruction follows the prefix** — `IF_C MOV x, y` conditionally runs *only* the **MOV**. The instruction after that runs unconditionally. To skip a block, use **SKIPF**.
4. **C flag inversion on subtract** — **SUB** sets C on *borrow*, not carry. So after `SUB a, b` `wc, C=1` means `a < b` (unsigned). Many P1 programmers expect the opposite — re-check.
5. **Skip patterns are LSB-first** — In **SKIPF**, bit 0 controls the next instruction, bit 1 the one after, etc. The numbering can feel backward; sketch it out.

## What We've Learned

You can now think like a P2 programmer:

- Every instruction can carry a condition prefix
- Flags are updated only when you ask (**WC**, **WZ**, **WCZ**)
- Unsigned and signed comparisons use different instructions (**CMP** vs. **CMPS**)
- Complex multi-way logic without a single jump using **IF\_x** prefixes
- **SKIP** and **SKIPF** for skipping arbitrary patterns of instructions

You've learned to write branchless code that preserves deterministic timing. That's the deepest part of P2 thinking.

## Coming Up Next

Chapter 7 reveals the CORDIC coprocessor in full — the math beast you glimpsed in Chapter 5. Trigonometry at the speed of logic gates:

- Rotate points in three instructions
- Polar Cartesian conversion
- Hardware square root and natural log
- Pipelined operations for graphics and signal processing

If you thought hardware multiply was nice, wait until you see what the CORDIC does.

**Have Fun!** And remember — every **JMP** you avoid is a pipeline flush you didn't pay for.

# Chapter 7: CORDIC Magic

*Trigonometry at the speed of logic gates*

## The Hook: Rotate a Point in 3 Lines

Let me show you something that, on most processors, would take a coffee break of instructions, a math library, and probably a tear or two:

```

1 ' Rotate point (x,y) by angle - that's it!
2     SETQ    y_coord      ' Set Y coordinate
3     QROTATE x_coord, angle ' Start rotation by angle
4     GETQX   new_x        ' Get rotated X (55 clocks later)
5     GETQY   new_y        ' Get rotated Y

```

Read that again. *Four lines*, and a 2D rotation is done. No floating-point library. No lookup tables. No iterative approximation. The P2 has a dedicated trigonometric coprocessor sitting next to every COG, just waiting for you to wake it up. You're about to learn how.

Let me show you something even more impressive:

```

1 ' Calculate sine and cosine simultaneously
2     QROTATE ##$7FFF_FFFF, angle ' D=radius (max), S=angle
3     GETQX   cosine             ' cos(angle) in 2.30 fixed pt
4     GETQY   sine               ' sin(angle) in 2.30 fixed pt
5     ' Both trig functions in 55 clocks total!

```

## What Just Happened?

CORDIC stands for COordinate Rotation DIgital Computer. It's a method invented in 1959 for calculating trigonometric functions using only shifts and **ADDS** - no multiplies needed! Each P2 COG has its own dedicated CORDIC unit built into the hardware.

Think of CORDIC as your mathematical co-processor that can:

- Rotate points around the origin
- Convert between rectangular and polar coordinates
- Calculate sine, cosine, tangent
- Compute square roots and magnitudes

- Find arctangent (angle between points)
- Even do logarithms and exponentials!

All of this in exactly 55 clock cycles. Every time. No variation.

## The CORDIC Pipeline - Your Mathematical Assembly Line

Here's the beautiful part: CORDIC operations are pipelined. While one calculation is running, you can start another. Let's see what that buys us:

```

1  ' Generate sine wave samples rapid-fire
2      MOV     angle, #0
3      MOV     count, #256
4
5  generate
6      QROTATE ##$7FFF_FFFF, angle    ' D=radius, S=angle
7      ADD     angle, ##$0100_0000    ' Increment angle (no wait!)
8
9      ' Do other work while CORDIC calculates
10     ADD     sample_ptr, #4
11     SUB     count, #1
12
13     GETQY   sample                    ' Get sine result
14     WRLONG  sample, sample_ptr        ' Store it
15
16     TJNZ    count, #generate          ' Test-jump-not-zero
17     ' Generated 256 samples with perfect overlap!

```

The pipeline means you're not really waiting 55 clocks - you're getting useful work done while CORDIC churns away in the background. Uff! That's free math!

## Core CORDIC Operations

### QROTATE - The Rotation Engine

Here's a subtle detail: CORDIC operations work on 2D coordinates (X, Y), but the **QROTATE** instruction only takes one coordinate directly. The solution? **SETQ** loads the Y coordinate into the Q register, then **QROTATE** takes X from its first operand. It's a two-instruction dance that becomes second nature:

```

1 ' Basic rotation: rotate point (x,y) by angle
2     SETQ    y                ' First: load Y into Q register
3     QROTATE x, angle        ' Then: X from operand, Y from Q
4     GETQX   new_x           ' Result: X' = X*cos() - Y*sin()
5     GETQY   new_y           ' Result: Y' = X*sin() + Y*cos()

```

The angle format is special: it's a 32-bit unsigned value where:

- \$0000\_0000 = 0 degrees
- \$4000\_0000 = 90 degrees
- \$8000\_0000 = 180 degrees
- \$C000\_0000 = 270 degrees
- \$FFFF\_FFFF = just under 360 degrees

This makes angle math incredibly easy - just use regular addition and subtraction!

### QVECTOR - From Rectangular to Polar

```

1 ' Convert (x,y) to polar (radius, angle)
2     QVECTOR x, y            ' X in D, Y in S (no SETQ for QVECTOR)
3     GETQX   radius         ' sqrt(x2 + y2)
4     GETQY   angle          ' atan2(y, x)

```

Perfect for:

- Finding distances between points
- Converting joystick input to angle/magnitude
- Radar and sonar applications

### The Power of 32-Bit Precision

CORDIC uses 32-bit precision throughout:

- Angles: 32 bits (0.00000008 degree resolution!)
- Coordinates: 32 bits signed
- Results: Full 32-bit or 64-bit when needed

## Real-World Example: Spinning a Sprite

Let's rotate a sprite around its center:

```

1  ' Rotate sprite vertices around center
2  ' (PTRA for vertex iteration - only PTRB support ++ post-increment.)
3  rotate_sprite
4      MOV    ptra, ##sprite_data
5      MOV    vertex_count, #4          ' 4 corners
6
7  next_vertex
8      RDLONG x, ptra++                ' Get X coordinate, advance pointer
9      RDLONG y, ptra++                ' Get Y coordinate, advance pointer
10
11     ' Center sprite at origin
12     SUB    x, center_x
13     SUB    y, center_y
14
15     ' Rotate by current angle
16     SETQ   y
17     QROTATE x, rotation_angle
18
19     ' While waiting, we can prepare
20     MOV    temp_x, center_x
21     MOV    temp_y, center_y
22
23     ' Get rotated coordinates
24     GETQX  x
25     GETQY  y
26
27     ' Translate back to position
28     ADD    x, temp_x
29     ADD    y, temp_y
30
31     ' Store rotated vertex (back up PTRB 8 to overwrite source longs)
32     SUB    ptrb, #8
33     WRLONG x, ptrb++
34     WRLONG y, ptrb++
35
36     DJNZ   vertex_count, #next_vertex
37

```

*continues on next page →*

```

38      ' Increment rotation for animation
39      ADD      rotation_angle, ##$0100_0000  ' ~1.4 deg (1/256 rotation)

```

## Your Turn: CORDIC Experiments

### Your Turn

**Your Turn:** Create a circular motion pattern

Starting code:

```

1      ORG      0
2
3      MOV      angle, #0
4      MOV      radius, ##100          ' 100 pixel radius
5
6 .loop  QROTATE radius, angle          ' D=X (radius), S=angle
7      ' Add code to:
8      ' 1. Get X,Y coordinates
9      ' 2. Add screen center offset
10     ' 3. Draw pixel at that position
11     ' 4. Increment angle
12     ' 5. Loop back to .loop

```

Goal: Make a dot trace a perfect circle on screen Hint: After **QROTATE**, use **GETQX**/getqy to get coordinates Success Check: Smooth circular motion, no gaps

### Your Turn

**Your Turn:** Distance calculator

Starting code:

```

1  ' Calculate distance between two points
2      MOV    x1, #10
3      MOV    y1, #20
4      MOV    x2, #40
5      MOV    y2, #60
6
7      ' Calculate differences
8      SUB    x2, x1        ' dx
9      SUB    y2, y1        ' dy
10
11     ' Your code here: use qvector to find distance

```

Goal: Calculate the distance between the two points Hint: **QVECTOR** dx, dy gives you radius (distance) in QX Success Check: Distance should be 50 units

### The Medicine Cabinet

Feeling overwhelmed by all this trigonometry? Here's your simplified prescription:

**Too Complex?** Just remember these three patterns:

**Pattern 1: Get sine/cosine**

```

1      QROTATE ##$7FFF_FFFF, angle    ' D=radius, S=angle
2      GETQX  cos_value
3      GETQY  sin_value

```

**Pattern 2: Rotate a point**

```

1      SETQ   y
2      QROTATE x, angle
3      GETQX  new_x
4      GETQY  new_y

```

**Pattern 3: Get distance**

```

1      QVECTOR dx, dy
2      GETQX  distance

```

Master these three and you can do 90% of what you need!

## Advanced CORDIC: The Pipeline Dance

Here's where CORDIC gets really powerful - overlapping operations:

```

1  ' Process multiple points while calculating
2  process_points
3      MOV     count, #16
4      MOV     ptra, ##point_array
5
6      ' Start first calculation
7      RDLONG  x, ptra++
8      RDLONG  y, ptra++
9      SETQ    y
10     QROTATE x, angle
11
12 process_loop
13     ' Start next calculation immediately
14     RDLONG  x, ptra++ wz    ' Z flag tells us if done
15     if_nz  RDLONG  y, ptra++
16     if_nz  SETQ    y
17     if_nz  QROTATE x, angle    ' New calculation starts
18
19     ' Get previous result
20     GETQX   prev_x
21     GETQY   prev_y
22
23     ' Store previous result
24     WRLONG  prev_x, ptrb++
25     WRLONG  prev_y, ptrb++
26
27     DJNZ    count, #process_loop
28
29     ' Don't forget last result!
30     GETQX   prev_x
31     GETQY   prev_y
32     WRLONG  prev_x, ptrb++
33     WRLONG  prev_y, ptrb++

```

See what happened? We started each new CORDIC operation immediately after the previous one, then retrieved results later. This pipeline approach means we're effectively getting one rotation every few instructions instead of waiting 55 clocks each time!

## CORDIC for Graphics

Want to draw a spiral? CORDIC makes it trivial:

```

1  ' Expanding spiral generator
2  spiral
3      MOV     angle, #0
4      MOV     radius, #1
5
6  draw_spiral
7      QROTATE radius, angle      ' D=X (radius), S=angle
8      GETQX   x
9      GETQY   y
10
11     ' Convert to screen coordinates
12     SAR     x, #16             ' Scale down
13     SAR     y, #16
14     ADD     x, #320           ' Center X
15     ADD     y, #240           ' Center Y
16
17     ' Plot pixel (simplified)
18     CALL    #plot_pixel
19
20     ' Expand spiral
21     ADD     angle, ##$0400_0000 ' Rotate 5.625 degrees (1/64 turn)
22     ADD     radius, ##100        ' Expand slowly
23
24     CMP     radius, ##30000 wcz
25     if_b   JMP     #draw_spiral

```

## CORDIC for Audio

Generate perfect sine waves for audio:

```

1  ' Audio tone generator using CORDIC
2  tone_generator
3      MOV     phase, #0
4      MOV     frequency, ##$0100_0000 ' ~1.4 deg/sample (1/256 rot)
5
6  sample_loop

```

*continues on next page →*

*↔ continued from previous page*

```

7      QROTATE ##$7FFF_FFFF, phase      ' D=radius, S=angle
8      ADD      phase, frequency        ' Increment phase
9
10     ' Do other audio processing while waiting
11     RDLONG   volume, ##volume_addr
12
13     GETQY    sample                    ' Get sine value
14     SAR      sample, #16                ' Scale to 16-bit
15     MULS     sample, volume            ' Apply volume
16
17     ' Output to DAC
18     WYPIN    sample, #AUDIO_PIN
19
20     ' Wait for sample period (48kHz)
21     WAITX    ##4166                    ' 200MHz / 48kHz
22
23     JMP      #sample_loop

```

## Common CORDIC Gotchas

Before you pull your hair out debugging, know these:

1. **One result at a time** - Each COG has its own CORDIC, but starting a new operation before retrieving your result overwrites it!
2. **55 clocks is exact** - Not 54, not 56. Always exactly 55 clocks from operation start to result ready.
3. **Don't forget SETQ** - For two-operand operations (QROTATE with X,Y), you must load Y into Q first.
4. **Results are scaled** - When rotating by unit circle ( $\$7FFF\_FFFF$ ), results are in 2.30 fixed point format.
5. **Angles wrap naturally** - Adding  $\$1\_0000\_0000$  to an angle is the same as adding 0. Use this!

## What About QLOG, QEXP?

Don't worry, we won't leave logarithms and exponentials behind. CORDIC handles those too:

```

1 ' Base-2 logarithm
2     QLOG    value
3     GETQX   result        ' log2(value) in 5.27 fixed point
4
5 ' Base-2 exponential (2^x)
6     QEXP    value
7     GETQX   result        ' 2^value

```

These are less commonly used but incredibly powerful for DSP and scientific calculations.

### Interlude

#### Jack Volder's Gift to Computing

In 1959, Jack Volder was working on navigation computers for aircraft. He needed to calculate trigonometric functions, but the computers of the day couldn't handle the complex math quickly enough.

His insight? Any angle can be decomposed into a sequence of smaller, fixed angles. By choosing these angles cleverly (arctan of powers of 2), he could rotate vectors using only shifts and **ADDS** - no multiplication needed!

The B-58 bomber's navigation computer was the first to use CORDIC. Today, it's in your P2, calculating sines and cosines faster than those room-sized computers could add two numbers.

From military navigation to your LED projects - quite a journey for an algorithm!

## What We've Learned

Let's celebrate your new CORDIC powers:

- Understood CORDIC's rotate and vector operations
- Generated sine and cosine values
- Calculated distances and angles
- Learned the pipeline technique for speed
- Created rotating graphics
- Built an audio tone generator

That's serious mathematical muscle!

## Coming Up Next

Chapter 8 brings us back to Earth with “Basic I/O” - the fundamental pin operations that make the real world connection. We’ll save Smart Pins for another manual and focus on the essentials: making pins go high and low, reading buttons, and basic timing.

But first, take a moment to appreciate what you just learned. CORDIC is unique to the Propeller 2 - most microcontrollers would need extensive software libraries to do what you just did in three instructions!

**Have Fun!** And remember - with CORDIC, you’re not just calculating trigonometry, you’re doing it at hardware speed. That’s magical!

# Chapter 8: Basic I/O

*Making the real world connection*

## The Hook: One Pin, Three Instructions, Infinite Possibilities

Watch this:

```
1 ' Complete button-and-LED program
2 .loop  TESTP  #BUTTON_PIN wc  ' Read button into C flag
3     if_c  DRVH  #LED_PIN      ' If pressed, LED on
4     if_nc DRVL  #LED_PIN      ' If not pressed, LED off
5         JMP   #.loop          ' Repeat forever
```

Four lines. Complete input/output program. No configuration registers, no data direction setup, no port manipulation. Just pure, simple I/O.

But wait, let me show you the same thing with even more elegance:

```
1 ' Even simpler - button controls LED directly
2 .loop  TESTP  #BUTTON_PIN wc  ' Read button
3         DRVC  #LED_PIN      ' Drive LED from C flag!
4         JMP   #.loop
```

Three lines! The **DRVC** instruction drives the pin to match the C flag. Input becomes output. Simple becomes simpler.

## Understanding P2 Pins

Let's unpack what makes those three lines so short. Every P2 pin is bidirectional and incredibly capable. Unlike older microcontrollers where you set data direction registers (remember TRIS bits? DDRA? Yeah, we don't miss those either), P2 pins change direction on the fly based on the instruction you use.

Here's the mental model:

- **Output instructions** automatically make the pin an output
- **Input instructions** automatically make the pin an input
  
- **Float instructions** make the pin high-impedance
- No setup required!

## Digital Output: Making Things Happen

### The Fundamental Four

```

1      DRVH   #56      ' Drive pin 56 HIGH (3.3V)
2      DRVL   #56      ' Drive pin 56 LOW (0V)
3      DRVNOT #56      ' Toggle pin 56
4      FLTL   #56      ' Float pin 56 (high-Z)

```

That's it. Four instructions, 90% of your output needs covered. We'll meet a few specialized cousins below, but if you only remember these four, you'll get a lot done.

### Conditional Driving

Here's where P2 gets clever:

```

1      DRVC   #56      ' Drive pin to match C flag
2      DRVNC  #56      ' Drive pin to NOT C flag
3      DRVZ   #56      ' Drive pin to match Z flag
4      DRVNZ  #56      ' Drive pin to NOT Z flag

```

And the really clever one:

```

1      DRVNOT #56 wcz   ' Toggle pin AND read old state to C
2      ' C now contains what the pin WAS before toggling

```

### Random and Pattern Outputs

```

1      DRVRND #56      ' Drive pin to a random level (hardware PRNG)
2      OUTL   #56      ' Set OUT bit low (dir unchanged)
3      OUTH   #56      ' Set OUT bit high (dir unchanged)

```

## Digital Input: Reading the World

### Basic Pin Reading

```

1          TESTP   #BUTTON_PIN wc ' Read pin into C flag
2    if_c JMP     #pressed         ' Branch if high
3    if_nc JMP    #not_pressed     ' Branch if low

```

Or read into Z flag for zero/non-zero testing:

```

1          TESTP   #SENSOR_PIN wz ' Read pin into Z flag
2    if_z JMP     #sensor_low      ' Jump if pin low (Z=1 when pin=0)
3    if_nz JMP    #sensor_high     ' Jump if pin is high

```

### Reading Multiple Pins

```

1 ' Read 8 pins at once (pins 0-7)
2     MOV     mask, #$FF         ' Pins 0-7
3     TESTB   ina, #0 wc        ' Test pin 0
4     RCL     result, #1        ' Rotate C into result
5     TESTB   ina, #1 wc        ' Test pin 1
6     RCL     result, #1
7     ' ... continue for all 8 pins

```

## Pin Timing: When Things Happen

### Waiting for Pin Changes

```

1 ' Wait for pin to go high
2 wait_high
3     TESTP   #SIGNAL_PIN wc
4     if_nc JMP    #wait_high
5
6 ' Wait for pin to go low
7 wait_low
8     TESTP   #SIGNAL_PIN wc
9     if_c JMP    #wait_low

```

But there's a better way - hardware-assisted waiting with Smart Events:

```

1      WAITSE1          ' Wait for event 1
2      WAITSE2          ' Wait for event 2
3      ' Configure events to watch pins - super efficient!
```

## Real-World Example: Button Debouncing

Mechanical buttons bounce. Here's how to handle it:

```

1  ' Debounced button reader
2  read_button
3      MOV      debounce, #0
4
5  check_button
6      TESTP   #BUTTON_PIN wc
7      if_c ADD      debounce, #1      ' Count high readings
8      if_nc MOV      debounce, #0      ' Reset on any low
9
10     CMP      debounce, #10 wcz ' Need 10 consecutive highs
11     if_ae JMP      #button_confirmed
12
13     WAITX   ##200_000          ' Wait 1ms at 200MHz
14     JMP      #check_button
15
16  button_confirmed
17     ' Button definitely pressed
18     DRVH    #LED_PIN
```

## Bit-Banged Serial (The Basics)

Yes, you'll usually reach for Smart Pins for serial — but it's worth seeing how to do it the hard way once, just so you appreciate what Smart Pins are doing for you. Here's how:

```

1  ' Bit-bang serial transmit at 115200 baud
2  tx_byte
3      OR      data, ##$100      ' Add stop bit
4      SHL     data, #1          ' Add start bit (0)
5      MOV     bits, #10         ' 1 start + 8 data + 1 stop
```

*continues on next page →*

```

6
7     GETCT    time           ' Get current time
8
9 tx_loop
10    SHR     data, #1 wc     ' Get next bit into C
11    DRVC    #TX_PIN        ' Output bit
12
13    ADDCT1  time, bit_time  ' Next bit time
14    WAITCT1                ' Wait for it
15
16    DJNZ    bits, #tx_loop
17    RET
18
19 bit_time LONG 100_000_000 / 115200 ' Clock cycles per bit

```

## Your Turn: I/O Experiments

### Your Turn

**Your Turn:** Create a light chaser

Starting code:

```

1     ORG     0
2
3     MOV     pattern, #1     ' Start with one LED
4
5 .loop MOV     pins, pattern  ' Your code here
6     ' Make pattern rotate through pins 56-63
7     ' Add delay between changes
8     ' Wrap around at the end, then jmp #.loop

```

Goal: Create a rotating light pattern on LEDs Hint: Use **SHL** and check for overflow

Success Check: Single lit LED rotating through all positions

### Your Turn

**Your Turn:** Reaction timer

Starting code:

```

1      ORG      0
2
3      ' Turn on LED after random delay
4      GETRND  delay
5      AND     delay, ##$3FFF_FFFF ' Limit range
6      WAITX   delay
7      DRVH    #LED_PIN
8
9      GETCT   start_time
10     ' Your code: wait for button press
11     ' Calculate reaction time

```

Goal: Measure reaction time between LED and button press Hint: Use **GETCT** after button detection Success Check: Time measured in clock cycles

### The Medicine Cabinet

Feeling overwhelmed by all these pin operations? Here's the simplified prescription:

**Just need something working?** Remember these patterns:

#### Output pattern:

```

1      DRVH    #PIN    ' Make it high
2      DRVL    #PIN    ' Make it low
3      DRVNOT  #PIN    ' Toggle it

```

#### Input pattern:

```

1      TESTP   #PIN wc ' Read it
2      if_c    JMP     #high ' It's high
3      if_nc   JMP     #low  ' It's low

```

#### Timed pattern:

```

1  .loop  DRVNOT  #LED
2        WAITX   ##50_000_000
3        JMP     #.loop

```

That's 80% of all I/O right there!

## Advanced Pin Control

### Pin Groups

You can control multiple pins at once:

```

1          DRVH    #LED_BASE addpins 3  ' Drive 4 pins high (base+3)
2          DRVL    #LED_BASE addpins 7  ' Drive 8 pins low

```

### Direct Pin Manipulation

For when you need absolute control:

```

1          MOV     outa, pattern    ' Set output register directly
2          MOV     dira, ##$FF     ' Set direction reg (rare in P2!)

```

But honestly? You'll rarely need these. The individual pin instructions are cleaner and clearer.

## Common I/O Gotchas

Before you pull your hair out wondering why a pin “won’t work,” save yourself debugging time and skim these:

1. **Pin numbers are 0-63** - Not port.bit notation like other MCUs
2. **No pullup/pulldown by default** - Use external resistors or configure Smart Pin modes (advanced topic)
3. **Pins float on reset** - All pins start as inputs (floating)
4. **Reading output pins** - You CAN read a pin you're driving (reads the actual pin state)
5. **3.3V logic levels** - P2 is 3.3V, not 5V tolerant!

## Timing Is Everything

Now here's something we'll keep coming back to: P2 I/O is deterministic. When you execute:

```

1          DRVH    #56
2          DRVL    #57

```

Pin 56 goes high and pin 57 goes low at EXACTLY the same clock cycle. No skew, no uncertainty. Uff! Try doing that on an interrupt-driven MCU. This determinism is what makes P2 perfect for precise timing applications.

## Real-World Example: Servo Control

Even without Smart Pins, controlling a servo is easy:

```
1 ' Standard servo control (1-2ms pulse every 20ms)
2 servo_control
3     MOV     position, ##300_000    ' 1.5ms = center at 200MHz
4
5 servo_loop
6     DRVH   #SERVO_PIN
7     WAITX  position                ' 1-2ms high pulse
8     DRVL   #SERVO_PIN
9     WAITX  ##4_000_000            ' Rest of 20ms at 200MHz
10
11     ' Adjust position as needed
12     RDLONG position, ##position_addr
13     FLE    position, ##200_000    ' Limit to 1ms min
14     FGE    position, ##400_000    ' Limit to 2ms max
15
16     JMP    #servo_loop
```

## The Power of Simple

Here's something beautiful about P2's I/O philosophy: it's transparent. Unlike microcontrollers with complex GPIO configurations, port multiplexing, and alternate functions, P2 pins just... work.

Want an output? Drive it. Want an input? Read it. Want it to float? Float it.

No setup, no configuration, no confusion.

## What We've Learned

Look at your new I/O skills:

- Understood P2's automatic pin direction
- Mastered the four fundamental output instructions
- Learned pin reading and conditional testing
- Created debounced inputs
- Built bit-banged serial
- Discovered deterministic timing

## A Note About Smart Pins

You might wonder - if basic I/O is this simple, why do we need Smart Pins?

Well, while you CAN bit-bang serial at 115200 baud, or generate PWM, or measure frequencies using the techniques in this chapter, Smart Pins do all of this in hardware, freeing your COG for more important work.

**For Smart Pin details:** See the dedicated “P2 Smart Pins Manual” which covers all 64 modes, from simple PWM to complex protocol generation. Smart Pins deserve their own complete treatment!

## Coming Up Next

Chapter 9 takes us into “Streaming Data” - the P2’s incredible FIFO system that can move megabytes of data without breaking a sweat. We’ll see how to stream video, audio, and massive data blocks at maximum speed.

**Have Fun!** Remember, every embedded system ultimately comes down to pins going high and low. You’ve just mastered the fundamentals that everything else builds upon!

# Chapter 9: Streaming Data

*Moving mountains of data without breaking a sweat*

## The Hook: 4KB in 4 Instructions

Watch this data transfer magic:

```

1 ' Copy 1000 longs (4KB) at maximum speed
2     SETQ    ##1000-1      ' Setup for 1000 longs
3     RDLONG  buffer, source ' Read them all!
4     SETQ    ##1000-1      ' Setup for 1000 longs
5     WRLONG  buffer, dest   ' Write them all!
6     ' 4KB moved in microseconds!
```

Four instructions. Four kilobytes. Faster than DMA on most processors. And we're just getting started...

## Block Transfers: The Power Move

The **SETQ** instruction is your gateway to block transfers:

```

1 ' Basic block read
2     SETQ    #16-1          ' Transfer 16 longs (minus 1!)
3     RDLONG  buffer, hubaddr ' Reads 16 consecutive longs
4
5 ' Basic block write
6     SETQ    #16-1          ' Transfer 16 longs
7     WRLONG  buffer, hubaddr ' Writes 16 consecutive longs
```

Here's the trick: **SETQ** tells the next hub instruction how many longs to transfer. The “-1” is because it's a count from 0 (yes, we'll trip over that off-by-one at least once — everyone does).

## The FIFO: Your Streaming Pipeline

The FIFO (First In, First Out) is P2's streaming engine. Think of it as a conveyor belt between hub memory and your COG:

```

1 ' Start FIFO reading
2     RDFAST #0, ##data_start ' Start fast read at data_start
3
4 ' Now read data at maximum speed
5 stream_loop
6     RFLONG value           ' Read from FIFO (no waiting!)
7     ' Process value here
8     ADD     accumulator, value
9     DJNZ   count, #stream_loop
10
11 ' The FIFO keeps feeding data automatically

```

The beauty? The FIFO reads ahead automatically. While you're processing one value, it's already fetching the next. No hub timing slots to worry about!

## Writing Through the FIFO

Reading was nice — writing is symmetric. We'll use **WRFAST** instead of **RDFAST**, and **WFLONG** instead of **RFLONG**, and that's it:

```

1 ' Start FIFO writing
2     WRFAST #0, ##dest_buffer
3
4 ' Stream data out
5 write_loop
6     ' Generate or process data
7     MOV     value, calculation
8     WFLONG value           ' Write to FIFO
9     DJNZ   count, #write_loop
10
11 ' Data streams to hub automatically

```

## Real-World Example: Screen Buffer Clear

Let's clear a 320x240x4 **BYTE** screen buffer (~307KB - fits in hub!):

```

1 ' Ultra-fast screen clear
2 clear_screen
3     MOV     color, ##$00_00_00_00 ' Black (4 bytes per pixel)

```

*continues on next page →*

```

4      MOV    pixels, ##320*240      ' 76,800 pixels
5
6      WRFast #0, ##screen_buffer    ' Start FIFO write
7
8  clear_loop
9      WFLong color                  ' Write 4-byte pixel
10     DJNZ   pixels, #clear_loop
11
12     ' 307KB cleared at maximum hub speed!
13     ' Note: Hub RAM is 512KB - plan buffer sizes accordingly

```

## Streaming with the Streamer

The Streamer is different from the FIFO - it's a dedicated DMA engine that can move data between hub memory and pins:

```

1  ' Configure streamer for video output
2      SETXFRQ ##PIXEL_FREQ      ' Set the pixel (NCO) output rate
3
4  ' Start streaming video data to pins
5      XINIT  ##STREAM_CMD, #0    ' Start streamer
6      ' Data flows from hub to pins automatically!

```

## FIFO and COG Execution

Here's something amazing - you can execute code from hub through the FIFO:

```

1  ' Execute large program from hub
2      ORGH   $1000                ' Code in hub memory
3
4  hub_code
5      ' This code is in hub but executes like it's in COG
6      ADD    x, y
7      SUB    a, b
8      ' Can be megabytes of code!

```

When you **CALL** or jump to hub code, the FIFO automatically feeds instructions to the COG. It's like having unlimited code space!

### The Medicine Cabinet

Feeling overwhelmed by all this streaming? Here's your prescription:

**Just need to move data?** Use these simple patterns:

**Block read pattern:**

```
1      SETQ    #SIZE-1
2      RDLONG  buffer, source
```

**Block write pattern:**

```
1      SETQ    #SIZE-1
2      WRLONG  buffer, dest
```

**FIFO read pattern:**

```
1      RDFAST  #0, ##source
2  .loop  RFLONG  value
3      ' Process value
4      DJNZ    count, #.loop
```

That's 90% of streaming right there!

## Advanced Streaming Techniques

### Circular Buffers with FIFO

```
1  ' Circular buffer reading
2      RDFAST  ##BUF_BLOCKS, ##buffer ' D[13:0]=block count; auto-wraps
3
4  circular_loop
5      RFLONG  value                ' Read from FIFO
6      ' Process value
7      ' FIFO automatically wraps at buffer end!
8      JMP     #circular_loop
```

## Processing Pipeline with FIFO

```

1  ' Read data with FIFO, process, write via PTRB
2      RDFAST #0, ##source      ' Set up FIFO for reading
3      MOV     dest_ptr, ##dest
4
5  pipeline
6      RFLONG  input            ' Get next input from FIFO
7
8      ' Scale using 16x16 multiply (result in input)
9      MUL     input, #SCALE_FACTOR
10
11     WRLONG  input, dest_ptr ' Write result via PTRB
12     ADD     dest_ptr, #4
13     DJNZ   count, #pipeline

```

Note: FIFO can only read OR write at a time, not both. Use PTRB/PTRB for the other direction.

## Your Turn: Streaming Experiments

### Your Turn

**Your Turn:** Fast memory fill

Starting code:

```

1      ORG     0
2
3      MOV     pattern, ##$DEADBEEF
4      MOV     dest, ##$1000
5      MOV     count, #256
6
7      ' Your code here: Fill 256 longs with pattern
8      ' Use SETQ and WRLONG

```

Goal: Fill memory with pattern using block transfer Hint: You'll need **SETQ #255** (not #256) Success Check: Memory filled in one operation

## Your Turn

### Your Turn: Data filter pipeline

Starting code:

```

1      ORG      0
2
3      RDFAST   #0, ##input_data      ' FIFO for reading
4      MOV     ptr, ##output_data     ' PTR for writing
5      MOV     count, #100
6
7  filter_loop
8      RFLONG  value                  ' Read from FIFO
9      ' Your code: Simple filter
10     ' Maybe average with previous value?
11     WRLONG  result, ptr++          ' Write via PTR
12     DJNZ   count, #filter_loop

```

Goal: Process streaming data through simple filter Hint: Keep previous value in a register

Success Check: Output is filtered version of input

## Common Streaming Gotchas

Before you pull your hair out wondering why your transfer is one **LONG** short, or why your FIFO won't cooperate, skim these:

1. **SETQ uses count-1** - For 16 longs, use SETQ #15, not SETQ #16
2. **FIFO is shared per COG** - Can't use FIFO for both code execution and data streaming simultaneously
3. **Write synchronization - WRFAST** writes complete in the background. To force a flush, issue the next **RDFAST**/**WRFAST** with D[31]=0 (it waits for the prior **WRFAST** to finish) rather than relying on a fixed delay
4. **Hub alignment** - Block transfers work best with long-aligned addresses
5. **FIFO depth** - The FIFO holds (cogs+11) = 19 longs. It refills automatically, so you rarely outrun it.

## Performance Numbers

Let's talk speed:

- **Block transfer:** Up to 1 **LONG** per clock (at 200MHz = 800MB/s!)
- **FIFO streaming:** Up to 1 **LONG** per clock sustained
- **Random hub access:** 9-16 clocks per access
- **Streamer to pins:** Up to sysclock/1 rate

Uff! That's seriously fast. Most microcontrollers need dedicated DMA controllers, peripheral coprocessors, and a stack of config registers to achieve what P2 does with two instructions. You're not paying for that DMA controller — it's already in the silicon.

## Real-World Example: Audio Buffer

Stream audio samples through processing:

```

1  ' Audio processing pipeline
2  audio_process
3      RDFAST  #0, ##input_buffer      ' FIFO for reading input
4      MOV    ptra, ##output_buffer    ' PTRA for writing output
5      MOV    samples, ##BUFFER_SIZE
6
7  process_loop
8      RFLONG  left_sample              ' Read left from FIFO
9      RFLONG  right_sample             ' Read right from FIFO
10
11     ' Apply simple low-pass filter
12     ADD    left_filtered, left_sample
13     SHR    left_filtered, #1         ' Average with previous
14
15     ADD    right_filtered, right_sample
16     SHR    right_filtered, #1
17
18     ' Apply volume
19     MULS   left_filtered, volume
20     MULS   right_filtered, volume
21
22     ' Output processed samples via PTRA
23     WRLONG  left_filtered, ptra++
24     WRLONG  right_filtered, ptra++
25
26     DJNZ   samples, #process_loop

```

## What We've Learned

Your streaming skills now include:

- Block transfers with **SETQ**
- FIFO reading and writing
- Streaming pipeline concepts
- Circular buffer techniques
- Parallel processing while streaming
- Real-world applications

## Coming Up Next

Chapter 10 explores “Hub Execution” - how to break free from the 496-instruction limit and run massive programs directly from hub memory. It's like having your cake and eating it too!

**Have Fun!** Remember, streaming is about throughput, not just speed. It's the difference between carrying one brick at a time and using a wheelbarrow!

# Chapter 10: Hub Execution

*Breaking free from the 496-instruction limit*

## The Hook: Unlimited Code Space

Remember fretting about fitting your code into 496 COG instructions? Watch this:

```
1      ORGH    $400          ' Place code in hub memory
2
3      ' This can be thousands of instructions!
4 main  MOV    x, #0
5      MOV    y, #0
6      CALL   #huge_function ' Can be massive
7      CALL   #another_big_one
8      CALL   #yet_another
9      ' Keep going... no limit!
10
11 huge_function
12      ' 1000 instructions? No problem!
13      ' 10000 instructions? Still fine!
14      RET
```

Your code now lives in hub memory's 512KB instead of COG memory's 2KB. That's 256 times more space!

## COG vs Hub Execution: The Trade-offs

Let's be honest about the differences:

**COG Execution** (traditional):

- Fast: exactly 2 clocks per instruction
- Deterministic: perfect for real-time
- Limited: only 496 instructions
- Self-contained: runs independently

**Hub Execution** (the new way):

- Fast sequential: 2 clocks per instruction (same as cog-exec, thanks to the 19-stage FIFO prefetch)

- Slower on branches: minimum 13 clocks per branch (the FIFO refill cost; +1 if target isn't long-aligned)
- Unlimited: 512KB of code space!
- Flexible: can call COG routines

The beauty? You can mix both in the same program! Sequential code in hub runs at full speed — only branches show the hub-execution penalty.

## How Hub Execution Works

Let's peek behind the curtain. When the processor encounters a jump or call to a hub address ( \$400), it automatically switches to hub execution mode. The FIFO starts streaming instructions from hub memory:

```

1      ORG      0          ' Start in COG
2
3  cog_code
4      ' This runs from COG RAM
5      CALL    #hub_function ' Call into hub
6      ' Back in COG mode here
7
8      ORGH    $1000      ' Switch to hub addresses
9
10 hub_function
11     ' This runs from hub RAM via FIFO
12     ' Can be huge!
13     RET              ' Returns to COG code

```

The magic happens automatically. No mode switching instructions needed!

## Real-World Example: Menu System

Here's something that would never fit in COG RAM:

```

1      ORGH    $2000
2
3  menu_system
4      CALL    #draw_menu_frame
5      CALL    #display_options
6      CALL    #get_user_input

```

*continues on next page →*

```
7         CALL    #process_selection
8
9         CMP     selection, #1 wcz
10      if_e CALL  #option_1_handler
11         CMP     selection, #2 wcz
12      if_e CALL  #option_2_handler
13         CMP     selection, #3 wcz
14      if_e CALL  #option_3_handler
15         ' ... many more options
16
17         JMP     #menu_system
18
19 draw_menu_frame
20         ' 200 instructions for fancy graphics
21         RET
22
23 display_options
24         ' 300 instructions for text rendering
25         RET
26
27 option_1_handler
28         ' 500 instructions for configuration
29         RET
30
31 ' Thousands of instructions total - no problem!
```

## The Hub Execution FIFO

You met the FIFO in the previous chapter as a data-streaming engine. Same hardware, different job here: it reads ahead, keeping a buffer of upcoming *instructions* ready for the COG. Think of it as a moving sidewalk for your code:

```

1  ' The FIFO maintains performance by reading ahead
2  hub_loop
3      ADD    x, y          ' FIFO has next instructions ready
4      SUB    a, b          ' No waiting for hub access
5      MUL    c, d          ' Instructions stream smoothly
6      ' FIFO automatically refills as needed

```

This read-ahead behavior is the whole reason sequential hub code matches cog-exec speed — the FIFO is doing the waiting for you, in parallel with the COG running instructions it already prefetched.

## Mixing COG and Hub Code

Here's the real power - combining both modes:

```

1      ORG    0
2
3  ' Critical timing code in COG
4  critical_loop
5      TESTP  #TRIGGER_PIN wz      ' Test trigger pin
6  if_nz JMP  #critical_loop        ' Loop until triggered
7      DRVH  #CRITICAL_PIN        ' Immediate response!
8      CALL  #hub_process          ' Do complex processing
9      JMP   #critical_loop
10
11     ORGH  $4000
12
13  ' Complex processing in hub
14  hub_process
15      ' Hundreds of instructions for data analysis
16      ' Not time-critical, so hub execution is fine
17      RET

```

Time-critical code stays in COG RAM for deterministic timing. Complex code lives in hub RAM for space.

## Your Turn: Hub Execution Experiments

### Your Turn

**Your Turn:** Build a simple calculator

Starting code:

```
1      ORG      0
2      JMP      #calculator      ' Jump to hub code
3
4      ORGH     $1000
5 calculator
6      ' Your code here:
7      ' 1. Display menu
8      ' 2. Get operation choice
9      ' 3. Get two numbers
10     ' 4. Call appropriate function
11     ' 5. Display result
12
13 add_function
14     ' Addition code
15     RET
16
17 subtract_function
18     ' Subtraction code
19     RET
20
21 ' Add more functions - no size limit!
```

Goal: Create a multi-function calculator Hint: Each function can be as complex as needed

Success Check: Multiple operations working

### The Medicine Cabinet

Overwhelmed by execution modes? Here's the simple version:

#### Keep it simple:

1. **Small, time-critical code** → Put in COG (org 0)
2. **Large, complex code** → Put in hub (orgh \$400+)
3. **Don't overthink it** → The processor handles the switch

#### Basic pattern:

```

1          ORG      0
2          JMP      #main      ' Jump to hub
3
4          ORGH     $400
5 main     ' Your big program here

```

That's it. Let the processor worry about the details!

## Advanced Hub Execution

### Long Jumps and Calls

Hub addresses need 20 bits, so jumping far requires special handling:

```

1 ' Jump to distant hub code
2     JMP      ##far_away      ' ## forces 32-bit immediate
3
4     ORGH     $40000          ' Far away in hub
5 far_away
6     ' Code here

```

### Hub Data Access from Hub Code

When executing from hub, you can still access hub data:

```

1     ORGH     $1000
2
3 hub_code
4     RDLONG  value, ##hub_data ' Read hub data

```

*continues on next page →*

```

5      ADD     value, #1
6      WRLONG  value, ##hub_data ' Write back
7
8      ORGH   $8000
9  hub_data
10     LONG   $12345678

```

## Performance Optimization

To maximize hub execution speed:

```

1  ' Align branch targets to 8-byte boundaries
2      ALIGNL                ' Align to long boundary
3  loop_start
4      ' Loop code here
5      DJNZ   count, #loop_start
6
7  ' Keep critical loops small
8  ' Consider moving inner loops to COG RAM

```

## Common Hub Execution Gotchas

Before you cram everything into hub and call it a day, know these:

1. **Speed variation** - Don't use hub execution for precise timing
2. **FIFO conflicts** - Can't stream data while executing from hub
3. **Address confusion** - Remember: <\$200 is COG, \$200-\$3FF is LUT, \$400 is hub
4. **Stack depth** - Still limited to 8-level hardware stack
5. **Relative jumps** - Work differently in hub mode

## Real-World Example: Command Parser

```

1      ORGH   $2000
2
3  command_parser
4      CALL   #get_command_line
5      CALL   #tokenize
6

```

```
7      ' Compare against commands
8      MOV    ptra, #command_string
9      MOV    ptrb, ##cmd_help
10     CALL   #string_compare
11     if_z  JMP    #help_command
12
13     MOV    ptrb, ##cmd_run
14     CALL   #string_compare
15     if_z  JMP    #run_command
16
17     ' Many more commands...
18
19     help_command
20     ' 500 instructions of help text display
21     RET
22
23     run_command
24     ' 1000 instructions of program execution
25     RET
26
27     string_compare
28     ' 50 instructions of string comparison
29     RET
30
31     ' Thousands of instructions total
32     ' Would need multiple COGs without hub execution!
```

## When to Use Hub Execution

### Perfect for:

- User interfaces and menus
- Command processors
- Complex algorithms
- String manipulation
- Protocol handlers
- Error handling and recovery

### Avoid for:

- Interrupt handlers (if you use them)

- Precise timing loops
- Bit-banged protocols
- Real-time control loops

## What We've Learned

You've mastered hub execution:

- Understanding COG vs hub trade-offs
- Automatic mode switching
- Mixing COG and hub code
- FIFO streaming of instructions
- When to use each mode
- Real-world applications

## Coming Up Next

Chapter 11 tackles the controversial topic: “Why No Interrupts?” We'll explore why the Propeller philosophy says you don't need them, and why that's actually a good thing!

**Have Fun!** Hub execution is like having a sports car that can also carry furniture - you get both speed and capacity when you need them!

# Chapter 11: Why No Interrupts?

*The most controversial P2 feature explained*

## The Hook: Interrupts Without Interrupts

Here's a traditional interrupt-driven button handler:

```
1 ' Traditional approach (not P2!) - pseudocode for the interrupt-driven style
2 ISR(BUTTON_INTERRUPT)
3     ' Interrupt service routine
4     buttonPressed = true
5     ' Return to interrupted code
6 END_ISR
```

And here's the P2 way:

```
1 ' Dedicated COG watching button
2 button_watcher
3     TESTP    #BUTTON_PIN wc
4     if_c WRLONG  ##1, ##button_flag
5     JMP      #button_watcher
6
7 ' Main COG doing important work
8 main_code
9     ' Never interrupted!
10    ' Checks button_flag when convenient
```

No interrupt latency. No context switching. No priority inversion. No critical sections. Just clean, deterministic, parallel processing.

## The Interrupt Problem

Let me tell you a story. You're concentrating on a complex problem when someone taps your shoulder. You stop, handle their request, then try to remember where you were. Now imagine this happening randomly, unpredictably, dozens of times per second.

That's interrupts.

Traditional processors need interrupts because they only have one processor. Something important happens? Stop everything and handle it! But this causes:

- **Latency:** Time to save context and jump to handler
- **Jitter:** Variable response time depending on what was interrupted
- **Priority inversion:** Low-priority task blocks high-priority
- **Race conditions:** Shared data access problems
- **Debugging nightmares:** Non-reproducible timing bugs

## The Propeller Solution

Eight processors. No sharing required.

```
1 ' COG 0: Main application
2 main_app
3     ' Complex calculations
4     ' Never interrupted
5     RDLONG  command, ##mailbox wz
6     if_nz CALL  #process_command
7     JMP     #main_app
8
9 ' COG 1: Serial port handler
10 serial_handler
11     ' Continuously monitors serial
12     TESTP   #RX_PIN wc
13     if_c CALL  #receive_byte
14     JMP     #serial_handler
15
16 ' COG 2: Motor control
17 motor_control
18     ' Precise timing loops
19     ' Never disrupted
20     WAITCNT motor_time
21     DRVNOT  #STEP_PIN
22     JMP     #motor_control
23
24 ' COG 3: Sensor monitor
25 sensor_monitor
26     ' Watches multiple sensors
27     ' Responds instantly
28     ' ... and so on
```

Each COG does one thing perfectly. No interruptions. No conflicts. Just pure, focused execution.

## Real-World Example: Perfect Servo Control

With interrupts, servo pulses jitter. With dedicated COGs, they're perfect:

```

1  ' COG dedicated to servo control
2  servo_cog
3      GETCT    pulse_time
4
5  servo_loop
6      ' Generate 8 servo pulses simultaneously
7      MOV      servo_mask, ##$FF      ' 8 servos
8      OR       outa, servo_mask      ' All high
9
10     MOV      index, #0
11  check_servos
12     RDLONG   width, ptra[index]     ' Get pulse width
13     ADDCT1   pulse_time, width     ' Set compare time
14
15     WAITCT1
16     BITL    outa, index             ' Turn off this servo
17
18     INCMOD   index, #7  wc          ' C set when index wraps 7->0
19  if_nc JMP    #check_servos        ' Loop until all 8 servos done
20
21     ' Wait for 20ms frame
22     WAITX    ##4_000_000
23     JMP      #servo_loop
24
25  ' Result: 8 servos with ZERO jitter!

```

Try that with interrupts. I'll wait. Actually, I won't - it's impossible to achieve this precision with interrupts.

### “But P2 HAS Interrupts!”

Yes, it does. And you probably shouldn't use them.

Well, let me be more nuanced. P2 has interrupts for those rare cases where you absolutely need them:

```

1 ' Setting up an interrupt (not recommended!)
2     SETSE1  %#001<<6 + PANIC_BUTTON  ' SE1 triggers when pin goes high
3     SETINT1 #EVENT_SE1                ' Enable INT1 on SE1 event
4
5 int1_handler
6     ' Interrupt code here
7     RETI1

```

When might you use them?

- Porting legacy code that requires interrupts
- Ultra-low-power designs where COGs must sleep
- Theoretical minimum latency response (but dedicated COG is usually faster!)

Uff! Even writing interrupt code feels wrong on a Propeller!

### The Medicine Cabinet

Still thinking you need interrupts? Here's your medicine:

**Think you need an interrupt for...**

**Fast response?**

```

1 ' Dedicated COG responds in ~4 clocks
2 watcher
3     TESTP  #INPUT_PIN wz            ' Test pin state
4     if_nz JMP  #watcher              ' Loop until pin high
5     DRVH   #RESPONSE_PIN           ' Instant response!

```

**Multiple events?**

```

1 ' One COG watches everything
2 monitor
3     TEST   sensors, #SENSOR1 wz
4     if_nz CALL #handle_sensor1
5     TEST   sensors, #SENSOR2 wz
6     if_nz CALL #handle_sensor2
7     ' Check all sensors every loop

```

*continues on next page →*

**Periodic tasks?**

```

1 ' Perfect timing without interrupts
2     GETCT    next_time
3 .loop  ADDCT1  next_time, ##PERIOD
4     WAITCT1                ' Exact timing
5     CALL    #periodic_task
6     JMP     #.loop

```

See? No interrupts needed!

**The Event System: Better Than Interrupts**

P2 has something better than interrupts - events:

```

1 ' Configure event to watch pin
2     SETSE1  %#01_000000 | BUTTON_PIN ' Rising edge event
3
4 ' Main code runs normally
5 main_loop
6     ' Do work...
7     POLLSE1 wc                ' Check if event occurred
8     if_c CALL    #handle_button ' Handle when convenient
9     ' Continue work...
10    JMP     #main_loop

```

Events are like interrupts that wait politely for you to check them. No rudeness!

**Interrupt Horror Stories**

Let me share why we avoid interrupts:

**Story 1: The Jittery Display**

Approach	Problem	Result
<b>With Interrupts</b>	Display updates interrupted by serial	Visible glitches, tearing, inconsistent timing
<b>With COGs</b>	Display COG runs uninterrupted	Perfect, smooth, glitch-free display

## Story 2: The Missed Pulse

Approach	Problem	Result
<b>With Interrupts</b>	Motor step interrupted by sensor read	Missed step, motor stalls, position lost
<b>With COGs</b>	Motor COG never misses a beat	Perfect positioning, no lost steps

## Story 3: The Debugging Nightmare

Approach	Problem	Result
<b>With Interrupts</b>	Bug only appears under specific timing	Days of debugging, hair loss, coffee overdose
<b>With COGs</b>	Deterministic timing, reproducible behavior	Bug found in minutes, sanity preserved

## Your Turn: COG vs Interrupt Challenge

### Your Turn

**Your Turn:** Build a reaction timer without interrupts

Starting code:

*continues on next page →*

```

1      ORG      0
2
3  ' COG 0: Main game logic
4      SETQ    ##button_flag          ' PTRA for new COG
5      COGINIT #1, @button_watcher
6
7  game_loop
8      ' Random delay
9      GETRND  delay
10     WAITX   delay
11
12     DRVH    #LED_PIN
13     WRLONG  #0, ##button_flag
14     GETCT   start_time
15
16  ' Wait for button (no interrupt!)
17  wait_press
18     RDLONG  pressed, ##button_flag wz
19     if_z   JMP    #wait_press
20
21     GETCT   end_time
22     SUB     end_time, start_time
23     ' Display reaction time
24
25  ' COG 1: Button watcher
26     ORGH    $400
27  button_watcher
28     ' Your code here

```

Goal: Implement button watcher COG Hint: Continuously monitor and set flag Success  
Check: Perfect timing without interrupts

## The Philosophy Deep Dive

The Propeller philosophy is about **determinism over responsiveness**.

Traditional processors optimize for average-case performance:

- Interrupts handle rare events
- Most code runs uninterrupted

- When events happen, everything stops

Propeller optimizes for worst-case determinism:

- Every COG runs predictably
- No surprises, ever
- Timing is guaranteed

It's the difference between a talented soloist who might miss a note and an orchestra where everyone plays their part perfectly.

## When Interrupts Actually Make Sense

I'll admit it - there are rare cases where interrupts are appropriate:

1. **Power-critical applications** where COGs must sleep
2. **Legacy code ports** that fundamentally require interrupts
3. **Single-COG designs** (but why waste the P2's power?)

But in 15 years of Propeller programming, I've needed interrupts exactly... never.

## Common “But What About...” Questions

**Q: “But what about interrupt priority?”** A: COGs don't have priority. They're all equal. Design your system accordingly.

**Q: “How do I handle critical events?”** A: Dedicate a COG to critical events. It will respond faster than any interrupt.

**Q: “Isn't dedicating a whole COG wasteful?”** A: You have eight! And a focused COG is simpler than interrupt-riddled code.

**Q: “What about power consumption?”** A: Use WAITSE/WAITCT for low-power waiting. COG sleeps until event.

## What We've Learned

You now understand the Propeller way:

- Why interrupts cause problems
- How COGs eliminate interrupt need
- Event system as polite alternative
- Real-world benefits of no interrupts
- Rare cases where interrupts might be used

- The philosophy of determinism

## Coming Up Next

Chapter 12 shows you “Optimization Mastery” - how to make your PASM2 code blazingly fast. We’ll explore the pipeline, instruction pairing, and timing tricks that squeeze every drop of performance from the P2.

**Have Fun!** And remember - in a world of interruptions, be a COG: focused, deterministic, and uninterruptible!

# Chapter 12: Optimization Mastery

*Making the fast faster*

## The Hook: Double Your Speed with One Change

Let me show you a loop that looks fine — until you realize you're paying for the same thing twice. Watch:

```

1  ' Before optimization: 13 clocks
2  .loop  RDLONG  value, ptra      ' 9-16 (cog-exec) / 9-26 (hub-exec)
3          ADD    value, #1       ' 2 clocks
4          WRLONG  value, ptra    ' 3-10 (cog-exec) / 3-20 (hub-exec)
5          ADD    ptra, #4        ' 2 clocks
6          DJNZ   count, #.loop   ' 2/4 (cog-exec) / 2/13-20 (hub-exec)
7
8  ' After optimization using PTR expressions:
9  .loop  RDLONG  value, ptra      ' Read from current address
10         ADD    value, #1       ' Process
11         WRLONG  value, ptra++   ' Write and increment in one!
12         DJNZ   count, #.loop   ' Saved the ADD instruction

```

Almost twice as fast! The secret? Understanding how P2 really works.

## Understanding the Pipeline

Before we hunt for clocks to save, let's understand where they go in the first place. P2 has a 2-stage pipeline:

1. **Fetch** - Get next instruction
2. **Execute** - Do the work

This means while one instruction executes, the next is already being fetched:

```

1          ADD    x, y           ' Executing while next inst fetches
2          SUB    a, b           ' Fetching while previous executes
3          ' Perfect overlap = maximum throughput

```

## Instruction Timing Basics

Not all instructions are created equal:

```

1 ' 2-clock instructions (most ALU operations)
2     ADD    x, y          ' 2 clocks
3     MOV    a, b          ' 2 clocks
4     AND    c, d          ' 2 clocks
5
6 ' Variable timing (hub access)
7     RDLONG value, hubaddr ' 9-16 clocks (hub slot wait)
8     WRLONG value, hubaddr ' 3-10 clocks (variable)
9
10 ' Long operations (CORDIC)
11     QROTATE x, angle    ' 2 clocks to start
12     GETQX  result      ' 2 clocks (but wait 55 for result)
13
14 ' Special cases
15     MUL    x, y          ' 2 clocks
16     QDIV   x, y          ' 2 clocks to start
17     GETQX  result      ' 2 clocks (but wait 30 for result)

```

## REP: The Speed Loop

**REP** creates hardware-accelerated loops with zero overhead:

```

1 ' Traditional loop: overhead per iteration
2 .loop  ADD    sum, value  ' 2 clocks
3        ADD    ptr, #4     ' 2 clocks
4        DJNZ   count, #.loop ' 2 or 4 clocks (4 if branch taken)
5
6 ' REP loop: 0 clocks overhead!
7        REP    #2, count   ' Repeat next 2 instructions
8        ADD    sum, value  ' 2 clocks
9        ADD    ptr, #4     ' 2 clocks = 4 total!

```

That's 33% faster just by using **REP**!

**Hub-Exec Note:** **REP** works in hub-exec too, but each iteration executes a hidden jump to loop back — and that hidden jump pays the 13+ clock hub-branch cost per iteration. So a 2-instruction **REP** loop that takes 4 clocks in cog-exec balloons to ~17+ clocks per iteration in hub-exec. For

time-critical inner loops, keep **REP** in COG or LUT memory. Hub-exec **REP** works correctly; it just isn't zero-overhead there.

## SKIP: Conditional Execution on Steroids

**SKIP** and **SKIPF** let you conditionally execute patterns of instructions:

```

1 ' Traditional: multiple jumps
2     CMP    x, #5 wcz
3 if_a    JMP    #greater
4 if_b    JMP    #less
5         JMP    #equal
6
7 ' With SKIP: cancel optional steps per a precomputed pattern
8 '   bit N (LSB first) skips the Nth instruction after SKIP
9     SKIP   config_mask    ' e.g. %010 runs steps 0 and 2, skips step 1
10    CALL  #setup_uart    ' Step 0
11    CALL  #setup_spi     ' Step 1
12    CALL  #setup_timer   ' Step 2
13    ' One pattern picks which steps run - no jumps, no stalls!

```

## Hub Access Optimization

Hub timing is critical for performance:

```

1 ' Unaligned hub access: variable timing
2     RDLONG v1, ##$1001    ' Not long-aligned, slower
3
4 ' Aligned hub access: predictable timing
5     RDLONG v1, ##$1000    ' Long-aligned, faster
6
7 ' Sequential access: maximum speed
8     RDLONG v1, ptr++      ' Hardware manages sequence
9     RDLONG v2, ptr++      ' Optimal hub slot usage
10    RDLONG v3, ptr++      ' Maximum throughput

```

## The FIFO Fast Path

For ultimate speed, use the FIFO:

```

1 ' Traditional hub reading: ~6 clocks average per long
2 .loop  RDLONG  value, ptra++
3       ADD    sum, value
4       DJNZ   count, #.loop
5
6 ' FIFO reading: 2 clocks per long!
7       RDFAST #0, ptra      ' Start FIFO
8 .loop  RFLONG  value      ' 2 clocks, always!
9       ADD    sum, value   ' 2 clocks
10      DJNZ   count, #.loop ' 2 clocks
11      ' 3x faster for sequential reads!

```

## Parallel Operations

CORDIC operations can overlap with other work:

```

1 ' CORDIC overlaps with other instructions
2     QMUL    x, y          ' Start 32x32->64 multiply (CORDIC)
3     ' 55 clocks to do other work!
4     ADD    a, b          ' These execute during CORDIC
5     SUB    c, d
6     MOV    index, #0
7     RDLONG data, ptra++
8     ' ... more work
9     GETQX  low_result    ' Get CORDIC result (lower 32 bits)
10    GETQY  high_result   ' Get CORDIC result (upper 32 bits)
11
12 ' QROTATE overlap
13    QROTATE x_coord, angle ' Start rotation (D=X, S=angle)
14    ' 55 clocks of other work!
15    GETQX  new_x         ' Get rotated X
16    GETQY  new_y         ' Get rotated Y

```

Note: **MUL**/**MULS** are 2-clock ALU instructions that complete immediately (16x16->32). Use **QMUL** for 32x32->64 with CORDIC overlap.

## Real-World Example: Fast Memory Copy

Let's optimize a memory copy routine:

```
1  ' Version 1: Basic (slow)
2  copy_basic
3      RDLONG  temp, source
4      WRLONG  temp, dest
5      ADD     source, #4
6      ADD     dest, #4
7      DJNZ   count, #copy_basic
8      ' ~13 clocks per long
9
10 ' Version 2: Better pointers
11 copy_better
12     RDLONG  temp, ptrb++
13     WRLONG  temp, ptrb++
14     DJNZ   count, #copy_better
15     ' ~8 clocks per long
16
17 ' Version 3: Block transfer (ultimate)
18 copy_ultimate
19     SUB     count, #1      ' SETQ needs count-1 (0 = 1 long)
20     SETQ   count          ' Setup block transfer
21     RDLONG buffer, source ' Read all at once
22     SETQ   count          ' (count already decremented)
23     WRLONG buffer, dest   ' Write all at once
24     ' ~1 clock per long for large blocks (hub maximum)!
```

### The Medicine Cabinet

Optimization overwhelming you? Start with these simple improvements:

#### Three easy wins:

1. Use **PTRA/PTRB** instead of manual pointer math

```

1 ' Slow
2     RDLONG x, addr
3     ADD   addr, #4
4
5 ' Fast
6     RDLONG x, ptra++

```

2. Align your data to **LONG** boundaries

```

1     ALIGNL           ' Force long alignment
2 data LONG   $12345678

```

3. Use **REP** for tight loops (note: hub-address expressions like `ptra++` only work with hub-access instructions, so we read first, then add)

```

1     REP   @.end, count
2     RDLONG val, ptra++   ' Fetch next long from hub
3     ADD   sum, val       ' Accumulate
4 .end

```

Just these three changes often double performance!

## Your Turn: Optimization Challenges

### Your Turn

**Your Turn:** Optimize a checksum calculator

Starting code:

*continues on next page →*

```

1 ' Slow version
2 checksum_slow
3     MOV     sum, #0
4     MOV     addr, ##buffer
5     MOV     count, #256
6
7 .loop  RDBYTE temp, addr
8     ADD     sum, temp
9     ADD     addr, #1
10    DJNZ   count, #.loop

```

Goal: Make it at least 4x faster Hint: Read longs instead of bytes, use FIFO Success Check: Same checksum, much faster

## Advanced Techniques

### Instruction Pairing

Some instruction pairs execute specially:

```

1 ' ## syntax handles AUGS automatically
2     MOV     x, ##$12345678 ' Assembler generates AUGS + MOV
3     ' Same result, cleaner code!
4
5 ' ALTD + instruction = indirect addressing
6     ALTD   index, #array
7     MOV    0-0, value     ' Stores to array[index]

```

### Pipeline-Aware Coding

Avoid pipeline stalls:

```

1 ' Bad: result needed immediately
2     ADD     x, y
3     CMP     x, #10 wcz     ' Stall waiting for x
4
5 ' Good: interleave operations

```

```

6      ADD    x, y
7      MOV    a, b      ' Do something else
8      CMP    x, #10 wcz  ' Now x is ready

```

## Unrolling Loops

Sometimes removing the loop is faster. (Remember: `ptr++` only works with hub-access instructions like **RDLONG**, so each iteration reads then **ADDS**.)

```

1  ' Looped version
2      REP    @.end, #4
3      RDLONG val, ptr++
4      ADD    sum, val
5  .end
6
7  ' Unrolled version (faster for small counts)
8      RDLONG val, ptr++
9      ADD    sum, val
10     RDLONG val, ptr++
11     ADD    sum, val
12     RDLONG val, ptr++
13     ADD    sum, val
14     RDLONG val, ptr++
15     ADD    sum, val

```

## Common Optimization Gotchas

Before you rewrite everything in **REP** and **SKIP**, a few sanity checks:

1. **Premature optimization** - Get it working first, then optimize
2. **Over-optimizing** - Sometimes clarity is worth 2 clocks
3. **Ignoring the big picture** - Optimize the bottleneck, not everything
4. **Breaking functionality** - Fast but wrong is useless
5. **Forgetting about power** - Faster isn't always better for battery life

## Profiling and Measurement

Always measure your optimizations:

```

1  ' Time your code
2      GETCT    start_time
3
4      ' Code to measure
5      CALL    #function_to_test
6
7      GETCT    end_time
8      SUB     end_time, start_time
9      ' end_time now contains exact clock cycles

```

**Pitfall — CT wraps:** **GETCT** reads a 32-bit free-running counter that wraps every ~21.5 seconds at 200 MHz ( $2^{32} \div 200$  MHz). For short measurements like the one above, the **SUB** trick masks the wrap correctly thanks to two’s-complement arithmetic. But for a *scheduler* or *timer* running over minutes, hours, or days, you need one of two strategies:

1. **Capture the full 64-bit count.** **GETCT D WC** returns the upper 32 bits (with **WC** set); a plain **GETCT D** returns the lower 32. Read upper-then-lower-then-upper-again and retry if the upper **WORD** changed, then build a 64-bit timestamp.
2. **Work with deltas, not absolute time.** Compare `(now - start)` rather than `now > deadline`. The subtraction wraps correctly even when the counter does.

The 21.5-second wrap is *fast* — long enough that you won’t see it in a basic example, short enough that real applications hit it constantly.

## What We’ve Learned

You’re now an optimization expert:

- Understanding the P2 pipeline
- Instruction timing knowledge
- **REP** and **SKIP** for zero-overhead loops
- FIFO for maximum throughput
- Parallel operation techniques
- Real-world optimization strategies

## Coming Up Next

Chapters 13-15 provide quick examples of Video Generation, Serial Protocols, and Signal Processing - with references to dedicated manuals for deep dives. Think of them as appetizers showing what’s possible!

---

**Have Fun!** Remember, the best optimization is often a better algorithm. But when you need every last cycle, you now know how to get them!

# Chapter 13: LUT Memory - Your Private Lookup Table

*512 longs of fast, deterministic storage in every COG*

## The Hook: A Lookup Table in 3 Cycles

Need fast data lookup without hub timing? Every COG has its own private 512-long Lookup RAM (LUT):

```

1 ' Sine table lookup - 3 clocks, every time
2 get_sine
3     AND    angle, #$FF    ' Mask to table index
4     RDLUT  value, angle    ' Read from LUT in 3 clocks!
5     RET

```

No hub timing to worry about. No waiting for the egg beater. Just 3 clock cycles, guaranteed. The LUT is like having a personal data assistant that never takes a coffee break.

## Why Another Memory?

You might be thinking, “Wait, I already have COG RAM and Hub RAM - why do I need a third memory?” Excellent question!

Memory	Size per COG	Access Time	Special Features
COG RAM	512 longs	2 clocks	Instructions live here
Hub RAM	512 KB shared	2-9 clocks (hub slot wait)	Shared by all COGs
<b>LUT RAM</b>	512 longs	<b>3 clocks</b>	<b>Private, deterministic, shareable with neighbor</b>

The LUT fills a sweet spot: faster than hub memory, doesn't compete with your instruction space, and has a trick up its sleeve - neighboring COGs can share LUTs!

## Reading and Writing the LUT

### Basic LUT Access

```

1 ' Write to LUT
2     WRLUT    ##$12345678, #100 ' Write 32-bit constant to LUT[100]
3     WRLUT    value, index      ' Write variable to LUT[index]
4
5 ' Read from LUT
6     RDLUT    result, #100      ' Read LUT[100] into result
7     RDLUT    data, index       ' Read LUT[index] into data

```

Notice the operand order: **WRLUT** writes its first operand to the address in the second, while **RDLUT** reads from its second operand into the first. A bit backwards from what you might expect, but you'll get used to it.

### Building a Lookup Table

Here's how to load a sine table into LUT:

```

1 ' Copy 256-entry sine table from hub to LUT
2 load_sine_table
3     MOV      index, #0
4     LOC      ptra, #\sine_data_hub ' Hub address of table
5
6 .loop  RDLONG  value, ptra++      ' Read from hub
7     WRLUT    value, index        ' Write to LUT
8     ADD      index, #1
9     CMP      index, #256 wz
10    if_nz   JMP      #.loop
11    RET
12
13 ' Now lookups are fast!
14 get_sine
15     RDLUT    sine_value, angle   ' 3 clocks!
16     RET

```

### Sidetrack

#### Bulk LUT Loading with SETQ2

For loading entire tables, **SETQ2** + **RDLONG** transfers hub data into LUT memory. The trick: when **SETQ2** is active, the **RDLONG** destination field (still 9-bit, 0-\$1FF) is interpreted as a LUT offset that maps physically to \$200-\$3FF:

```

1      SETQ2   #256-1           ' 256 longs to LUT
2      RDLONG  $000, hub_table_ptr ' Dest field 0 = physical LUT $200

```

The destination operand is the LUT offset, not the absolute address — so you write \$000 (LUT base), not \$200 (which would overflow the 9-bit field). Remember the -1 in **SETQ2** (same rule as **SETQ** for hub block transfers).

## LUT Sharing Between COGs

Here's something clever: adjacent COG pairs can share LUT data! When you enable LUT sharing with **SETLUTS**, writes your neighbor makes to their LUT are automatically *copied* to your LUT too.

```

1 ' --- COG 1 (consumer) - MUST enable sharing FIRST ---
2     SETLUTS #1           ' Enable LUT write copying FROM COG 0
3     ' Now when COG 0 writes to its LUT, data is COPIED to our LUT
4
5 ' --- COG 0 (producer) - writes AFTER consumer enables sharing ---
6     WRLUT  message, #10  ' Write MY LUT[10] (copies to COG 1)
7     WRLUT  #1, #0        ' Set ready flag (copies to COG 1)
8
9 ' --- COG 1 (consumer) - reads its OWN LUT (which contains copies) ---
10 .wait  RDLUT  flag, #0   ' Read MY LUT[0] (contains copy from COG 0)
11      CMP    flag, #1 wz
12  if_nz  JMP    #.wait
13      RDLUT  message, #10 ' Read MY LUT[10] (copied from COG 0)

```

The key instruction is:

- **SETLUTS**: Enable write copying - when neighbor writes with **WRLUT**, data is copied to YOUR LUT
- **RDLUT**: Read your own LUT (which now contains copied data)

Important: The consumer COG must enable **SETLUTS** *before* the producer writes, otherwise the writes won't be copied!

This gives you a 512-long shared buffer between COG pairs without touching hub memory. Perfect for high-bandwidth data passing!

### Sidetrack

#### Which COGs Are Neighbors?

The LUT sharing pairs are fixed:

- COG 0 COG 1
- COG 2 COG 3
- COG 4 COG 5
- COG 6 COG 7

When you enable sharing, your odd/even companion COG's LUT *writes* are copied into your own LUT — so you read the copies from your own LUT. Sharing works only within the fixed pair; non-adjacent COGs cannot share.

## Practical Examples

Enough theory — let's see what people actually use the LUT for in real code:

### Fast Data Transformation

```
1 ' Gamma correction table in LUT
2 ' Input: 8-bit value in 'pixel'
3 ' Output: Gamma-corrected value
4 gamma_correct
5     AND    pixel, #$FF    ' Mask to 8 bits
6     RDLUT  pixel, pixel   ' Transform via table
7     RET
8
9 ' Initialize gamma table (power law curve)
10 ' Would be pre-calculated and loaded from hub
```

## Circular Buffer in LUT

```

1  ' Fast circular buffer using LUT
2  ' 256-entry buffer at LUT addresses 0-255
3
4  buf_write_ptr  LONG    0
5  buf_read_ptr   LONG    0
6
7  put_byte
8      WRLUT    data, buf_write_ptr
9      ADD     buf_write_ptr, #1
10     AND     buf_write_ptr, #$FF    ' Wrap at 256
11     RET
12
13 get_byte
14     RDLUT   data, buf_read_ptr
15     ADD     buf_read_ptr, #1
16     AND     buf_read_ptr, #$FF    ' Wrap at 256
17     RET

```

## Fast Stack in LUT

```

1  ' Stack implementation in LUT
2  ' Grows downward from $1FF
3  stack_ptr      LONG    $1FF
4
5  PUSH
6      WRLUT   value, stack_ptr
7      SUB    stack_ptr, #1
8      RET
9
10 POP
11     ADD    stack_ptr, #1
12     RDLUT  value, stack_ptr
13     RET

```

## LUT with the Streamer

Here's where LUT gets really interesting. The Streamer can read directly from LUT to generate waveforms without any COG intervention:

```

1  ' Fill LUT with waveform data
2  ' Then let Streamer output it to DAC
3
4  load_waveform
5      MOV    index, #0
6
7  .fill  MOV    value, index
8      SHL    value, #24      ' Scale for DAC
9      WRLUT  value, index
10     ADD    index, #1
11     CMP    index, #512 wz
12  if_nz JMP    #.fill
13
14  ' Now configure Streamer to read from LUT
15  ' Streamer handles the rest - no COG cycles needed!

```

The Streamer configuration for LUT reading is covered in detail in the Video and Audio manuals - but the key point is that your LUT becomes a 512-sample waveform buffer that plays automatically.

## Common Gotchas

### WRONG: Confusing LUT addresses

```

1  ' WRONG - This reads COG RAM, not LUT!
2      MOV    value, $200      ' $200 is COG RAM address

```

### RIGHT: Use RDLUT for LUT access

```

1  ' RIGHT - RDLUT addresses the LUT space
2      RDLUT  value, #0        ' LUT address 0

```

### WRONG: Reading LUT before neighbor writes

```

1  ' WRONG - No data to read yet!
2      SETLUTS #1              ' Enable sharing
3      RDLUT  data, #10        ' Empty - neighbor hasn't written!

```

**RIGHT: Wait for neighbor's write signal**

```

1 ' RIGHT - Wait for data to be copied
2     SETLUTS #1           ' Enable sharing BEFORE neighbor writes
3 .wait RDLUT  ready, #0   ' Check flag in MY LUT
4     TJZ      ready, #.wait ' Wait until neighbor writes
5     RDLUT   data, #10    ' Now MY LUT has copied data

```

**Medicine Cabinet****The Medicine Cabinet****LUT Memory Quick Reference**

Instruction	Operation	Cycles
<b>RDLUT</b> D, S	Read LUTS into D	3
<b>WRLUT</b> D, S	Write D to LUTS	2
<b>SETLUTS</b> D	Enable LUT write copying (D[0]=1)	2

**Memory Map:**

- LUT addresses: 0-511 (512 longs = 2KB)
- Neighbor pairs: 0 1, 2 3, 4 5, 6 7

**Best Uses:**

- Lookup tables (sine, gamma, encoding)
- Fast circular buffers
- COG-pair data sharing
- Streamer waveform source

**Your Turn****Your Turn****Exercise 1: Build an 8-bit Encoder**

Create a LUT-based ASCII to 7-segment display encoder. Load a 128-entry table where each entry maps an ASCII code to the 7-segment pattern for that character.

*continues on next page →*

```
1 ' Your code here:
2 ' 1. Load segment patterns into LUT
3 ' 2. Write encode_char routine
4 ' Hint: rdlut segment_pattern, ascii_char
```

## Your Turn

### Exercise 2: High-Speed COG Communication

Use LUT sharing to create a message passing system between COG 2 and COG 3:

- COG 2 writes 8-long messages
- COG 3 reads them without hub access
- Use a simple ready/ack protocol

```
1 ' Hint: Use LUT[0] as ready flag, LUT[1-8] as message buffer
```

## What We've Learned

The LUT in your toolbox:

- 512 longs of fast, private memory in every COG
- 3-clock deterministic access via **RDLUT** / **WRLUT**
- Bulk loading via **SETQ2** + **RDLONG**
- COG-pair LUT sharing for high-bandwidth data passing
- Streamer source for waveform generation

## Coming Up Next

Chapter 14 hands you the keys to 64 autonomous I/O processors — Smart Pins. We'll cover the universal configuration pattern and the modes you'll reach for most often, then point you at the dedicated Smart Pins Manual for the deep dive.

**Have Fun!** And remember — a 3-clock private lookup table is a luxury most chips don't give you. Use it!

# Chapter 14: Smart Pins Orientation

*64 autonomous I/O processors waiting to do your bidding*

## The Hook: A UART in 4 Lines

Remember that tedious bit-bang serial from Chapter 8? Watch this:

```

1  ' Configure pin as UART transmitter - done!
2      DIRL      #TX_PIN          ' Reset pin first!
3      WRPIN     ##P_ASYNC_TX, #TX_PIN ' Configure as async TX
4      WXPIN     ##BAUD_115200, #TX_PIN ' Set baud rate
5      DIRH      #TX_PIN          ' Enable - runs on its own

```

That's it. The pin is now a fully autonomous UART transmitter. It handles start bits, stop bits, timing - everything. You just feed it bytes with **WYPIN** and it sends them. The pin has become a state machine.

And here's the mind-bending part: *every single one of the 64 pins can do this*. Or PWM. Or ADC. Or quadrature decoding. Or 28 other modes.

## What Are Smart Pins, Really?

Each of the P2's 64 I/O pins contains its own little processor - a state machine that can operate completely independently of the COGs. This means:

- A pin configured as UART keeps sending/receiving without COG intervention
- A PWM output keeps running its duty cycle automatically
- An ADC samples continuously in the background
- A quadrature decoder tracks position even while your COG does other things

The COG only needs to configure the pin and occasionally read/write data. The pin does the rest.

## The Universal Smart Pin Pattern

Every Smart Pin follows the same configuration pattern. This is **the most important thing to remember**:

```

1  ' === THE SMART PIN RECIPE ===
2

```

*continues on next page →*

```

3 ' Step 1: RESET the pin (CRITICAL!)
4     DIRL    pin           ' Always start by resetting
5
6 ' Step 2: CONFIGURE the mode
7     WRPIN  mode, pin     ' What should this pin do?
8
9 ' Step 3: SET parameters
10    WXPIN  x_value, pin  ' Mode-specific parameter X
11    WYPIN  y_value, pin  ' Mode-specific parameter Y
12
13 ' Step 4: ENABLE the pin
14    DIRH   pin           ' Start the magic!

```

### Sidetrack

#### Why DIRL First?

The **DIRL** at the start isn't optional politeness - it's *required*. Smart Pins must be re-set before configuration to ensure they're in a known state. **SKIP** this and you'll get unpredictable behavior as old settings conflict with new **ONES**.

Think of it like power-cycling a misbehaving device. Always start fresh.

## The Core Instructions

### Configuration Instructions

Instruction	Purpose
<b>WRPIN</b> mode, pin	Set the operating mode
<b>WXPIN</b> value, pin	Set X parameter (mode-specific)
<b>WYPIN</b> value, pin	Set Y parameter (mode-specific)
<b>DIRH</b> pin	Enable the Smart Pin
<b>DIRL</b> pin	Disable/reset the Smart Pin

## Data Instructions

Instruction	Purpose
<b>WYPIN</b> data, pin	Write data to Smart Pin (same instruction!)
<b>RDPIN</b> data, pin	Read result, clear “ready” flag
<b>RQPIN</b> data, pin	Read result, keep “ready” flag
<b>AKPIN</b> pin	Acknowledge (clear “ready” flag only)

## Status Instructions

Instruction	Purpose
<b>TESTP</b> pin WC	Check if IN flag is set (data ready)
<b>TESTPN</b> pin WC	Check if IN flag is clear

## Understanding the IN Flag

Every Smart Pin has an IN flag that signals “something happened.” What that something is depends on the mode:

- **UART TX**: IN high = ready for another **BYTE**
- **UART RX**: IN high = **BYTE** received
- **ADC**: IN high = new sample ready
- **PWM**: IN high = period complete
- **Counter**: IN high = threshold reached

You check this flag with **TESTP** and clear it by reading with **RDPIN** (or explicitly with **AKPIN**).

```

1 ' Wait for Smart Pin to be ready
2 wait_ready
3     TESTP    #PIN wc        ' Check IN flag
4     if_nc   JMP    #wait_ready    ' Loop if not ready
5     RDPIN   data, #PIN        ' Read and clear flag

```

## Sidetrack

### Event-Driven Alternative

Instead of polling with **TESTP**, you can use the event system:

```

1 SETSE1  #%001<<6 + PIN    ' Event when IN rises
2 WAITSE1                               ' Sleep until ready - no polling!
3 RDPIN   result, #PIN      ' Read the result

```

This is more efficient because your COG sleeps instead of spinning. See Chapter 15 for the full event story.

## Common Smart Pin Modes

Here are the modes you'll use most often:

### Asynchronous Serial (UART)

```

1 ' Transmit mode
2     DIRL    #TX_PIN
3     WRPIN   ##P_ASYNC_TX | P_OE, #TX_PIN
4     WXPIN   ##(clkfreq/ baud)<<16 | 7, #TX_PIN ' Baud + 8 bits
5     DIRH    #TX_PIN
6
7 ' Send a byte
8 send     TESTP   #TX_PIN wc      ' Wait for ready
9   if_nc  JMP     #send
10        WYPIN   BYTE, #TX_PIN    ' Send it

```

```

1 ' Receive mode
2     DIRL    #RX_PIN
3     WRPIN   ##P_ASYNC_RX, #RX_PIN
4     WXPIN   ##(clkfreq/ baud)<<16 | 7, #RX_PIN
5     DIRH    #RX_PIN
6
7 ' Get a byte
8 recv     TESTP   #RX_PIN wc      ' Check for received byte
9   if_nc  JMP     #recv

```

*continues on next page →*

```

10      RDPIN   BYTE, #RX_PIN   ' Get it
11      SHR     BYTE, #24      ' Shift to low byte

```

## PWM Output

```

1  ' PWM mode - period + duty cycle
2      DIRL     #PWM_PIN
3      WRPIN   ##P_PWM_SAWTOOTH | P_OE, #PWM_PIN
4      WXPIN   ##period, #PWM_PIN      ' Period in clocks
5      WYPIN   ##duty, #PWM_PIN       ' High time in clocks
6      DIRH     #PWM_PIN
7
8  ' Change duty cycle on the fly
9      WYPIN   ##new_duty, #PWM_PIN   ' Just update Y parameter

```

## ADC Input

```

1  ' ADC mode - continuous sampling
2      DIRL     #ADC_PIN
3      WRPIN   ##P_ADC | P_ADC_GIO, #ADC_PIN
4      WXPIN   ##13, #ADC_PIN         ' 14-bit mode (period = 2^13 clocks)
5      DIRH     #ADC_PIN
6
7  ' Read ADC value
8  read_adc
9      RDPIN   adc_value, #ADC_PIN    ' Get N-bit ADC count (LSB-aligned)

```

## Quadrature Encoder

```

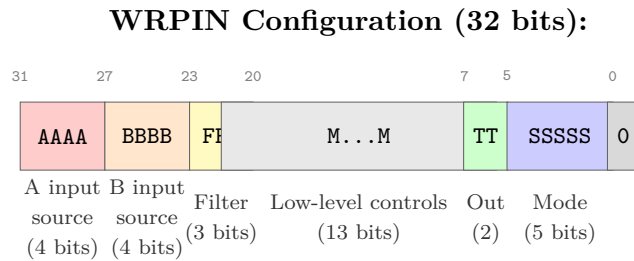
1  ' Quadrature decoder - A on pin, B on pin+1
2      DIRL     #ENC_PIN
3      WRPIN   ##P_QUADRATURE, #ENC_PIN
4      DIRH     #ENC_PIN
5
6  ' Read position
7      RDPIN   position, #ENC_PIN    ' Get accumulated count

```

## Configuration Values Demystified

Don't worry, you don't have to memorize all 32 mode bit-patterns. The mode values like `P_ASYNC_TX` are constants defined by the assembler. But here's what's happening behind the scenes, in case you're curious:

The **WRPIN** D value is a 32-bit configuration:



For most common modes, you'll use predefined constants like `P_ASYNC_TX`, `P_PWM_SAWTOOTH`, `P_ADC`. The P2 assembler knows all of them.

## Common Gotchas

### WRONG: Forgetting to reset before configure

```

1 ' WRONG - Pin may be in unknown state!
2     WRPIN    ##P_PWM_SAWTOOTH, #PIN
3     WXPIN    ##1000, #PIN
4     DIRH     #PIN

```

### RIGHT: Always DIRL first

```

1 ' RIGHT - Start clean
2     DIRL     #PIN                ' Reset first!
3     WRPIN    ##P_PWM_SAWTOOTH, #PIN
4     WXPIN    ##1000, #PIN
5     DIRH     #PIN

```

**WRONG: Enabling before configuring**

```
1 ' WRONG - Pin enabled with partial config!
2     DIRL    #PIN
3     DIRH    #PIN                ' Enabled too early!
4     WRPIN   ##P_ASYNC_TX, #PIN
```

**RIGHT: DIRH comes last**

```
1 ' RIGHT - Configure completely, then enable
2     DIRL    #PIN
3     WRPIN   ##P_ASYNC_TX, #PIN
4     WXPIN   ##BAUD, #PIN
5     DIRH    #PIN                ' Enable last!
```

**Medicine Cabinet**

## The Medicine Cabinet

### Smart Pin Quick Reference

#### The Recipe:

1. **DIRL** pin — Reset the pin first
2. **WRPIN** mode, pin — Set the operating mode
3. **WXPIN** x, pin — Set X parameter
4. **WYPIN** y, pin — Set Y parameter
5. **DIRH** pin — Enable the Smart Pin

#### Common Modes:

- **UART TX:** P\_ASYNC\_TX — Serial transmit
- **UART RX:** P\_ASYNC\_RX — Serial receive
- **PWM:** P\_PWM\_SAWTOOTH — Sawtooth wave output
- **PWM:** P\_PWM\_TRIANGLE — Triangle wave output
- **ADC:** P\_ADC — Analog input
- **Quadrature:** P\_QUADRATURE — Encoder
- **NCO:** P\_NCO\_FREQ — Frequency output

#### Data Flow:

- **WYPIN** = Write data TO Smart Pin
- **RDPIN** = Read data FROM Smart Pin (clears IN)
- **TESTP** = Check if IN flag set

**Golden Rule:** **DIRL** before **WRPIN**, **DIRH** after **WXPIN**/**WYPIN**

## Your Turn

### Your Turn

#### Exercise 1: PWM LED Dimmer

Create a PWM output that dims an LED:

1. Configure a pin for PWM sawtooth mode
2. Set a 1 kHz period (at 160 MHz: period = 160,000)
3. Vary duty cycle from 0% to 100%

```

1 ' Your code here:
2 ' Hint: Change duty with WYPIN new_duty, #LED_PIN

```

## Your Turn

### Exercise 2: Simple Serial Echo

Set up UART at 115200 baud:

1. Configure RX on pin 63
2. Configure TX on pin 62
3. Echo every received **BYTE** back

```
1 ' Your code here:
2 ' At 160 MHz: baud_divisor = 160_000_000 / 115200 = 1389
3 ' WXPIN format: (divisor << 16) | (bits - 1)
```

**Going Deeper:** This chapter covered the Smart Pin essentials - the configuration pattern and common modes. For complete coverage of all 32 modes, timing diagrams, and advanced techniques, see the dedicated “P2 Smart Pins Manual.”

## What We’ve Learned

Smart Pin essentials:

- Every pin contains its own state machine (32 modes available)
- The universal recipe: **DIRL** → **WRPIN** → **WXPIN** → **WYPIN** → **DIRH**
- The IN flag signals “something happened” (mode-specific)
- UART, PWM, ADC, quadrature — same configuration pattern
- Smart Pins free the COG for other work

## Coming Up Next

Chapter 15 explores the event system — how to stop polling and start waiting, so your COG sleeps until something interesting happens. The companion to Smart Pins: when one tells the other a byte is ready, you want to be notified, not spinning.

**Have Fun!** And remember — every Smart Pin you configure is a coprocessor you don’t have to babysit. That’s leverage!

# Chapter 15: Event-Driven Programming

*Stop spinning, start waiting*

## The Hook: No More Polling Loops

Remember all those busy loops waiting for things to happen?

```
1 ' OLD WAY: Spin waiting for serial data (burns CPU cycles!)
2 wait_rx TESTP   #RX_PIN wc      ' Check over and over
3   if_nc JMP     #wait_rx        ' Spin spin spin...
4       RDPIN    data, #RX_PIN
5
6 ' NEW WAY: Sleep until data arrives (zero CPU cycles!)
7       SETSE1   #%001<<6 + RX_PIN ' Wake on IN rise
8       WAITSE1                               ' Sleep until event
9       RDPIN    data, #RX_PIN
```

The event system lets your COG sleep while waiting. When the event happens, it wakes up instantly. No cycles wasted, and you respond the moment something happens.

## Why Events Matter

Polling loops have two problems:

1. **They waste cycles** - The COG spins doing nothing useful
2. **They add latency** - You check periodically, so there's delay between "thing happened" and "you noticed"

The event system solves both. Your COG *sleeps* and *wakes the instant* something happens. It's like having a personal assistant tap your shoulder instead of constantly looking up to check.

## The Four Selectable Events

Let's meet the cast. Every COG has four configurable event channels: SE1, SE2, SE3, and SE4. Each can be configured to trigger on different conditions:

Event	Configuration	Wait	Poll
SE1	<b>SETSE1</b>	<b>WAITSE1</b>	<b>POLLSE1</b>
SE2	<b>SETSE2</b>	<b>WAITSE2</b>	<b>POLLSE2</b>
SE3	<b>SETSE3</b>	<b>WAITSE3</b>	<b>POLLSE3</b>
SE4	<b>SETSE4</b>	<b>WAITSE4</b>	<b>POLLSE4</b>

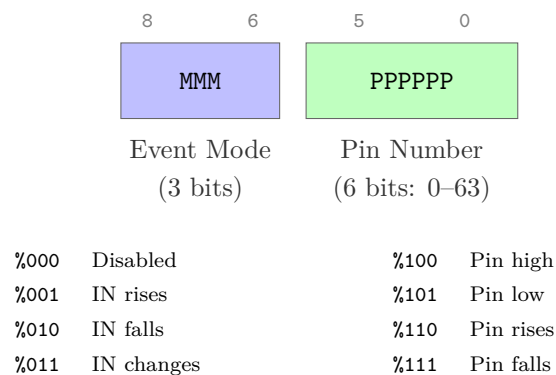
Plus there are built-in timer events:

Timer	Wait	Poll
CT1	<b>WAITCT1</b>	<b>POLLCT1</b>
CT2	<b>WAITCT2</b>	<b>POLLCT2</b>
CT3	<b>WAITCT3</b>	<b>POLLCT3</b>

## Configuring an Event

The **SETSE1** through **SETSE4** instructions take a 9-bit configuration value:

### SETSE1/2/3/4 Configuration (9 bits):



## Event Modes

Mode	Meaning
%000	Never (disabled)
%001	IN rises (Smart Pin ready)
%010	IN falls
%011	IN changes
%100	Pin is low (level)
%101	Pin is low (level)
%110	Pin is high (level)
%111	Pin is high (level)

## EVENT\_\* Constants: When You Need Interrupts

While dedicated COGs are usually better than interrupts (see Chapter 11), sometimes you need them. The **SETINT1/2/3** instructions select which event triggers an interrupt using these constants:

Table 15.1.

Constant	Value	Description
EVENT_INT	%0000	Pin matches interrupt configuration
EVENT_CT1	%0001	CT equals CT1 (timer 1 target)
EVENT_CT2	%0010	CT equals CT2 (timer 2 target)
EVENT_CT3	%0011	CT equals CT3 (timer 3 target)
EVENT_SE1	%0100	Selectable event 1 triggered
EVENT_SE2	%0101	Selectable event 2 triggered
EVENT_SE3	%0110	Selectable event 3 triggered
EVENT_SE4	%0111	Selectable event 4 triggered
EVENT_PAT	%1000	SETPAT pattern detected
EVENT_FBW	%1001	Hub FIFO wrapped around
EVENT_XMT	%1010	Streamer needs data
EVENT_XFI	%1011	Streamer operation complete
EVENT_XRO	%1100	NCO frequency counter rolled

Continued on next page

Table 15.1. (Continued)

Constant	Value	Description
EVENT_XRL	%1101	Streamer read last LUT location (\$1FF)
EVENT_ATN	%1110	Another COG signaled attention
EVENT_QMT	%1111	CORDIC/PIX math complete

### Using EVENT\_\* with SETINT:

```

1 ' Enable INT1 when SE1 event occurs
2     SETSE1  %#001<<6 + RX_PIN      ' SE1 = IN rise on RX_PIN
3     SETINT1 #EVENT_SE1             ' INT1 fires when SE1 triggers
4
5 ' Enable INT2 on timer match
6     ADDCT2  target, ##200_000      ' Set timer 2 target
7     SETINT2 #EVENT_CT2             ' INT2 fires when CT = CT2
8
9 ' Enable INT3 when another COG signals
10    SETINT3 #EVENT_ATN             ' INT3 fires on COGATN

```

**Pro tip:** You can also use these constants with WAITSE/POLLSE by first triggering them with the appropriate hardware condition, then waiting. But for most purposes, the SETSE mode bits (the table above this one) are what you'll configure directly.

### Smart Pin Events

The most common use is waiting for a Smart Pin to have data:

```

1 ' Wait for Smart Pin on pin 15 to be ready
2     SETSE1  %#001<<6 + 15      ' IN rise on pin 15
3     WAITSE1                               ' Sleep until ready
4     RDPIN   data, #15          ' Get the data

```

### Pin Edge Events

You can also wait for raw pin edges (without Smart Pin):

```

1 ' Wait for rising edge on pin 5
2     SETSE1  %#001<<6 + 5      ' Rising edge on pin 5
3     WAITSE1                               ' Sleep until edge

```

*continues on next page →*

```

4         ' Edge detected!
5
6 ' Wait for falling edge on pin 10
7     SETSE2  %#010<<6 + 10    ' Falling edge on pin 10
8     WAITSE2
9     ' Edge detected!

```

## Timer Events

For precise timing, use the counter comparison events:

```

1 ' Wait exactly 1 millisecond (at 200 MHz)
2     GETCT  target            ' Current time
3     ADD    target, ##200_000 ' +1ms at 200MHz
4     ADDCT1 target, #0        ' Set CT1 target
5     WAITCT1                    ' Sleep until CT >= CT1
6
7 ' Alternative using WAITX (simpler but less precise)
8     WAITX  ##200_000        ' Wait ~1ms at 200MHz

```

The timer events are:

- **ADDCT1/ADDCT2/ADDCT3:** Set the comparison target
- **WAITCT1/WAITCT2/WAITCT3:** Wait until CT reaches target
- **POLLCT1/POLLCT2/POLLCT3:** Check (non-blocking) if target reached

## Waiting vs Polling

Two ways to use events:

### WAIT - Sleep Until Event

```

1     WAITSE1                    ' COG sleeps here
2     ' Wakes instantly when event occurs

```

- COG sleeps, uses no cycles
- Wakes immediately when event fires
- Can't do anything else while waiting

## POLL - Check and Continue

```

1      POLLSE1 wc          ' Check event, clear if set
2  if_c JMP      #event_handler ' Handle if occurred
3      ' Continue with other work...
```

- COG keeps running
- Checks event flag, clears it
- Returns result in C flag
- Good for servicing multiple events

## Multiple Events

With four SE channels, you can monitor multiple sources:

```

1 ' Setup multiple events
2     SETSE1  %#001<<6 + RX_PIN      ' Serial data ready
3     SETSE2  %#001<<6 + BUTTON_PIN ' Button pressed (rising edge)
4     SETSE3  %#001<<6 + ADC_PIN     ' ADC sample ready
5
6 event_loop
7     POLLSE1 wc          ' Check serial
8  if_c CALL    #handle_serial
9
10    POLLSE2 wc          ' Check button
11  if_c CALL    #handle_button
12
13    POLLSE3 wc          ' Check ADC
14  if_c CALL    #handle_adc
15
16    JMP      #event_loop
```

## Practical Examples

Time to put events to work. Three patterns you'll reach for again and again:

## Timeout with Fallback

```

1  ' Wait for serial data, but give up after 100ms
2  wait_with_timeout
3      SETSE1  %#001<<6 + RX_PIN      ' Serial ready event
4
5      GETCT   timeout
6      ADD     timeout, ##16_000_000  ' 100ms at 160MHz
7      ADDCT1  timeout, #0
8
9  .wait     POLLSE1 wc                ' Check serial
10  if_c     JMP     #.got_data
11         POLLCT1 wc                ' Check timeout
12  if_c     JMP     #.timed_out
13         JMP     #.wait
14
15 .got_data
16         RDPIN   data, #RX_PIN
17         RET
18
19 .timed_out
20         MOV     data, #-1          ' Return error
21         RET

```

## Debounced Button Press

```

1  ' Wait for clean button press with debounce
2  debounced_button
3      SETSE1  %#001<<6 + BUTTON      ' Rising edge
4      WAITSE1                                ' Wait for press
5
6      WAITX   ##2_000_000            ' 10ms debounce at 200MHz
7
8      TESTP   #BUTTON wc            ' Verify still pressed
9  if_nc     JMP     #debounced_button ' Bounce - try again
10         RET     ' Clean press!

```



**RIGHT: Poll first to clear any pending event**

```
1 ' RIGHT - Clear any stale event
2     SETSE1  %#001<<6 + PIN
3     POLLSE1                ' Clear if already set
4     ' ... do other stuff ...
5     WAITSE1                ' Now wait cleanly
```

**WRONG: Using WAIT when you need to handle multiple sources**

```
1 ' WRONG - Can only wait for one event at a time
2     WAITSE1                ' Stuck here until SE1
3     ' SE2 might fire and be missed!
```

**RIGHT: Use POLL loop for multiple events**

```
1 ' RIGHT - Check all sources
2 .loop  POLLSE1  wc
3     if_c  CALL    #handle_se1
4         POLLSE2  wc
5     if_c  CALL    #handle_se2
6         JMP      #.loop
```

## Medicine Cabinet

## The Medicine Cabinet

### Event System Quick Reference

#### Configure Events:

```
1      SETSE1/2/3/4  #%MMM_PPPPPP  ' Mode and pin
```

#### Event Modes:

---

%MMM	Trigger
------	---------

---

%001	IN rises (Smart Pin ready)
------	----------------------------

%010	IN falls
------	----------

%011	IN changes
------	------------

%10x	Pin is low (level)
------	--------------------

%11x	Pin is high (level)
------	---------------------

---

#### Wait (blocking):

```
1      WAITSE1/2/3/4  ' Sleep until event
2      WAITCT1/2/3    ' Sleep until timer
3      WAITATN        ' Sleep until attention
```

#### Poll (non-blocking):

```
1      POLLSE1/2/3/4 WC ' Check event, clear flag, C=occurred
2      POLLCT1/2/3 WC  ' Check timer, C=reached
3      POLLATN WC      ' Check attention, C=received
```

#### Timer Setup:

```
1      ADDCT1/2/3 target, #delta  ' Set comparison target
```

#### Inter-COG:

```
1      COGATN #mask  ' Signal COGs (bit per COG)
```

## Your Turn

### Your Turn

#### Exercise 1: Event-Driven Serial

Rewrite a serial receive loop to use events instead of polling:

1. Configure SE1 for UART RX Smart Pin ready
2. Use WAITSE1 instead of **TESTP** loop
3. Measure the cycle count difference

```
1 ' Your code here:
2 ' Hint: setse1 #001<<6 + RX_PIN
```

### Your Turn

#### Exercise 2: Dual Event Monitor

Create a loop that monitors both a button (pin edge event) and a timer (periodic event):

1. SE1 = button press (rising edge)
2. CT1 = 1 second heartbeat
3. On button: toggle LED
4. On timer: print timestamp

```
1 ' Your code here:
2 ' Use POLL for both, handle whichever fires
```

## What We've Learned

The event toolkit you now command:

- Four selectable event channels (SE1-SE4) per COG
- Three timer events (CT1-CT3) for precise scheduling
- **WAIT** (sleep) vs **POLL** (check-and-continue) tradeoffs
- The ATN inter-COG signaling system
- Common patterns: timeout-with-fallback, debounce, periodic sampling

## Coming Up Next

Chapter 16 brings the whole journey together — orchestrating eight COGs in parallel harmony to build complete systems. It's where the P2 philosophy really shines.

**Have Fun!** And remember — every spin loop you replace with **WAITSE** is CPU cycles you've handed back to your design. Be generous with events!

# Chapter 16: Multi-COG Orchestration

*Bringing it all together in parallel harmony*

## The Hook: A Complete System in 8 COGs

Watch this system architecture come alive:

```
1 ' Main orchestrator (COG 0)
2 main_orchestrator
3     ' Launch the orchestra (SETQ sets PTRA for new COG)
4     SETQ    @sensor_params
5     COGINIT #1, @sensor_cog
6     SETQ    @motor_params
7     COGINIT #2, @motor_cog
8     SETQ    @comms_params
9     COGINIT #3, @comms_cog
10    SETQ    @display_params
11    COGINIT #4, @display_cog
12    SETQ    @safety_params
13    COGINIT #5, @safety_cog
14    SETQ    @logger_params
15    COGINIT #6, @logger_cog
16    SETQ    @debug_params
17    COGINIT #7, @debug_cog
18
19    ' Now coordinate them all
20 orchestrate
21     RDLONG  sensor_data, ##SENSOR_MAILBOX wz
22     if_nz CALL  #process_sensor_data
23
24     RDLONG  command, ##COMMAND_MAILBOX wz
25     if_nz CALL  #execute_command
26
27     CALL    #update_system_state
28     WRLONG  state, ##STATE_MAILBOX
29
30     JMP     #orchestrate
```

Eight independent processors, each with a specific job, all working in perfect coordination. This is the true power of P2!

## Communication Patterns

Eight processors running in parallel sounds wonderful — until you realize they need to talk to each other. Let's meet the three patterns you'll use 95% of the time.

### The Mailbox Pattern

The simplest and most common — a single hub **LONG** that one COG writes and another reads:

```

1  ' Producer COG
2  producer
3      ' Generate data
4      CALL    #calculate_result
5      WRLONG  result, ##MAILBOX_ADDR
6
7  ' Consumer COG
8  consumer
9      RDLONG  data, ##MAILBOX_ADDR wz
10     if_z JMP    #consumer          ' Wait for data
11     WRLONG  #0, ##MAILBOX_ADDR    ' Clear mailbox
12     CALL    #process_data

```

### The Ring Buffer Pattern

For streaming data between COGs:

```

1  ' Writer COG
2  writer_cog
3      RDLONG  wr_ptr, ##WRITE_PTR
4      WRLONG  data, wr_ptr
5      ADD     wr_ptr, #4
6      AND     wr_ptr, ##BUFFER_MASK ' Wrap around
7      WRLONG  wr_ptr, ##WRITE_PTR
8
9  ' Reader COG
10 reader_cog
11     RDLONG  rd_ptr, ##READ_PTR

```

*continues on next page →*

*↔ continued from previous page*

```

12         RDLONG  wr_ptr, ##WRITE_PTR
13         CMP     rd_ptr, wr_ptr wz
14     if_z JMP     #reader_cog          ' Buffer empty
15
16         RDLONG  data, rd_ptr
17         ADD     rd_ptr, #4
18         AND     rd_ptr, ##BUFFER_MASK
19         WRLONG  rd_ptr, ##READ_PTR

```

## The Command Queue Pattern

For sending commands between COGs:

```

1  ' Command structure in hub
2  ' +0: Command ID
3  ' +4: Parameter 1
4  ' +8: Parameter 2
5  ' +12: Result/Status
6
7  ' Commander COG
8  send_command
9         WRLONG  cmd_id, ##CMD_BUFFER+0
10        WRLONG  param1, ##CMD_BUFFER+4
11        WRLONG  param2, ##CMD_BUFFER+8
12        WRLONG  ##$FFFF, ##CMD_BUFFER+12  ' Mark as pending
13
14  wait_complete
15        RDLONG  status, ##CMD_BUFFER+12
16        CMP     status, ##$FFFF wz
17        if_z JMP     #wait_complete
18
19  ' Worker COG
20  process_commands
21        RDLONG  status, ##CMD_BUFFER+12
22        CMP     status, ##$FFFF wz
23        if_nz JMP     #process_commands    ' No command
24
25        RDLONG  cmd_id, ##CMD_BUFFER+0
26        RDLONG  param1, ##CMD_BUFFER+4
27        RDLONG  param2, ##CMD_BUFFER+8

```

*continues on next page →*

```

28
29     CALL    #execute_command
30     WRLONG  result, ##CMD_BUFFER+12    ' Signal complete

```

## Synchronization Techniques

Sometimes communication isn't enough — you need two or more COGs to *agree* on what happens when. That's where synchronization comes in.

### Using Locks

When multiple COGs need atomic access to the same piece of data, P2 gives you 16 hardware locks. They're tiny and they're fast:

```

1  ' Atomic increment using lock
2  atomic_increment
3      LOCKTRY #COUNTER_LOCK wc      ' C=1 if we got the lock
4  if_nc JMP    #atomic_increment    ' Retry until we get it
5
6      RDLONG  value, ##COUNTER
7      ADD     value, #1
8      WRLONG  value, ##COUNTER
9
10     LOCKREL #COUNTER_LOCK

```

### Event Synchronization

COGs waiting for specific events:

```

1  ' COG 1: Signal event
2      WRLONG  ##EVENT_FLAG, ##EVENT_ADDR
3
4  ' COG 2: Wait for event
5  wait_event
6      RDLONG  flag, ##EVENT_ADDR wz
7  if_z JMP    #wait_event
8      WRLONG  #0, ##EVENT_ADDR      ' Clear event

```

## Real-World Example: Robot Controller

Let's build a complete robot control system:

```

1  ' COG 0: Main Controller
2  main_controller
3      CALL    #init_system
4
5  main_loop
6      ' Read sensor hub
7      RDLONG  distance, ##DISTANCE_SENSOR wz
8      if_z JMP    #too_close
9
10     ' Check for commands
11     RDLONG  cmd, ##SERIAL_COMMAND wz
12     if_nz CALL  #process_command
13
14     ' Update motor speeds
15     CALL    #calculate_motion
16     WRLONG  left_speed, ##LEFT_MOTOR
17     WRLONG  right_speed, ##RIGHT_MOTOR
18
19     JMP     #main_loop
20
21  ' COG 1: Ultrasonic Sensor
22  sensor_cog
23     ' Trigger ultrasonic pulse
24     DRVH    #TRIGGER_PIN
25     WAITX   ##1000
26     DRVL    #TRIGGER_PIN
27
28     ' Measure echo time - wait for rising edge
29  .wait_hi
30     TESTP   #ECHO_PIN wz
31     if_nz JMP    #.wait_hi
32     GETCT   start_time
33  .wait_lo                                     ' Wait for falling edge
34     TESTP   #ECHO_PIN wz
35     if_z   JMP    #.wait_lo
36     GETCT   end_time
37

```

*continues on next page →*

```
38     ' Calculate distance
39     SUB     end_time, start_time
40     ' Convert to distance...
41     WRLONG distance, ##DISTANCE_SENSOR
42
43     WAITX  ##10_000_000      ' 50ms at 200MHz
44     JMP     #sensor_cog
45
46 ' COG 2: Left Motor Driver
47 left_motor_cog
48     RDLONG speed, ##LEFT_MOTOR wz
49     if_z JMP     #left_motor_cog      ' No speed set
50
51     ' Generate motor control signals
52     ' ... PWM generation code
53     JMP     #left_motor_cog
54
55 ' COG 3: Right Motor Driver
56 ' (Similar to left motor)
57
58 ' COG 4: Serial Communications
59 serial_cog
60     ' Check for incoming commands
61     TESTP  #RX_PIN wc
62     if_nc JMP     #serial_cog
63
64     CALL  #receive_byte
65     ' Build command...
66     WRLONG command, ##SERIAL_COMMAND
67     JMP     #serial_cog
68
69 ' COG 5: LED Status Display
70 status_cog
71     RDLONG system_state, ##STATE_MAILBOX
72
73     ' Display state on LEDs
74     CMP     system_state, #STATE_RUNNING wz
75     if_z DRVH  #GREEN_LED
76     if_nz DRVL #GREEN_LED
77
78     CMP     system_state, #STATE_ERROR wz
```

```

79     if_z DRVH     #RED_LED
80     if_nz DRVL   #RED_LED
81
82         WAITX    ##10_000_000
83         JMP      #status_cog
84
85 ' COG 6: Safety Monitor
86 safety_cog
87     ' Monitor critical systems
88     RDLONG  battery, ##BATTERY_VOLTAGE
89     CMP     battery, ##MIN_VOLTAGE wcz
90     if_b  WRLONG  ##STATE_SHUTDOWN, ##STATE_MAILBOX
91
92     ' Check temperature
93     RDLONG  temp, ##TEMPERATURE
94     CMP     temp, ##MAX_TEMP wcz
95     if_a  WRLONG  ##STATE_OVERHEAT, ##STATE_MAILBOX
96
97         JMP      #safety_cog
98
99 ' COG 7: Debug Output
100 debug_cog
101     ' Output system state for debugging
102     ' ... debug code

```

Eight COGs, each doing one job perfectly, creating a responsive, reliable robot!

## Your Turn: Multi-COG Project

### Your Turn

#### Exercise: Traffic Light Controller

Requirements:

- COG 0: Main sequencer
- COG 1: North-South lights
- COG 2: East-West lights
  
- COG 3: Pedestrian button watcher
- COG 4: Timer/scheduler

Starting structure:

```

1      ORG      0
2  ' COG 0: Main sequencer
3      ' Launch other COGs
4      ' Coordinate light changes
5      ' Handle pedestrian requests
6
7  ' Your implementation here

```

Goal: Working traffic light with pedestrian crossing Hint: Use mailboxes for COG communication Success Check: Lights change correctly, pedestrian button works

### The Medicine Cabinet

Multi-COG systems overwhelming? Start simple:

**Start with just 2 COGs:**

```

1  ' Main + Helper pattern
2  main  SETQ    @params          ' PTRA for new COG
3      COGINIT #1, @helper
4      ' Main work
5
6  helper ' Support work

```

**Use simple mailboxes:**

```

1  ' Fixed hub addresses for communication
2  MAILBOX_1 = $1000
3  MAILBOX_2 = $1004

```

**Debug one COG at a time: TEST** each COG in isolation before combining!

## Design Principles for Multi-COG Systems

The hardware gives you eight processors. Whether your *design* survives the journey is up to you. A few rules of thumb we've learned the hard way:

1. **Single Responsibility:** Each COG does ONE thing well
2. **Loose Coupling:** COGs communicate through hub, not direct dependencies

3. **Clear Ownership:** Each piece of data has one writer
4. **Predictable Timing:** Real-time tasks get dedicated COGs
5. **Graceful Degradation:** System continues if one COG fails

## Common Multi-COG Gotchas

Before you pull your hair out wondering why the eight-COG dream turned into a debugging nightmare, skim these:

1. **Race conditions** - Use locks for shared write access
2. **Deadlocks** - Avoid circular dependencies
3. **Starvation** - Ensure all COGs get resources
4. **Communication overhead** - Don't over-communicate
5. **Debugging complexity** - Use LED indicators for each COG

## What We've Learned

You've mastered multi-COG orchestration:

- Communication patterns (mailbox, ring buffer, queue)
- Synchronization techniques
- Real-world system architecture
- Design principles
- Common pitfalls and solutions

This is it - you now understand the full power of the Propeller 2!

## Your Journey Continues

You've completed this manual, but your P2 journey has just begun:

1. **Build something amazing** - Put your knowledge to work
2. **Share with the community** - Your projects inspire others
3. **Explore other manuals** - Smart Pins, Video, I/O await
4. **Push boundaries** - P2 can do things we haven't imagined yet

## Chapter Summary

### Chapter Summary

**Congratulations!** You've mastered multi-COG orchestration!

You now understand:

- How to coordinate 8 parallel processors
- Communication patterns between COGs
- Synchronization techniques
- Real-world system design

**You did it!** You're now fluent in PASM2 and ready to build incredible parallel systems!

### Have Fun!

Remember what you've learned:

- Eight COGs working together are more powerful than any interrupt-driven system
- Parallel processing isn't harder, it's different
- The P2 way is about determinism and elegance
- Every complex system is just simple parts working together

Now go forth and create something amazing with your Propeller 2!

## Epilogue: The Journey Forward

Well, here we are at the end... or should I say, at the beginning?

You've traveled from blinking your first LED to orchestrating eight parallel processors. You've mastered CORDIC mathematics, tamed the FIFO, and learned why interrupts are usually the wrong answer. That's quite a journey!

But here's the secret: everything you've learned is just the foundation. The P2 community continues to discover new techniques, new optimizations, new ways to use this remarkable chip. Every project pushes the boundaries a little further.

### What Makes You Different Now

You're not just another embedded programmer anymore. You think in parallel. You see solutions that others miss. When someone says "that's impossible in real-time," you know better - you just dedicate a COG to it.

## The Community Awaits

The Parallax forums are filled with fellow travelers on this journey. Share your projects. Ask questions. Help newcomers. The community that inspired this manual continues to grow because people like you contribute back.

## One Last Story

I remember my first P2 project. I was trying to control 16 servos with perfect timing while reading sensors and communicating over serial. On my previous microcontroller, it was a nightmare of interrupts and jitter.

On the P2? Three COGs. Clean, simple, perfect timing. That's when I truly understood - this isn't just a different processor, it's a different philosophy of computing.

## Your Challenge

Build something that wouldn't be possible without parallel processing. Something that would be a nightmare of interrupts on other processors. Then share it with the world.

Show them what eight COGs can do.

Show them the Propeller way.

*“The best way to predict the future is to invent it.”*

— Alan Kay

And with your Propeller 2, you have everything you need to invent amazing futures.

## Have Fun!

— The ghosts of deSilva, the P2 community, and one very enthusiastic AI

*P.S. - Don't forget to blink an LED once in a while, just for old times' sake. It's still magical, even after all you've learned.*

THE END

(But really, just the beginning...)

## Further Reading

This teaching manual focuses on concepts, patterns, and building your understanding. For complete technical specifications, refer to these companion documents:

**Propeller 2 Assembly Language (PASM2) Manual** Complete PASM2 instruction details including syntax, timing, and flag effects for all 300+ instructions. Quick lookup reference for day-to-day development.

**Parallax Propeller 2 Documentation** (*v35, Rev B/C silicon, 2021-05-18*) Official silicon documentation from Parallax covering hardware specifications, electrical characteristics, and detailed register maps.

# Appendix A: Platform Comparison

*How P2 compares to other microcontrollers*

If you're coming from another embedded platform, this appendix shows how the P2's approach differs and when those differences matter.

## The Landscape

The embedded world is dominated by a handful of architectures:

Platform	Architecture	Typical Cores	Peripherals	Timing
<b>STM32</b>	ARM Cortex-M	1-2	Fixed location	Cache-dependent
<b>ESP32</b>	Xtensa/RISC-V	2	Fixed location	FreeRTOS scheduled
<b>Arduino/AVR</b>	AVR	1	Fixed location	Deterministic but slow
<b>PIC32</b>	MIPS	1	Fixed location	Interrupt-driven
<b>P2 Propeller</b>	Custom	<b>8</b>	<b>Any pin</b>	<b>Deterministic</b>

## What Makes P2 Different

### Eight Real Processors, Not Time-Slicing

On ARM, ESP32, or PIC, you typically have 1-2 cores that share time between tasks using interrupts or an RTOS. The P2 gives you eight complete, identical processors that run truly in parallel.

**Traditional approach:**

```

1 ' Everyone fights for the same CPU
2 ISR(TIMER1_vect) { motor_control(); } ' Might delay...
3 ISR(UART_RX_vect) { serial_handler(); } ' ...this
4 main() { while(1) { sensor_loop(); } } ' Hope we get time

```

## P2 approach:

```
1 ' Each task owns its own processor
2 COG0: COGINIT(1, @motor_control)   ' Coordinator launches workers
3 COG1: motor_control()              ' Dedicated - always on time
4 COG2: serial_handler()             ' Dedicated - never misses byte
5 COG3: sensor_loop()               ' Dedicated - consistent sample
6 COG4-7: ready for more
```

No interrupt priority juggling. No RTOS configuration. Each task owns its processor.

## Smart Pins: Peripherals on Every Pin

Traditional MCUs have fixed peripheral assignments: UART1 is on PA9/PA10, SPI1 is on PB3/PB4/PB5, and if you need those pins for something else, you're stuck rerouting your PCB.

On P2, every pin contains a programmable state machine. Any pin can become a UART, SPI, PWM, ADC, quadrature decoder, or 27 other modes. The peripheral comes to your pin, not the other way around.

## Deterministic Timing

ARM MCUs with cache have unpredictable timing. A memory read might take 1 cycle (cache hit) or 50+ cycles (cache miss). Even instruction timing varies—ARM instructions take 1-3+ cycles depending on the operation. This makes cycle-accurate timing extremely difficult.

P2 takes a different approach: nearly all instructions execute in exactly **2 clock cycles**. Want to know how long a code sequence takes? Count the instructions and multiply by 2. Hub memory uses round-robin access that gives every COG predictable, guaranteed access slots. Your timing loops work identically every time—no cache luck required.

## Coming From ARM/STM32

You're used to configuring HAL structures, writing interrupt handlers, and managing DMA. Here's how P2 solves those problems:

Instead of...	On P2...	The Benefit
<code>HAL_UART_Transmit()</code>	Configure Smart Pin once, then <b>WYPIN</b> bytes	Pin handles all timing autonomously
<code>HAL_TIM_PWM_Start()</code>	Configure Smart Pin once, update with <b>WYPIN</b>	Pin runs independently—your COG is free
NVIC priority configuration	Nothing needed	All COGs equal, no priority inversion ever
<code>HAL_DMA_Start()</code>	Use built-in FIFO/Streamer	Simpler API, integrated into each COG
<code>arm_sin_f32()</code> library	<b>QROTATE</b> instruction	Hardware trig in exactly 55 clocks
FreeRTOS <code>xTaskCreate()</code>	<b>COGINIT</b>	True parallel execution, not scheduled

**The result:** Deterministic timing, zero interrupt conflicts, and I/O configuration that just works.

## Coming From ESP32

You're used to WiFi/Bluetooth convenience and FreeRTOS abstractions. P2 takes a different approach:

ESP32 Way	P2 Way	The Benefit
Built-in WiFi/BT	Add WizNet or ESP module	You choose your connectivity—or skip it entirely
<code>xTaskCreate()</code>	<b>COGINIT</b>	Not scheduled—truly parallel, guaranteed timing
GPIO matrix routing	Smart Pins	32 modes per pin, far more capability
FreeRTOS timing	Deterministic hub	Cycle-accurate timing guaranteed
Arduino framework	Spin2/PASM2	Deeper control, deeper understanding

**The result:** 8 real cores running simultaneously, timing you can count on, I/O flexibility that eliminates peripheral conflicts.

## Coming From Arduino/AVR

You'll find P2 familiar but dramatically more powerful:

Arduino Way	P2 Way	The Upgrade
<code>digitalWrite()</code>	<b>DRVH/DRVL</b> or Smart Pins	Similar syntax, vastly more capability
<code>delay()</code> blocks everything	<b>WAITX</b> or dedicated COG	Timing without blocking other tasks
One thing at a time	8 things truly parallel	Real concurrency, not fake multitasking
8-bit math limits	32-bit + hardware CORDIC	No more overflow worries, hardware trig
Libraries for everything	Growing ecosystem + OBEX	More control, deeper understanding

**The result:** Graduate from 8-bit limitations to 8 parallel 32-bit processors with hardware math and Smart Pins on every I/O.

## When P2 Is the Right Choice

P2 excels when you need:

- **Multiple real-time tasks** running simultaneously without conflicts
- **Precise timing** that cache misses and interrupts can't disrupt
- **Video or audio generation** requiring cycle-accurate output
- **Flexible I/O** where any pin can become any peripheral
- **Hardware math** for motor control, signal processing, or robotics
- **Multiple motor/servo control** with dedicated COGs per channel
- **Protocol implementation** where Smart Pins handle timing autonomously

## Platform Trade-offs

Every platform makes trade-offs. P2 optimizes for **determinism, parallelism, and flexibility** rather than:

If you need...	P2's answer
Built-in WiFi/Bluetooth	Add WizNet or ESP module—you choose connectivity
Massive library ecosystem	Growing OBEX + helpful community
Ultra-low-power sleep	External modules or different platform
Lowest unit cost at 100K+ volumes	P2 targets flexibility over commodity pricing

**The honest reality:** If your project is “connect to WiFi and display data,” an ESP32 does that with less effort. But if your ESP32 project is fighting timing jitter, missing deadlines, or running out of peripheral pins—that’s exactly what P2 solves.

## Community Resources

While P2’s ecosystem is smaller than ARM or Arduino, it’s active and welcoming:

**Parallax Forums** - The heart of the P2 community. Chip Gracey (P2’s designer) participates actively, answering questions and discussing design decisions. You’ll find help from experienced developers who’ve solved problems you haven’t encountered yet.

**P2 Object Exchange (OBEX)** - A library of reusable Spin2 and PASM2 objects covering drivers, protocols, display interfaces, and more. Before writing something from scratch, check OBEX—someone may have already done the work.

**Community Support** - Unlike large platforms where your question disappears in a sea of posts, the P2 community is small enough that questions get noticed and answered. Many community members have decades of Propeller experience.

Coming from Arduino’s library-for-everything culture, you’ll write more code yourself—but you’ll understand it deeply, and help is always available when you get stuck.

## The P2 Hardware Ecosystem

P2 isn’t just a chip - it’s a platform with expansion options:

### Video & Audio:

- A/V Breakout Board: VGA, RCA, 80mW stereo, microphone input
- Digital Video Out: HDMI-type with differential signaling
- Built-in 8-bit DAC per pin (16-bit with dithering)

### Connectivity:

- USB Host Board: Two USB-A ports
- USB Device Board: HID or CDC modes
- Serial Host/Device: RS-232 interfaces
- WizNet/ESP modules: Ethernet or WiFi

### Development:

- P2 Eval Board: Complete development environment
- Edge Modules: 4MB or 32MB flash for embedding
- Breakout Boards: All 64 pins accessible

You add what you need - no paying for peripherals you won't use.

## Summary

P2 represents a fundamentally different approach to embedded computing—one that eliminates entire categories of problems:

- **Eight processors** means your motor control never delays your serial handler
- **64 Smart Pins** means peripheral conflicts become impossible
- **Deterministic timing** means your code works the same way every time
- **Hardware CORDIC** means real-time math without floating-point libraries

Engineers who've fought interrupt priority inversions, missed timing deadlines, and PCB rework due to peripheral conflicts find P2 refreshing. You spend your time solving your actual problem, not fighting your MCU.

**Welcome to the P2 community.** You've got 8 processors, 64 Smart Pins, and a community that's been building amazing things since the original Propeller. Time to see what you can build.

# Index

## A

- ADC operations: Ch14
- **ADD** instruction: Ch3, Ch5
- ADDCT1/2/3: Ch15
- Address modes: Ch3
- **ALTD**/**ALTS**: Ch3
- Architecture: Ch2
- Assembly basics: Ch3

## B

- Bit manipulation: Ch3
- Block transfers: Ch4, Ch9
- Booleans: Ch3

## C

- C flag: Ch6
- **CALL**/**RET**: Ch3
- Clock timing: Ch2
- **CMP** instruction: Ch6
- COG anatomy: Ch2
- COG communication: Ch2, Ch16
- Conditional execution: Ch3, Ch6
- CORDIC: Ch7
- Counters: Ch2

## D

- DAC operations: Ch14
- Debugging: Ch12
- Division: Ch5
- **DRVH**/**DRVL**: Ch1

## E

- Egg beater: Ch2, Ch4
- Event system: Ch15

**F**

- FIFO: Ch4, Ch9
- Flags: Ch6
- Flow control: Ch3

**G**

- **GETQX**/GETQY: Ch5, Ch7

**H**

- Hardware multiply: Ch5
- Hub execution: Ch10
- Hub memory: Ch2, Ch4

**I**

- IN flag: Ch14, Ch15
- Immediate values: Ch3
- Instruction format: Ch3
- Interrupts: Ch11

**J**

- **JMP** instruction: Ch3

**L**

- LED control: Ch1
- Locks: Ch2, Ch16
- Logic operations: Ch3
- LUT memory: Ch13
- LUT sharing: Ch13

**M**

- Mailboxes: Ch2, Ch16
- Mathematics: Ch5
- Memory access: Ch4
- **MOV** instruction: Ch3
- **MUL**/MULS: Ch5

- Multi-COG: Ch16

## O

- Optimization: Ch12

## P

- Parallel processing: Ch2
- Pipeline: Ch7, Ch12
- Pins, Smart: Ch8, Ch14
- Platform comparison: Appendix A
- POLLSE1-4: Ch15
- PTR A/PTR B: Ch3, Ch4
- PWM: Ch8, Ch14

## Q

- Q flag: Ch7
- **QDIV**: Ch5
- **QROTATE**: Ch7

## R

- **RDBYTE**/**RDWORD**/**RDLONG**: Ch4
- **RDLUT**/**WRLUT**: Ch13
- **RDPIN**: Ch14, Ch15
- **REP** instruction: Ch3
- Rotation: Ch3, Ch7

## S

- SETSE1-4: Ch15
- Shift operations: Ch3
- **SKIP** instruction: Ch3, Ch6
- Smart Pins: Ch8, Ch14
- Streamer: Ch9, Ch13

**T**

- Timer: Ch2
- Timing: Ch2, Ch12
- Trigonometry: Ch7

**U**

- UART: Ch8, Ch14

**V****W**

- WAITSE1-4: Ch15
- WAITCT1-3: Ch15
- **WAITX**: Ch1
- **WRBYTE**/**WRWORD**/**WRLONG**: Ch4
- **WRPIN**/**WXPIN**/**WYPIN**: Ch14

**Z**

- Z flag: Ch6