

P2 Documentation and Code Project

You are viewing draft content
created with the use of Claude.AI



P2 Assembly Programming

A Human-Centered Approach to Parallel Processing

June 2026

Version 3.0.1

Tutorial Philosophy

Learn by doing, celebrate progress, have fun!

Code Block Colors:

- **Green** – PASM2
- **Blue** – Spin2
- **Purple** – CORDIC
- **Teal** – Multi-Cog
- **Red** – Antipattern

Teaching Elements:

- **Sidetracks** – deeper dives
- **Medicine Cabinet** – simpler alternatives
- **Your Turn** – hands-on exercises
- **Interludes** – stories & context

Contents

Copyright and License	8
Dedication	9
Acknowledgments	10
Preface: Welcome to the Journey	12
Chapter 1: Your First Spin	15
Why P2?	15
The Hook: Making Light	15
What’s Really Happening	16
Sidetrack: The Clock Preamble	17
Let’s Make It Better	18
Understanding Cogs	18
Your Turn: Experiments	19
Sidetrack: Why Start at Address 0?	20
Common Gotchas	21
What We’ve Learned	21
Coming Up Next	21
Chapter 2: Architecture Safari	23
The Propeller Philosophy	23
Cog Anatomy 101	24
Meet the Hub: The Meeting Place	25
Let’s See Cogs in Action	26
Cog Communication: How They Talk	26
The Timer: Everyone Gets One	27
Why No Interrupts? (Usually)	28
Real-World Example: Parallel Sensors	28
Sidetrack: The Philosophy of Parallel	30
Common Gotchas	31
What We’ve Learned	31
Your Turn: Experiments	31
Coming Up Next	32
Chapter 3: Speaking PASM2	33

The Hook: One Instruction, Many Powers	33
Instruction Anatomy 101	33
The Basic Vocabulary	34
Flow Control: Jump!	36
Labels: Naming Your Places	38
Data in DAT Blocks: Your Program's Pantry	41
The Flags: C and Z (and Q!)	47
Special Instructions That Will Blow Your Mind	47
Real-World Example: Fast Memory Copy	48
Common Gotchas	49
Your Turn: Experiments	50
Sidetrack: Why PASM2 Is Different	51
What We've Learned	51
Coming Up Next	52
Chapter 4: The Hub Connection	53
The Hook: Instant Communication	53
Reading from Hub	53
Writing to Hub	53
The FIFO Pipeline	54
Real-World Example: Video Buffer	54
Your Turn: Experiments	54
Common Gotchas	55
What We've Learned	55
Coming Up Next	56
Chapter 5: Mathematics Unleashed	57
The Hook: Hardware Multiply	57
The Multiplication Revolution	57
Division Without Tears	57
64-Bit Operations	58
Real-World Example: Fixed-Point Math	58
Your Turn: Experiments	59
Common Gotchas	59
What We've Learned	59
Coming Up Next	60
Chapter 6: Flags and Decisions	61
The Hook: Any Instruction Can Be Conditional	61
The C and Z Flags	61
Complex Conditions	61
Skip Patterns - Conditional Blocks	62
Your Turn: Experiments	62

Common Gotchas	63
What We've Learned	63
Coming Up Next	63
Chapter 7: CORDIC Magic	64
The Hook: Rotate a Point in 3 Lines	64
What Just Happened?	64
The CORDIC Pipeline - Your Mathematical Assembly Line	65
Core CORDIC Operations	65
Real-World Example: Spinning a Sprite	66
Your Turn: CORDIC Experiments	67
Advanced CORDIC: The Pipeline Dance	69
CORDIC for Graphics	70
CORDIC for Audio	71
Common CORDIC Gotchas	71
What About QLOG, QEXP?	72
What We've Learned	72
Coming Up Next	73
Chapter 8: Basic I/O	74
The Hook: One Pin, Three Instructions, Infinite Possibilities	74
Understanding P2 Pins	74
Digital Output: Making Things Happen	75
Digital Input: Reading the World	76
Pin Timing: When Things Happen	76
Real-World Example: Button Debouncing	77
Bit-Banged Serial (The Basics)	77
Your Turn: I/O Experiments	78
Advanced Pin Control	79
Common I/O Gotchas	80
Timing Is Everything	80
Real-World Example: Servo Control	80
The Power of Simple	81
What We've Learned	81
A Note About Smart Pins	81
Coming Up Next	81
Chapter 9: Streaming Data	83
The Hook: 4KB in 4 Instructions	83
Block Transfers: The Power Move	83
The FIFO: Your Streaming Pipeline	83
Writing Through the FIFO	84
Real-World Example: Screen Buffer Clear	84

Streaming with the Streamer	85
FIFO and Cog Execution	85
Advanced Streaming Techniques	86
Your Turn: Streaming Experiments	87
Common Streaming Gotchas	88
Performance Numbers	88
Real-World Example: Audio Buffer	89
What We've Learned	89
Coming Up Next	90
Chapter 10: Hub Execution	91
The Hook: Unlimited Code Space	91
Cog vs Hub Execution: The Trade-offs	91
How Hub Execution Works	92
Real-World Example: Menu System	92
The Hub Execution FIFO	93
Mixing Cog and Hub Code	93
Your Turn: Hub Execution Experiments	94
Advanced Hub Execution	95
Common Hub Execution Gotchas	96
Real-World Example: Command Parser	96
When to Use Hub Execution	97
What We've Learned	97
Coming Up Next	98
Chapter 11: Why No Interrupts?	99
The Hook: Interrupts Without Interrupts	99
The Interrupt Problem	99
The Propeller Solution	100
Real-World Example: Perfect Servo Control	100
"But P2 HAS Interrupts!"	101
The Event System: Better Than Interrupts	102
Interrupt Horror Stories	103
Your Turn: Cog vs Interrupt Challenge	104
The Philosophy Deep Dive	104
When Interrupts Actually Make Sense	105
Common "But What About..." Questions	105
What We've Learned	105
Coming Up Next	106
Chapter 12: Optimization Mastery	107
The Hook: Double Your Speed with One Change	107
Understanding the Pipeline	107

Instruction Timing Basics	107
REP: The Speed Loop	108
SKIP: Conditional Execution on Steroids	108
Hub Access Optimization	109
The FIFO Fast Path	109
Parallel Operations	110
Real-World Example: Fast Memory Copy	110
Your Turn: Optimization Challenges	111
Advanced Techniques	112
Common Optimization Gotchas	113
Profiling and Measurement	113
What We've Learned	114
Coming Up Next	114
Chapter 13: LUT Memory - Your Private Lookup Table	115
The Hook: A Lookup Table in 3 Cycles	115
Why Another Memory?	115
Reading and Writing the LUT	116
LUT Sharing Between Cogs	117
Practical Examples	118
LUT with the Streamer	119
Common Gotchas	120
Medicine Cabinet	120
Your Turn	121
What We've Learned	122
Coming Up Next	122
Chapter 14: Smart Pins Orientation	123
The Hook: A UART in 4 Lines	123
What Are Smart Pins, Really?	123
The Universal Smart Pin Pattern	123
The Core Instructions	124
Understanding the IN Flag	125
Common Smart Pin Modes	126
Configuration Values Demystified	128
Common Gotchas	128
Medicine Cabinet	129
Your Turn	130
What We've Learned	130
Coming Up Next	131
Chapter 15: Event-Driven Programming	132
The Hook: No More Polling Loops	132

Why Events Matter	132
The Four Selectable Events	133
Configuring an Event	133
Timer Events	136
Waiting vs Polling	136
Multiple Events	137
Practical Examples	137
ATN - Inter-Cog Events	139
Common Gotchas	139
Medicine Cabinet	140
Your Turn	141
What We've Learned	142
Coming Up Next	142
Chapter 16: Multi-Cog Orchestration	143
The Hook: A Complete System in 8 Cogs	143
Communication Patterns	143
Synchronization Techniques	145
Real-World Example: Robot Controller	146
Your Turn: Multi-Cog Project	148
Design Principles for Multi-Cog Systems	149
Common Multi-Cog Gotchas	150
What We've Learned	150
Your Journey Continues	150
Chapter Summary	150
Epilogue: The Journey Forward	151
Further Reading	152
Appendix A: Platform Comparison	153
The Landscape	153
What Makes P2 Different	153
Coming From ARM/STM32	155
Coming From ESP32	155
Coming From Arduino/AVR	156
When P2 Is the Right Choice	156
Platform Trade-offs	156
Community Resources	157
The P2 Hardware Ecosystem	157
Summary	158
Index	159

Copyright and License

Copyright © 2025-2026 Iron Sheep Productions, LLC and Parallax Inc.

This work is licensed under the Creative Commons Attribution–NonCommercial–NoDerivatives 4.0 International License (CC BY-NC-ND 4.0).

You are free to:

- **Share** — copy and redistribute the material in any medium or format

Under the following terms:

- **Attribution** — You must give appropriate credit, provide a link to the license, and indicate if changes were made (for example, formatting or excerpting).
- **NonCommercial** — You may not use the material for commercial purposes.
- **NoDerivatives** — If you remix, transform, translate, or build upon the material, you may not distribute the modified material.

Commercial use: For uses that may be commercial (including paid courses, kits, or redistribution with products), please contact Iron Sheep Productions, LLC and Parallax Inc. (info@ironsheep.biz) for separate permission.

To view the full license, visit: <https://creativecommons.org/licenses/by-nc-nd/4.0/>

Trademarks

Propeller, Propeller 2, P2, Spin, and the Parallax logo are trademarks of Parallax Inc.

Disclaimer

The information in this manual is subject to change without notice. While every effort has been made to ensure accuracy, the authors and publishers assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

Dedication

To deSilva — *Whose legendary P1 assembly tutorial taught a generation of programmers that assembly language could be approachable, enjoyable, and even fun. Your unique voice—combining technical precision with human warmth—showed us that great documentation teaches not just the mind, but speaks to the spirit of discovery.*

To the Propeller Community — *Who have spent countless hours exploring, documenting, and sharing their knowledge. From the early P1 pioneers to today's P2 innovators, your collective wisdom makes this manual possible.*

To Future Makers — *May you find in these pages the same joy of discovery that we experienced. The Propeller 2 is more than a microcontroller—it's an invitation to think differently about computing. Welcome to the journey.*

"The best way to predict the future is to invent it." — Alan Kay

Acknowledgments

This manual stands on the shoulders of giants. We gratefully acknowledge:

Primary Contributors

deSilva - For creating the gold standard of microcontroller documentation with the P1 Assembly Tutorial. Your pedagogical approach, combining technical depth with human empathy, remains unmatched. This manual attempts to honor your legacy while adapting to the P2's capabilities.

Iron Sheep Productions LLC (Stephen M Moraco) - For extensive P2 documentation efforts, community tools, and the vision of creating an AI-optimized knowledge base. Your systematic approach to extracting and organizing P2 knowledge made this comprehensive manual possible.

Chip Gracey - Creator of the Propeller architecture. Thank you for giving us a microcontroller that thinks differently and challenges us to do the same.

Community Contributors

The Parallax Forums Community - Your questions, answers, code examples, and endless experimentation have created a living knowledge base that no single author could match.

Early P2 Adopters - Who dealt with evolving documentation, changing specifications, and still produced amazing projects that showed us what was possible.

Technical Reviewers

Special thanks to those who reviewed drafts, tested code examples, and provided invaluable feedback:

- The P2 Documentation Team at Parallax
- Community members who beta-tested examples
- Everyone who reported errors and suggested improvements

Inspiration

The MIT AI Lab - For showing us that technical documentation can have personality

Donald Knuth - For proving that programming texts can be literature

The Demoscene Community - For pushing hardware beyond its limits and inspiring us to do the same

Production Notes

This manual was created using:

- Knowledge extracted from official Parallax technical documentation and OBEX (Object Exchange) community contributions
- AI-assisted content generation trained on deSilva's writing style
- Community validation and real-world testing
- A commitment to making parallel processing accessible to everyone

"If I have seen further, it is by standing on the shoulders of giants." — Isaac Newton

Any errors, omissions, or dad jokes that fell flat are entirely the responsibility of the authors, not our distinguished contributors.

Preface: Welcome to the Journey

Well, here we are! You're about to embark on a journey into the heart of the Propeller 2, and I promise you, it's going to be quite different from what you might expect.

A Different Kind of Processor

The Propeller 2 isn't just another microcontroller. Oh no, it's something far more interesting. Imagine, if you will, eight independent processors (we call them cogs) all working together in perfect harmony, sharing a common memory space, yet each running their own programs at full speed. No interrupts fighting for attention, no complex priority schemes, just eight brains working in parallel.

And if you think this sounds terribly complicated, you're probably right... but here's the secret: it's actually simpler than traditional architectures once you understand the philosophy.

About This Manual

This manual follows in the footsteps of deSilva's legendary P1 tutorial. What does that mean? It means we're going to:

1. **Start with working code** - You'll be blinking LEDs before you know what hit you
2. **Learn by doing** - Theory is important, but nothing beats hands-on experience
3. **Have some fun** - Yes, assembly language can actually be enjoyable!
4. **Be honest about complexity** - When something is hard, we'll admit it and then show you how to handle it

Who Is This For?

Are you a complete beginner to assembly language? Welcome! We'll take good care of you.

Are you a grizzled veteran who's been writing assembly since the 6502? Welcome! The P2 will still surprise you.

Are you somewhere in between? Perfect! This is exactly where you want to be.

The only requirement is curiosity and a willingness to think a bit differently about how computers can work.

How to Use This Manual

The Fast Track — If you're impatient (and who isn't?), jump straight to Chapter 1. Get that LED blinking. Feel the satisfaction. Then come back here when you're ready for more.

The Scenic Route — Read the chapters in order. Each builds on the previous one, and I’ve hidden little gems of knowledge throughout that will make later chapters easier.

The Reference Approach — Already know what you’re looking for? The table of contents and index are your friends. The appendices contain every instruction, every smart pin mode, every CORDIC operation.

What Makes the P2 Special?

Let me count the ways:

- **8 symmetric cogs** - No master/slave relationships, all cogs are equal
- **64 smart pins** - Each pin has its own processor for I/O operations
- **CORDIC engine** - Hardware trigonometry and coordinate transformations
- **Hardware multiply/divide** - Finally! Real math at hardware speed
- **512KB of RAM** - Shared by all cogs with deterministic access timing
- **No interrupts** - Well, actually there are interrupts, but we’ll talk about why you probably don’t want them

A Note on Our Approach

The best technical documentation remembers you’re human. You’ll get frustrated. You’ll make mistakes. Your code won’t work the first time (or the second, or sometimes even the third).

That’s normal. That’s learning. And that’s why this manual provides plenty of “medicine” along the way - simpler alternatives, working examples, and the occasional moment of levity to keep your spirits up.

The deSilva Spirit

Throughout this manual, you’ll encounter the teaching spirit of deSilva. When you see phrases like:

- “Well, ...” - We’re about to correct a common assumption
- “Uff!” - We just got through something complex
- “Have Fun!” - We mean it, this stuff is actually enjoyable

These aren’t just quirks; they’re signals that we remember you’re human and we’re on this journey together.

Ready?

Take a deep breath. Pour yourself your favorite beverage. Open your development environment.

Let's make some magic happen with the Propeller 2!

"The Propeller architecture is based on the simple idea that the best way to avoid the complexity of interrupts is to have enough processors that you don't need them."

— Chip Gracey, creator of the Propeller

Turn the page, and let's blink that LED! →

Chapter 1: Your First Spin

Let's blink an LED and change your life

Why P2?

Before we dive into code, let me tell you why you're in for something different.

If you've fought with interrupt priority conflicts on an ARM, watched your timing jitter because of cache misses, or discovered that the UART you need is only available on pins you're already using... well, the P2 was designed by someone who got tired of those problems too.

Here's the P2 philosophy in a nutshell:

Instead of one processor fighting with interrupts, you get eight complete, identical processors (cogs) that run truly in parallel. Your serial handler never delays your motor control. Your sensor sampling never misses a deadline. Each task owns its own processor.

Instead of fixed peripherals, every one of the 64 pins contains its own programmable state machine. Any pin can become a UART, PWM output, quadrature encoder, ADC - whatever you need, wherever you need it.

Instead of timing that depends on cache luck, the hub memory has deterministic access. Your timing loops work the same way every time.

Instead of calling math libraries, there's a hardware CORDIC that gives you sine, cosine, and arctangent (via vector rotate and convert) with results ready in exactly 55 clocks. Every time.

Does this mean P2 is perfect for everything? Of course not. But if your projects involve multiple real-time tasks, precise timing, video or audio generation, or just running out of peripheral pins - you're in the right place.

For a full comparison to ARM, ESP32, Arduino, and PIC platforms, see Appendix A. But you probably want to blink that LED first, don't you?

The Hook: Making Light

I know you're absolutely crazy to have your first instruction executed, so let's not waste any time. Here's a complete PASM2 program that blinks an LED on pin 56 (that's the built-in LED on the P2 Eval board):

```
CON
    _clkfreq = 200_000_000      ' 200 MHz system clock
```

continues on next page →

```

DAT
' LED Blinker - Your first PASM2 program!
    ORG      0          ' Start at cog address 0

    DRVH     #56        ' Drive pin 56 high (LED on)
    WAITX    ##50_000_000 ' Wait 0.25 seconds at 200MHz
    DRVL     #56        ' Drive pin 56 low (LED off)
    WAITX    ##50_000_000 ' Wait 0.25 seconds
    JMP      #$-6       ' Jump back 6 longs (## adds hidden AUGD)

```

That's it! Five lines of code and you have a blinking LED. Load this into any cog and watch the magic happen.

What's Really Happening

Well, now that you've seen it work (you did try it, right?), let's talk about what's actually going on here.

The Instructions Decoded

ORG 0 - This tells the assembler to start placing code at cog address 0. Every cog has its own private 512 longs (2KB) of memory, and execution always starts at address 0 when a cog is loaded.

DRVH #56 - This drives pin 56 high (3.3V). The 'h' means high. The '#' means we're using an immediate value (the actual number 56) rather than the contents of register 56. One instruction, and your LED is on!

WAITX ##50_000_000 - This waits for 50 million clock cycles. At 200 MHz, that's 0.25 seconds. Notice the double '##'? That means this is a 32-bit immediate value. Single '#' only gives us 9 bits.

DRVL #56 - Drive low. LED off. You get the pattern.

JMP #\$-6 - Jump back 6 longs. The '\$' means "current address", so '\$-6' means "6 addresses back from here". Why 6? Each ## immediate generates a hidden **AUGD** instruction (it augments the WAITX destination operand), so we have 6 longs total: **DRVH, AUGD, WAITX, DRVL, AUGD, WAITX**. Infinite loop achieved!

But Wait, There's More!

"Hold on," you might say, "how does this even get into the cog?"

Ah, excellent question! In the real world, you'd typically launch this from Spin2 (the high-level language) like this:


```

CON
    _clkfreq = 200_000_000 ' 200 MHz system clock

PUB main()
    COGINIT(COGEXEC_NEW, @blink_code, 0) ' Start PASM2 in new cog
    repeat ' Keep the main cog alive

DAT
    ORG      0
blink_code
    DRVH     #56
    WAITX    ##50_000_000
    DRVL     #56
    WAITX    ##50_000_000
    JMP      #$.-6

```

ch01-first-blink.spin2

The **COGINIT** instruction loads your PASM2 code from hub memory into a fresh cog and starts it running. Meanwhile, your Spin2 code keeps running in its own cog. You now have parallel processing!

Sidetrack

The Clock Preamble

Notice the **CON** section at the top of that example? Every P2 program needs to configure its system clock:

```

CON
    _clkfreq = 200_000_000 ' 200 MHz system clock

```

This tells the P2 to run at 200 MHz using your board's crystal oscillator. Without it, the chip runs at a sluggish ~20 MHz on its internal RC oscillator—and timing-dependent code (including **DEBUG** output) won't behave as expected.

At 200 MHz with most instructions taking 2 clocks, each cog executes approximately 100 million instructions per second (100 MIPS). With 8 cogs running in parallel, that's 800 MIPS of total processing power—and that's before smart pins start handling I/O autonomously.

From here on, we'll omit this preamble from examples to keep them focused on the concept being taught. When you create your own files, always include it at the top before your **PUB** or **DAT** sections.

Let's Make It Better

The blinker works, but it's a bit rigid, isn't it? What if we want to change the blink rate? Let's use a register:

```

        ORG      0

        MOV      delay, ##50_000_000    ' Set delay to 0.25 sec
                                         ' (at 200 MHz)

.blink  DRVH     #56                     ' LED on
        WAITX    delay                  ' Wait
        DRVL     #56                     ' LED off
        WAITX    delay                  ' Wait
        JMP      #.blink                 ' Repeat forever

delay   LONG     0                       ' Storage for delay value

```

Uff! Look at that - we're using a register now! The **MOV** instruction copies our delay value into a register (which we cleverly named 'delay'). Now we can change the blink rate by modifying just one value.

A note on terminology: P2 documentation often uses "register" to refer to any long in cog RAM. Unlike ARM or x86 where registers are a small, special set (R0-R15, EAX, etc.), every cog RAM location can be used as a general-purpose register. However, the last 16 locations (496-511) have special functions: addresses 496-503 are dual-purpose (usable as RAM if interrupts aren't used), and 504-511 are special-purpose registers (PTRA, PTRB, DIRA, DIRB, OUTA, OUTB, INA, INB). When you see "register" in P2 context, think "cog RAM location."

Understanding Cogs

Here's something important: each cog is a complete processor with its own memory. When we loaded our blink program, it was copied from hub memory into cog memory. The cog then executes it independently, without any further connection to hub memory (unless we explicitly read or write to it).

Think of it like this:

- **Hub memory** is the meeting place (512KB shared by all)
- **Cog memory** is private workspace (2KB per cog)
- Loading a cog is like making a photocopy - the cog gets its own copy to run

This is why our blinker keeps running even after the Spin2 code that launched it goes into an infinite repeat loop. The cog is independent!

Your Turn: Experiments

Now for the fun part. Try these modifications:

Experiment 1: Different Patterns

Make the LED blink in a pattern: short-short-long (like SOS):

```

        ORG      0

        MOV      short, ##20_000_000    ' 0.1 second (at 200 MHz)
        MOV      long_d, ##60_000_000    ' 0.3 seconds (at 200 MHz)

.pattern DRVH     #56                    ' Short pulse 1
        WAITX    short
        DRVL     #56
        WAITX    short

        DRVH     #56                    ' Short pulse 2
        WAITX    short
        DRVL     #56
        WAITX    short

        DRVH     #56                    ' Long pulse
        WAITX    long_d
        DRVL     #56
        WAITX    long_d

        JMP      #.pattern

short    LONG     0
long_d   LONG     0

```

Experiment 2: Multiple LEDs

Blink LEDs on pins 56 and 57 alternately:

```

        ORG      0

.loop   DRVH     #56                    ' LED 56 on
        DRVL     #57                    ' LED 57 off
        WAITX    ##50_000_000          ' 0.25 sec at 200 MHz

        DRVL     #56                    ' LED 56 off
        DRVH     #57                    ' LED 57 on
        WAITX    ##50_000_000          ' 0.25 sec at 200 MHz

        JMP      #.loop

```

Experiment 3: Fading (Advanced)

This one's a bit tricky - we'll use PWM to fade the LED:

```

        ORG      0

        DIRL     #56                ' Reset the pin before configuring
        WRPIN    ##P_PWM_TRIANGLE, #56 ' Configure pin 56 for PWM
        WXPIN    ##$0100_0001, #56   ' Frame = 256 base periods
        DIRH     #56                ' Enable the pin

.fade   WYPIN    level, #56          ' Set duty cycle
        WAITX    ##100_000          ' Small delay
        ADD      level, #1           ' Increment brightness
        AND      level, #$FF         ' Wrap at 256
        JMP      #.fade

level   LONG     0

```

Don't worry if the PWM example seems complex - we'll cover smart pins in detail in Chapter 14!

The Medicine Cabinet

Feeling overwhelmed? Here's the simplified prescription:

Minimum viable blinker - Just 3 instructions:

```

.loop   DRVNOT   #56                ' Toggle pin 56
        WAITX    ##50_000_000      ' 0.25s at 200MHz
        JMP      #.loop            ' Repeat

```

The **DRVNOT** instruction toggles a pin - if it's high, make it low; if it's low, make it high. Sometimes simpler is better!

Sidetrack

Why Start at Address 0?

You might wonder why cog code always starts at address 0. It's actually quite elegant:

When a cog is started with `COGINIT`, the hardware:

1. Stops the cog (if it was running)
2. Copies 504 longs from hub to cog memory (addresses 0-503); the top 8 (504-511) are special hardware registers, not loaded
3. Starts execution at cog address 0

continues on next page →

This means every cog program starts fresh, with a clean slate. No residual state, no confusion. It's like each cog gets a fresh brain transplant every time it starts!

The last 16 longs (addresses 496-511) have special functions: 496-503 are dual-purpose (usable as RAM if interrupts not used), and 504-511 are special-purpose registers (PTRA, PTRB, DIRA, DIRB, OUTA, OUTB, INA, INB). We'll explore these later.

Common Gotchas

Before we move on, let me save you some debugging time:

1. **Forgetting the ##** - Using `WAITX #25_000_000` will NOT wait for 0.25 seconds! Single `#` only allows values up to 511.
2. **Wrong pin number** - The P2 Eval board's eight LEDs are on pins 56-63. The P2 Edge Standard Module has two LEDs on pins 56-57; the 32MB Edge Module uses 56-57 for its PSRAM and places its two LEDs on pins 38-39 instead.
3. **Clock setup required** - P2 boots at ~20MHz (internal RC oscillator). Most programs configure 200MHz with a crystal. Our examples assume 200MHz - adjust **WAITX** values if your clock differs.
4. **Cog already running** - If you `COGINIT` to a specific cog that's already running something else, it will be stopped and replaced. Use `COGEXEC_NEW` to automatically find a free cog.

What We've Learned

Let's celebrate what you've accomplished:

- ✓ Written your first PASM2 program
- ✓ Controlled hardware (LED) directly
- ✓ Used immediate values (`#` and `##`)
- ✓ Created loops with **JMP**
- ✓ Understood cog independence
- ✓ Modified code for different patterns

That's quite a lot for Chapter 1!

Coming Up Next

In Chapter 2, we'll take our "Architecture Safari" and explore:

- How 8 cogs really work together

- The hub memory system and the “egg beater”
- Why the P2 doesn’t need interrupts
- How to make cogs talk to each other

But for now, enjoy your blinking LED. You’ve just taken your first step into parallel processing!

Have Fun! And remember, every expert was once a beginner who kept their LED blinking when everyone else gave up.

Chapter 2: Architecture Safari

Eight brains are better than one

The Propeller Philosophy

Before we dive into the technical details, let's talk philosophy. Why would anyone design a micro-controller with eight processors?

The answer is beautifully simple: to avoid complexity.

“Wait,” you might say, “eight processors sounds MORE complex, not less!”

Well, consider the traditional approach:

- One processor trying to do everything
- Interrupts constantly breaking your flow
- Priority levels to juggle
- Context switching overhead
- Race conditions and timing nightmares


Now consider the Propeller way:

- Eight processors, each doing one thing well
- No interrupts needed (why interrupt when you have a dedicated processor?)
- No priorities (all cogs are equal)
- Deterministic timing (you know EXACTLY when things happen)
- True parallel processing (not time-slicing)

It's like the difference between one stressed-out juggler trying to keep eight balls in the air versus eight relaxed people each tossing one ball. Which seems simpler?

Cog Anatomy 101

Let's dissect a cog and see what makes it tick:



But here's the beautiful part: Cogs are identical. There's no "master" cog or "special" cog. Any cog can do anything any other cog can do. Democracy in silicon!

The 512-Long Limit

Each cog has exactly 512 longs (2048 bytes) of memory. The first 496 longs (\$000–\$1EF) are yours to use for code and data. The last 16 (\$1F0–\$1FF) have special roles: \$1F0–\$1F7 can serve as ordinary registers or hold the interrupt vectors and the PA/PB parameters, and \$1F8–\$1FF are the special-purpose registers (PTRA/PTRB, the I/O registers, and friends) – but not like P1; we'll get to that.

"Only 496 instructions?" you might cry. "That's tiny!"


Well, yes and no. Remember:

1. PASM2 instructions are powerful - one instruction often does what takes several in other processors
2. You have EIGHT of these cogs
3. There's hub execution mode for larger programs (Chapter 10)
4. Most real-time tasks **FIT** easily in 496 instructions

Think of it like haiku - the constraint forces elegance.

Meet the Hub: The Meeting Place


The hub is where cogs come together. It's 512KB of RAM shared by all cogs, and it's where the magic of cooperation happens.



The Egg Beater Revolution

Now here's where P2 gets clever. In P1, cogs took turns accessing the hub in a round-robin fashion. If you missed your slot, you waited for the wheel to come around again.

P2 uses what we call the “egg beater” model. Imagine eight beaters (cogs) all whipping through the same bowl (hub) simultaneously, but their paths are cleverly arranged so they never collide:



The practical result? Hub access is MUCH faster and more predictable. Instead of waiting up to 16 clocks (P1), you wait at most 8 clocks (P2), and often less if you align your accesses properly.

Let's See Cogs in Action

Here's a simple demonstration of multiple cogs working together:

```
' Multi-cog LED Pattern Demo
PUB main() | i
  repeat i from 0 to 3
    COGINIT(COGEXEC_NEW, @cog_code, 56 + i) ' Start 4 cogs
  repeat ' Main cog just watches

DAT
  ORG      0
cog_code
  MOV      pin_num, ptra          ' Pin number was passed in via PTRa

.loop     DRVNOT pin_num          ' Toggle our LED
          SHL      pin_num, #24    ' Pin number to bits 24-31
          OR       pin_num, ##10_000_000 ' Combine with delay
          WAITX    pin_num         ' Wait (varies per cog!)
          SHR      pin_num, #24    ' Restore pin number
          JMP      #.loop

pin_num LONG 0

                                     ch02-multicog-blink.spin2
```

What's happening here:

1. The main Spin2 code starts 4 cogs
2. Each cog gets a different pin number (56, 57, 58, 59)
3. Each cog blinks its LED at a slightly different rate
4. All four LEDs blink independently and simultaneously!

Cog Communication: How They Talk

Cogs are independent, but they're not isolated. They can communicate through hub memory:

Method 1: Simple Variables

```
' cog 1: Writer
    MOV     value, #42
    WRLONG  value, ##$1000 ' Write to hub address $1000

' cog 2: Reader
    RDLONG  result, ##$1000 ' Read from hub address $1000
```

Method 2: Locks (When It Matters)

When multiple cogs might write to the same location, we need locks:

```
' Get a lock
.try_lock
    LOCKTRY lock_id wc      ' Try to get lock (C=1 if we got it)
    if_nc JMP     #.try_lock ' Keep trying until we get it

    ' Critical section - we have the lock!
    RDLONG  value, ##shared_addr
    ADD     value, #1
    WRLONG  value, ##shared_addr

    LOCKREL lock_id        ' Release the lock

lock_id LONG    0          ' Lock 0-15
```

Method 3: Mailboxes (Elegant)

A mailbox is just a hub location where cogs leave messages:

```
' cog A: Leave a message
    WRLONG  message, ##mailbox

' cog B: Check for messages
.check RDLONG data, ##mailbox wz
    if_z JMP     #.check      ' Keep checking if empty
    WRLONG  #0, ##mailbox    ' Clear mailbox
    ' Process the message in 'data'
```

The Timer: Everyone Gets One

Every cog reads the same free-running 64-bit system counter (with **GETCT**), and each cog has its own CT1/CT2/CT3 compare targets to schedule timed events against it. This is incredibly useful:

```

' Method 1: Simple delay
    GETCT    start_time
    ADDCT1   start_time, ##1_000_000
    WAITCT1                      ' Wait exactly 1,000,000 clocks

' Method 2: Periodic events
    GETCT    time
.loop      ADDCT1   time, ##10_000_000
           WAITCT1                      ' Wait for next 10M clock interval
           DRVNOT   #56                  ' Toggle LED
           JMP      #.loop                ' Perfectly periodic!

```

The beauty? Each cog has its own CT1/CT2/CT3 compare targets, so your timed waits never collide with another cog's. No shared-resource conflicts!

Why No Interrupts? (Usually)

Here's a controversial P2 feature: it HAS interrupts, but you probably shouldn't use them. Why?

Because with 8 cogs, you don't need interrupts! Instead of interrupting important work, just dedicate a cog to monitoring whatever would have triggered the interrupt:

```

' Traditional (with interrupts):
' Main code runs, gets interrupted, handles event, returns

' Propeller way:
' cog 1: Main code runs uninterrupted
' cog 2: Watches for event continuously
pin_watcher
    TESTP    #BUTTON_PIN wc
    if_c     JMP      #button_pressed
           JMP      #pin_watcher

button_pressed
    WRLONG   ##1, ##button_flag ' Signal other cogs
    JMP      #pin_watcher

```

No interrupt latency, no context switching, no priority inversion. Just dedicated, deterministic monitoring.

Real-World Example: Parallel Sensors

Let's read four different sensors simultaneously:

```
' Each cog runs this with different parameters
sensor_reader
    RDLONG  sensor_pin, ptra[0]    ' Get pin assignment
    RDLONG  hub_addr, ptra[1]      ' Where to store results

read_loop
    ' Read sensor (simplified - real sensors need protocols)
    TESTP   sensor_pin wc
    RCL     value, #1              ' Accumulate bits

    ' Every 32 reads, store to hub
    INCMOD  counter, #31 wc
    if_c WRLONG value, hub_addr
    if_c ADD  hub_addr, #4

    JMP     #read_loop

sensor_pin LONG 0
hub_addr   LONG 0
value      LONG 0
counter    LONG 0
```

Four cogs running this code = four sensors being read truly simultaneously. Try doing that with a single processor and interrupts!

The Medicine Cabinet

Feeling overwhelmed by all this parallel processing? Here's your prescription:

Start simple: Use just one or two cogs at first

```
' Just two cogs - main program and one helper
PUB main()
  COGINIT(COGEXEC_NEW, @helper, 0)
  ' Your main code here
```

Debug one cog at a time: Get each cog working alone before combining

```
' Test cog in isolation first
debug_cog
  DRVH    #MY_DEBUG_LED ' Visual confirmation it's running
  ' Your actual code here
```

Use Spin2 for coordination: Let the high-level language handle the complex stuff

```
' Spin2 manages cogs, PASM2 does the real-time work
PUB orchestrator()
  startSensorCog(0)
  startMotorCog(1)
  startCommsCog(2)
  ' Spin2 coordinates, PASM2 executes
```

Sidetrack

The Philosophy of Parallel

The Propeller's design philosophy comes from a simple observation: in the real world, things happen in parallel, not in sequence.

Consider your car:

- The engine runs continuously
- The radio plays independently
- The climate control maintains temperature
- The dashboard updates displays
- The ABS monitors wheel speed

These aren't taking turns - they're all happening simultaneously. The Propeller models this reality directly. Instead of one processor frantically time-slicing between tasks, you have eight processors each focused on their job.

It's not just different - it's more natural.

Common Gotchas

Save yourself some debugging time:

1. **Cog RAM is copied, not shared** - Changes in cog RAM don't affect hub RAM unless you explicitly write them back
2. **Cogs start at 0** - Always! Your code better be there.
3. **Hub addresses are byte addresses** - cog addresses are long addresses. Don't mix them up!

```
RDLONG value, ##$1000 ' Reads from hub byte address $1000
MOV     value, $100    ' Moves from cog long address $100 (256)
' Note: cog RAM is only 512 longs ($000-$1FF)!
```

4. **PTRA/PTRB are your friends** - These special registers make hub access much easier
5. **Cogs are truly independent** - Stopping one cog doesn't affect others (unless they're waiting for it)

What We've Learned

Look at what you now understand:

- ✓ Why eight processors is simpler than one with interrupts
- ✓ How cogs are structured and limited
- ✓ The hub memory system and egg beater access
- ✓ Multiple ways for cogs to communicate
- ✓ Why interrupts are usually unnecessary
- ✓ How to think in parallel

Your Turn: Experiments

Experiment 1: Cog Counter

Start cogs to increment different hub locations. With **COGEXEC_NEW**, the loop will start up to 7 new cogs (since cog 0 runs Spin2):

```
PUB main() | i
  repeat i from 0 to 7
    COGINIT(COGEXEC_NEW, @counter, $1000 + (i * 4))
  repeat
    ' Monitor the counters in hub RAM
```

continues on next page →

```

DAT
    ORG      0
counter MOV   hub_ptr, ptra      ' Our hub address arrived in PTRa
.loop  RDLONG value, hub_ptr
      ADD    value, #1
      WRLONG value, hub_ptr
      WAITX  ##1_000_000
      JMP    #.loop

hub_ptr LONG   0
value   LONG   0

```

ch02-hub-counters.spin2

Experiment 2: Parallel Pattern

Make 8 LEDs display a moving pattern, with each cog controlling one LED:

```

' Each cog gets LED pin in ptra
    ORG      0
    RDLONG   pin, ptra
    RDLONG   delay, ptrb      ' Different delay per cog

.flash  DRVH   pin
        WAITX  delay
        DRVL   pin
        WAITX  delay
        SHL    delay, #1      ' Double the delay
        CMP    delay, ##100_000_000 wcz
    if_a MOV    delay, ##1_000_000 ' Reset if too long
        JMP    #.flash

```

Coming Up Next

In Chapter 3, “Speaking PASM2”, we’ll dive deep into the instruction set:

- The anatomy of an instruction
- Conditional execution that will blow your mind
- Math operations that actually make sense
- Why PASM2 is unlike any assembly you’ve used

But for now, appreciate what you’ve learned: you understand the Propeller’s parallel philosophy. That’s not just technical knowledge - it’s a new way of thinking about computing.

Have Fun! Remember, parallel processing isn’t harder - it’s different. And different can be wonderful.

Chapter 3: Speaking PASM2

Learning the native tongue

The Hook: One Instruction, Many Powers

Look at this single PASM2 instruction:

```
ADD    value, #1 wc
```

This one line:


- Adds 1 to 'value'
- Optionally sets the carry flag
- Executes in exactly 2 clock cycles
- Can be conditional
- Can even modify itself!

In most processors, that would take multiple instructions. In PASM2, it's just one. Let's learn to speak this powerful language.

Instruction Anatomy 101

Every PASM2 instruction follows the same basic pattern:

PASM2 Instruction Format:



Let's dissect a real instruction:

Example: **if_z add total, #10 wc**



The Basic Vocabulary

Moving Data Around

The **MOV** family - your bread and butter:

```
' Basic move
MOV    dest, source    ' dest = source
MOV    x, #42          ' x = 42 (immediate)
MOV    x, ##70000      ' x = 70000 (32-bit immediate)

' But wait, there's more!
NOT     dest, source    ' dest = NOT source (inverted)
ABS     dest, source    ' dest = |source|
NEG     dest, source    ' dest = -source

' And the mind-blowing ones
ALTD    dest, source    ' Modify NEXT inst's dest field!
ALTS    dest, source    ' Modify NEXT inst's source field!
```

Well, that escalated quickly! Don't worry about **ALTD**/ALTS yet - just know they exist and they're amazing.

Math Without Tears

P2 has hardware multiply and divide. Let that sink in. Hardware. Multiply. And. Divide.

```
' Addition and subtraction
ADD     x, y            ' x = x + y
SUB     x, y            ' x = x - y
ADDS    x, y            ' Signed add
SUBS    x, y            ' Signed subtract

' The revolution: hardware multiply!
MUL     x, y            ' x = x * y (low 32 bits)
MULS    x, y            ' Signed multiply

' And hardware divide!
QDIV    x, y            ' Start division x/y
GETQX   result          ' Get quotient
GETQY   remainder       ' Get remainder
```

Here's a complete multiply example:

```
' Simple 16x16->32 multiply (2 clocks)
MOV     x, #123
MOV     y, #456
```

continues on next page →

↪ continued from previous page

```

    MUL    x, y          ' Result: 123 * 456 = 56088 in x
    ' Uses lower 16 bits of each operand!

    ' For full 32x32->64 multiply, use CORDIC:
    QMUL    x, y          ' Start multiply (uses full 32 bits)
    ' ... 55 clocks of other work ...
    GETQX   low           ' Lower 32 bits of result
    GETQY   high          ' Upper 32 bits of result

```

Uff! In the old days, we'd write loops for this. Now hardware does it!

Logic Operations

Your Boolean friends:

```

    AND    x, mask        ' x = x AND mask
    OR     x, bits        ' x = x OR bits
    XOR    x, toggle      ' x = x XOR toggle
    NOT    x              ' x = NOT x (XOR with $FFFFFFFF)

    ' Bit manipulation
    BITL    x, #5         ' Clear bit 5 of x
    BITH    x, #5         ' Set bit 5 of x
    BITNOT  x, #5         ' Toggle bit 5 of x
    TESTB   x, #5 wc      ' Test bit 5, result in C flag

```

Shifting and Rotating

Moving bits around:

```

    SHL     x, #3         ' Shift left 3 bits
    SHR     x, #3         ' Shift right 3 bits
    SAR     x, #3         ' Arithmetic shift right (signed)
    ROL     x, #3         ' Rotate left 3 bits
    ROR     x, #3         ' Rotate right 3 bits

    ' Variable shifts (amount in register)
    SHL     x, y          ' Shift x left by y bits

    ' Fancy ones
    REV     x              ' Reverse bit order (!!)
```

```

    MERGEB  x              ' Merge bits, not bytes (inverse of SPLITB)

```

Flow Control: Jump!

Unconditional Jumps

```

        JMP    #target      ' Jump to target
        JMP    target      ' Jump to address in target register

' Relative jumps
        JMP    #$$-4        ' Jump back 4 longs (addresses)
        JMP    #$$+8        ' Jump forward 8 instructions

```

Conditional Execution (The Magic)

Here's where PASM2 gets beautiful. ANY instruction can be conditional:

```

if_z    ADD    x, #1        ' Only add if Z flag set
if_nz   ADD    x, #1        ' Only add if Z flag clear
if_c    ADD    x, #1        ' Only add if C flag set
if_nc   ADD    x, #1        ' Only add if C flag clear

```

The basic conditions:

Condition	Meaning
if_z	If Z flag set (result was zero)
if_nz	If Z flag clear (result not zero)
if_c	If C flag set (carry/borrow occurred)
if_nc	If C flag clear
if_c_and_z	If both C and Z set
if_c_or_z	If either C or Z set
if_c_eq_z	If C equals Z
if_c_ne_z	If C not equal to Z

And the comparison conditions (use after **CMP**):

Condition	Meaning
if_a	If above (unsigned greater)
if_b	If below (unsigned less)
if_ae	If above or equal
if_be	If below or equal

The Call/Return Dance

```

CALL    #subroutine    ' Call subroutine
RET                                ' Return from subroutine

' But here's the P2 twist - CALL uses internal stack
subroutine
    ' Do something useful
    RET                    ' Returns to instruction after CALL

' You get 8 levels of hardware stack!
```

The _RET_ Prefix: Return With Benefits

Here's a clever trick the P2 offers. What if you could execute an instruction *and* return from a subroutine in one go? That's exactly what the `_RET_` prefix does.

```

' Normal way: Two instructions
add_and_return
    ADD    x, y            ' Do the add
    RET                                ' Then return (RET is a ~4-cycle branch)

' _RET_ way: One instruction!
add_and_return
    _ret_  ADD    x, y      ' Add AND return (saves 2 cycles)
```

The `_RET_` prefix says: “Execute this instruction, then return.” It’s like getting a free return ticket with your instruction. The add happens, flags get set normally, and then—pop!—you’re back at the caller.

When does `_RET_` NOT return?

Here’s the catch: if the instruction itself branches, no return happens. The branch wins:

```

_ret_  JMP    #somewhere    ' JMP wins - no return
_ret_  CALL    #helper      ' CALL wins - no return
_ret_  DJNZ   count, #loop   ' Branch? No return. Zero? Return!

```

That last one is interesting! If count isn't zero, **DJNZ** branches and no return. But when count hits zero, no branch occurs, so you get your return. Clever, right?

One-Instruction Subroutines

This is where `_RET_` really shines:

```

' Toggle pin 0 - entire subroutine is ONE instruction!
toggle_led
    _ret_    DRVNOT    #0                ' Toggle and return

' Read all inputs - also just one instruction
read_inputs
    _ret_    MOV      result, ina        ' Copy INA and return

' Usage:
    CALL    #toggle_led                ' Blink!
    CALL    #read_inputs                ' result now has INA

```

A normal subroutine needs at least two instructions (the work + **RET**). With `_RET_`, you can have genuinely single-instruction subroutines. Your code gets smaller and faster.

The Medicine: `_RET_` Quick Reference

Pattern	What Happens
<code>_ret_ add x, y</code>	ADD executes, then return
<code>_ret_ jmp #label</code>	JMP executes, NO return (branch wins)
<code>_ret_ djnz n, #loop</code>	If n>0: branch, no return. If n=0: no branch, return

Important: Unlike `RET wcz`, the `_RET_` prefix does NOT restore C and Z flags from the stack. If you need flag restoration, use the regular `RET wcz` instruction.

Labels: Naming Your Places

You've been using labels throughout this chapter without us properly introducing them. How rude of me! Let's fix that.

Global Labels: The Big Signposts

A global label is just a name at the start of a line:

```

DAT                                ORG

send_byte      RDBYTE  x, ptr      ' Global label
               WYPIN   x, tx_pin
               RET

receive_byte   TESTP   rx_pin      wc  ' Another global label
               RDPIN   x, rx_pin
               RET

```

Global labels are visible everywhere in your DAT block. You can jump to them, call them, reference them from Spin2 - they're your main signposts.

Local Labels: The Little Helpers

But here's a problem. What if every routine needs a loop? You can't have two labels called `loop` - the assembler would be terribly confused.

Enter local labels. Prefix a name with a dot (`.`) and it becomes local:

```

DAT                                ORG

send_byte      RDBYTE  x, ptr
.loop         TESTP   tx_pin      wc  ' Local to send_byte
              if_nc   JMP    #.loop
              WYPIN   x, tx_pin
              RET

receive_byte   TESTP   rx_pin      wc  ' New scope begins here
              if_nc   JMP    #.wait
.wait         TESTP   rx_pin      wc  ' Local to receive_byte
              if_nc   JMP    #.wait
              RDPIN   x, rx_pin
.loop         SHR     x, #24        ' Different .loop, OK!
              RET

```

Each global label starts a new “scope”. The `.loop` under `send_byte` is completely separate from the `.loop` under `receive_byte`. You can reuse `.loop`, `.done`, `.retry`, `.exit` to your heart's content.

The Colon Alternative

You might also see local labels with a colon prefix:

```
:loop      DJNZ      count, #:loop    ' Same as .loop
```

Both `:` and `.` work identically. I prefer the dot - it's what modern convention has settled on - but you'll see both in the wild.

Reference Operators: Finding Your Labels

When you reference a label, you need to tell the assembler what you want:

```
' In cog code (after ORG):
    JMP      #my_routine    ' # = immediate cog address
    CALL     #.helper       ' # works for local labels too
    MOV      x, #data_table ' Get cog address of data

' For hub addresses (used with Spin2):
    MOV      ptr, @hub_data ' @ = hub address of label
```

The `#` means “immediate value” - use this for jumps and calls within cog code. The `@` means “hub address” - use this when passing addresses to Spin2 or for hub memory operations.

Scope Boundaries: When Local Labels Reset

Here's the rule: **every global label or data definition starts a new local scope.**

```
func_a      MOV      x, #1          ' Scope #1 begins
.loop       DJNZ     x, #.loop      ' .loop in scope #1

data_block  LONG     0, 0, 0, 0    ' Scope #2 begins (data!)

func_b      MOV      y, #2          ' Scope #3 begins
.loop       DJNZ     y, #.loop      ' .loop in scope #3, OK!
.done       RET
```

This is wonderfully useful - your utility routines can all use `.loop` and `.done` without stepping on each other's toes.

The Medicine: Quick Reference

What	Syntax	Example
Global label	name	my_routine
Local label	.name or :name	.loop, :done
Jump to label	#label	jmp #.loop
Hub address	@label	mov ptr, @data

Common Gotchas

1. **Forgetting the dot:** `loop` is global, `.loop` is local. If you accidentally create a global `loop`, you'll get conflicts.
2. **Scope surprise:** Data definitions (`LONG`, `WORD`, `BYTE`) also start new scopes. If you put data between two parts of a routine, your local labels won't work!
3. **The 30-character limit:** For compatibility with all tools, keep label names under 30 characters. `this_is_a_really_long_label_name` might cause trouble.

Data in DAT Blocks: Your Program's Pantry

Speaking of data definitions starting new scopes... we should probably talk about how to actually declare data! You've seen snippets like `counter LONG 0` scattered through our examples, but there's a whole world of data declaration waiting for you.

The Three Sizes

Just like Goldilocks, you have three choices:

DAT			
my_byte	BYTE	\$FF	' 1 byte (8 bits)
my_word	WORD	\$1234	' 2 bytes (16 bits)
my_long	LONG	\$DEADBEEF	' 4 bytes (32 bits)

When do you use each? Well:

- **BYTE** for characters, small counters, flags, or when every byte of memory counts
- **WORD** for medium values, 16-bit peripherals, or when `BYTE` is too small but `LONG` is wasteful
- **LONG** for everything else - addresses, large numbers, and "I don't want to think about it"

If you're unsure, use LONG. Memory is cheap, debugging overflow errors is not.

Multiple Values on One Line

Here's a convenience - you can list multiple values after a single type:

```
DAT
primes      LONG    2, 3, 5, 7, 11, 13, 17, 19
gpio_pins   BYTE    16, 17, 18, 19, 20, 21
message     BYTE    "Hello, P2!", 0      ' String with null term
```

The assembler just lays them out consecutively in memory. That string? It's just bytes - each character followed by a zero at the end.

Arrays: The Repetition Trick

Need 100 bytes of zeros? Don't type them all out:

```
DAT
buffer      BYTE    0[100]           ' 100 zero bytes
lookup_table WORD    $FFFF[64]       ' 64 words, all $FFFF
scratch_pad LONG    0[32]           ' 32 longs of zeros
```

The [count] syntax repeats the value. This is your friend for buffers, tables, and anywhere you need initialized storage.

You can even combine values and repetition:

```
DAT
mixed_init   BYTE    $AA, $BB, 0[10], $CC, $DD
              '      ^   ^   ^^^^^   ^   ^
              '      Two vals, then 10 zeros, then two more
```

BYTEFIT and WORDFIT: The Safety Net

Here's a subtle trap. What if you accidentally write:

```
DAT
oops         BYTE    1000             ' 1000 won't fit in a byte!
```

The assembler will silently truncate 1000 to 232 (the low 8 bits). Your program will run, but with wrong values. Debugging that is no fun at all.

Enter the safety net:

```
DAT
safe_byte      bytefit 100, 200, 255    ' OK - all fit in a byte
danger_byte    bytefit 100, 200, 300    ' ERROR! 300 > 255, error!

safe_word      wordfit 1000, 50000      ' OK - all fit in a word
danger_word    wordfit 1000, 70000      ' ERROR! 70000 > 65535
```

Use **BYTEFIT** and **WORDFIT** when you want the assembler to check that your constants actually **FIT**. It's like having a compiler catch your mistakes before they become 3 AM debugging sessions.

Alignment: Keeping Things Tidy

The P2 is quite forgiving about alignment - it can read a LONG from any address. But aligned access is faster and cleaner. Sometimes you need to force alignment:

```
DAT
some_bytes     BYTE    1, 2, 3          ' 3 bytes
               ALIGNW          ' Align to word boundary
next_word      WORD    $1234            ' Now properly aligned

more_bytes     BYTE    "ABC"            ' 3 bytes
               ALIGNL          ' Align to long boundary
next_long      LONG    $DEADBEEF        ' Now on a 4-byte boundary
```

When does alignment matter? Mostly when you're:

- Mixing data sizes in the same DAT block
- Creating structures that Spin2 code will access
- Optimizing for maximum hub access speed

If you're just starting out, don't worry about it. Add alignment when you hit problems.

A Complete Example

Let's put it all together:

```
DAT          ORG

' Your code goes here
entry        MOV    ptr, ##buffer_addr
             MOV    count, BUFFER_SIZE
.fill        WRBYTE fill_value, ptr++
             DJNZ   count, $.fill
             JMP    #$

' Constants (read-only, effectively)
```

continues on next page →

↪ continued from previous page

```

fill_value    BYTE    $55
              ALIGNL
BUFFER_SIZE   LONG    256

' Lookup tables
              ALIGNL
sin_table     WORD    0, 1608, 3212, 4808, 6393    ' Sine table
              WORD    7962, 9512, 11039, 12540    ' ... continues

' Working storage
              ALIGNL
buffer_addr   LONG    0                        ' Set by Spin2 at startup
temp         LONG    0
result       LONG    0

' Reserve uninitialized space
              ALIGNL
scratch      RES     16                        ' Reserve 16 longs

```

Notice the pattern: code first, then constants, then working storage, then reserved space. This keeps things organized and makes your DAT block readable.

The Medicine: Quick Reference

Declaration	Meaning	Example
byte val	8-bit value	byte \$FF
word val	16-bit value	word \$1234
long val	32-bit value	long \$DEADBEEF
val[n]	Repeat n times	byte 0[100]
bytefit	Byte with range check	bytefit 100, 200
wordfit	Word with range check	wordfit 1000, 50000
alignw	Align to word	(no value)
alignl	Align to long	(no value)
res n	Reserve n longs	res 16

Common Gotchas

1. **Forgetting the label:** Every piece of data needs a label if you want to reference it. Anonymous data just wastes space.
2. **String termination:** BYTE "Hello" doesn't include a null terminator. Add , 0 if you need one!

3. **RES is in longs:** RES 10 reserves 10 *longs* (40 bytes), not 10 bytes. This trips up everyone at least once.
4. **Alignment after RES:** RES doesn't affect alignment. If you need alignment after reserved space, add an explicit ALIGNL or ALIGNW.

Including External Files: FILE

Sometimes you have binary data sitting in a file - a font bitmap, a sound sample, a pre-computed lookup table. Rather than manually converting it to hex values (ugh!), you can pull it straight in:

```
DAT
font_data      file    "myfont.bin"      ' Import entire file
sound_sample   file    "beep.raw"        ' Raw audio data
lut_table      file    "precalc.dat"     ' Pre-computed values
```

The assembler reads the file at compile time and drops its raw bytes right into your DAT block. The label gives you a way to reference where the data starts.

Where does it search?

1. Same folder as your source file
2. Library paths (if configured)

Practical example - embedded bitmap:

```
DAT          ORG
entry        MOV      ptr, ##@splash_screen
              ' ... display routine using ptr

' The image data lives right here in your code
splash_screen file    "logo_128x64.bin"    ' 1024 bytes of pixel data
```

No conversion scripts, no copy-paste errors. Just reference the file and it's part of your program. The icing on the cake? If you update the file, recompile, and your program has the new data.

String and Data Generation Methods

Beyond manually typing values, you have some helpers for creating data:

String Addresses in PASM

Need a string's address? In PASM you point at a DAT label with `##@:`

```

        MOV    ptra, ##@hello      ' ptra points to "Hello!" in hub
        CALL   #print_string

        MOV    ptra, ##@err_msg    ' Another string
        CALL   #print_string

        ALIGNL
hello    BYTE   "Hello!", 0
        ALIGNL
err_msg  BYTE   "Error: ", 0

```

The string lives in your DAT block, and @label gives its hub address. Spin2 has an even shorter shortcut - @"Hello!" composes the string and hands back its address with no separate label - but that one is Spin2-only, not PASM.

STRING("text") and LSTRING("text")

These work similarly but in different contexts:

```

{Spin2_v43}                ' enable LSTRING (a v43+ keyword)
' In Spin2 code (not PASM):
DEBUG(STRING("Temperature: ")) ' Zero-terminated string address
DEBUG(LSTRING("Status"))      ' Length byte first, then string

```

STRING() returns the hub address of a zero-terminated string - same as what C programmers expect. LSTRING() puts a length byte at the front, which is handy when you need to know the string length without scanning for null. LSTRING() is a newer keyword, so it needs a {Spin2_v43} (or later) directive at the top of your file before the compiler will recognize it.

BYTE[], WORD[], LONG[] - Data Arrays

In Spin2, you can compose inline data and get its address:

```

{Spin2_v43}                ' BYTE/WORD/LONG composers (v43+)
tbl := BYTE(10, 20, 30, 40, 50) ' Returns address of byte array
config := LONG($DEAD_BEEF, $CAFE_BABE)

```

These are primarily Spin2 features, but they generate hub data that your PASM code can access if you know the addresses.

The Pattern:

Method	Result	Use Case
file "name"	Raw binary data	Images, audio, lookup tables
##@label	String address	Inline strings in PASM (DAT label)
@ "text"	String address	Quick inline strings (Spin2 shortcut)
STRING("text")	Zero-terminated string address	Spin2 string constants
LSTRING("text")	Length-prefixed string	Spin2, needs {Spin2_v43}

The Flags: C and Z (and Q!)

Flags are your friends. They remember things:

```

' Z flag - was the result zero?
      SUB    x, y wz      ' Set Z if x-y equals zero
if_z   JMP    #equal      ' Jump if they were equal

' C flag - did we carry/borrow?
      ADD    x, y wc      ' Set C if addition overflowed
if_c   JMP    #overflow   ' Handle overflow

' Both at once!
      CMP    x, y wcz     ' Compare and set both flags
if_a   JMP    #x_greater  ' Jump if x > y (unsigned)

```

The Q flag is special - it's used by CORDIC operations (Chapter 7).

Special Instructions That Will Blow Your Mind

SKIP - The Instruction Skipper

```

SKIP    ##%00001010    ' LSB first: 1=skip, 0=execute
ADD     x, #1           ' Executed (bit 0 = 0)
ADD     y, #1           ' Skipped! (bit 1 = 1)
ADD     z, #1           ' Executed (bit 2 = 0)
SUB     a, #1           ' Skipped! (bit 3 = 1)
SUB     b, #1           ' Executed (bit 4 = 0)
' ... bit N controls the Nth instruction after SKIP

```

This is like having conditional execution on steroids!

REP - Hardware Loops

```

REP    #4, #5          ' Repeat next 4 instructions 5 times
ADD    x, #1
SUB    y, #1
ROL    z, #1
ROR    w, #1
' These 4 instructions execute 5 times total
' No loop overhead!

```

ALTD/ALTS - Instruction Modification

```

' Modify the next instruction's destination
MOV    index, #10
ALTD   index, #array  ' Next instruction's dest = array+10
MOV    0-0, value     ' Actually moves to array[10]!

```

This replaces self-modifying code from P1. Much cleaner!

Sidetrack

Hub-Exec Note: All ALTx instructions — **ALTI, ALTS, ALTD, ALTR, ALTB, ALTSN, ALTSB, ALTSW, ALTGN, ALTGB, ALTGW** — work identically in cog-exec and hub-exec modes. They act on the next pipelined instruction regardless of whether it came from cog/LUT memory or the hub-prefetch FIFO. So when you graduate to hub execution in Chapter 10, every ALTx trick you learn here still works.

Real-World Example: Fast Memory Copy

Let's combine what we've learned:

```

' Fast block copy using REP
fast_copy
MOV    ptrA, ##source_addr  ' Source pointer
MOV    ptrB, ##dest_addr    ' Destination pointer

REP    #2, ##256            ' Repeat 256 times
RDLONG temp, ptrA++          ' Read and increment
WRLONG temp, ptrB++          ' Write and increment
' 1024 bytes (256 longs) copied with no loop overhead!

temp   LONG    0

```


The Medicine Cabinet

Feeling overwhelmed? Here's your simplified prescription:

Minimum Instructions to Know

```
' Moving data
    MOV    dest, source    ' Copy data

' Math
    ADD    dest, source    ' Addition
    SUB    dest, source    ' Subtraction

' Logic
    AND    dest, source    ' AND operation
    OR     dest, source    ' OR operation

' Flow
    JMP     #label         ' Jump
    CALL    #label         ' Call subroutine
    RET                               ' Return

' Flags
    CMP     x, y wcz       ' Compare and set flags
    if_z    JMP    #label  ' Conditional jump
```

Master these 10 instructions and you can write real programs!

Common Gotchas

1. Immediate values:

- #value for 9-bit immediates (0-511)
- ##value for 32-bit immediates
- Forgetting # uses the register at that address!

2. Flag confusion:

- wz sets Z flag, wc sets C flag, wcz sets both
- No flag update means flags unchanged

3. PTR A/PTR B are special:

```
RDLONG x, ptr++    ' Read and auto-increment
RDLONG x, ++ptr    ' Pre-increment then read
RDLONG x, ptr--    ' Read and auto-decrement
RDLONG x, ptr[5]   ' Read from ptr + 5*4
```

4. Address confusion:

- Cog addresses are in longs (0-511)
- Hub addresses are in bytes (0-524287)

Your Turn: Experiments

Experiment 1: Conditional Counter

Count up if button pressed, down if not:

```

        ORG      0

.loop   TESTP    #BUTTON_PIN wz ' Test button
if_c    ADD      counter, #1      ' Increment if pressed
if_nc   SUB      counter, #1      ' Decrement if not

        WRLONG   counter, ##HUB_ADDR ' Display count
        WAITX    ##1_000_000
        JMP      #.loop

counter LONG      0

```

Experiment 2: Pattern Matcher

Find a pattern in data:

```

        ORG      0

        MOV      pattern, ##$DEADBEEF
        MOV      ptr, ##data_start

.search RDLONG    value, ptr++
        CMP      value, pattern wz
if_z     JMP      #.found
        CMP      ptr, ##data_end wcz
if_b     JMP      #.search
        JMP      #.not_found

.found   ' Pattern found!
        DRVH     #SUCCESS_LED
        JMP      #$

.not_found
        DRVH     #FAIL_LED
        JMP      #$

```

Experiment 3: Speed Test

Compare multiply methods:

```
' Method 1: Hardware multiply
    GETCT    start_time
    MUL      x, y
    GETCT    end_time
    SUB      end_time, start_time
    ' Result: 2 clocks!

' Method 2: Shift and add (old school)
    GETCT    start_time
    ' ... shift/add loop here
    GETCT    end_time
    ' Result: Many more clocks!
```

Sidetrack

Why PASM2 Is Different

Most assembly languages are thin wrappers over hardware. PASM2 is different - it's designed for humans:

1. **Symmetry:** Every instruction can use every addressing mode
2. **Orthogonality:** Features combine predictably
3. **Conditional everything:** Not just jumps, ANY instruction
4. **No special cases:** General-purpose registers, no accumulator

This isn't accident - it's philosophy. The P2 was designed to make assembly programming pleasant.

What We've Learned

- ✓ Instruction anatomy and structure
- ✓ Basic data movement and math
- ✓ Hardware multiply and divide (!)
- ✓ Conditional execution on any instruction
- ✓ Special instructions (SKIP, **REP**, ALT*)
- ✓ flag operations and testing
- ✓ Why PASM2 is human-friendly

Coming Up Next

Chapter 4, “The Hub Connection”, explores:

- Reading and writing hub memory
- The FIFO and fast block transfers
- Hub execution mode
- Sharing data between cogs

You now speak basic PASM2. Time to learn how cogs communicate!

Have Fun! Remember, PASM2 isn't like other assembly languages - it's actually enjoyable!

Chapter 4: The Hub Connection

How cogs share and care

The Hook: Instant Communication

```
' cog 1: Leave a message
    WRLONG  ##$DEADBEEF, ##$1000

' cog 2: Get the message
    RDLONG  message, ##$1000
    ' message now contains $DEADBEEF!
```

That's it - cogs talking through hub memory. But there's so much more...

Reading from Hub

The basics are simple:

```
RDBYTE  value, hubaddr  ' Read 1 byte
RDWORD  value, hubaddr  ' Read 2 bytes (word)
RDLONG  value, hubaddr  ' Read 4 bytes (long)

' With PTRB/PTRB magic
RDLONG  value, ptr++    ' Read and increment pointer
RDLONG  value, ++ptr    ' Increment then read
RDLONG  value, ptr[5]   ' Read from ptr + 5*4
```

Writing to Hub

Just as easy:

```
WRBYTE  value, hubaddr  ' Write 1 byte
WRWORD  value, hubaddr  ' Write 2 bytes
WRLONG  value, hubaddr  ' Write 4 bytes

' The mighty block transfer
SETQ    #16-1           ' Transfer 16 longs
RDLONG  buffer, hubaddr  ' Reads 16 longs in one go!
```

The FIFO Pipeline

Here's where P2 gets serious about speed:

```
' Start the FIFO
    RDFAST  #0, ##data_start  ' Start fast read

' Now read at maximum speed
.loop  RFLONG  value          ' Read from FIFO
      ' Process value
      DJNZ    count, #.loop   ' Decrement and jump if not zero

' No hub timing worries - FIFO handles it all!
```

Real-World Example: Video Buffer

```
' Fast screen clear using block transfer
' Note: SETQ/SETQ2 maximum is 511 (for 512 longs)
' For larger fills, we loop in 512-long chunks
clear_screen
    MOV      hub_ptr, ##screen_buffer
    MOV      chunks, ##640*480/512  ' Number of 512-long chunks
    MOV      color, ##$00_00_00_00  ' Black

.loop  SETQ   #512-1              ' Transfer 512 longs (max)
      WRLONG  color, hub_ptr      ' Fill this chunk
      ADD     hub_ptr, ##512*4    ' Advance 512 longs (2KB)
      DJNZ    chunks, #.loop
      ' Full screen cleared with minimal loop overhead!
```

The Medicine Cabinet

Simple hub access pattern:

```
' Just use PTR for everything
    MOV      ptr, ##hub_address
    RDLONG   value, ptr++
    ' That's all you really need!
```

Your Turn: Experiments

Now you know the moves. Try these to make them stick:

1. **Round-trip:** Write a long to hub at address \$1000, then read it back into a different register. Verify the values match using a comparison and a flag.
2. **Block fill:** Use **SETQ + WRLONG** to fill 64 consecutive longs in hub memory with the value \$DEAD_BEEF. Then read them back one by one and confirm.
3. **FIFO race:** Use **RDFAST** to stream a sequence of longs from hub. Compare the throughput to a loop of individual **RDLONG** instructions. The difference should astonish you.
4. **Pointer arithmetic:** Set **PTRA** to a hub address, then use **RDLONG** `val, ptr[3]` to read the third long ahead — without moving the pointer. Experiment with `++`, `--`, and indexed forms.

Common Gotchas

Before you pull your hair out debugging hub access:

1. **Forgetting the ##** — hub addresses are 20-bit. A bare `#address` only encodes 9 bits, so you'll hit the wrong memory. Always use `##` for hub addresses outside the 0–511 range.
2. **Long-aligned access** — **RDLONG** and **WRLONG** work on long boundaries. If your address isn't a multiple of 4, the P2 silently masks the low bits. For byte-precise access, use **RDBYTE** / **WRBYTE**.
3. **SETQ block size** — **SETQ** `##N-1` transfers `N` longs (not `N-1`). The `-1` is because the encoded field is “count minus one.” Off-by-one bugs love this one.
4. **PTRA vs. PTRB** — Both are pointer registers, but they're independent. Don't assume **PTRA** holds anything if you've been working with **PTRB**.
5. **Hub-exec vs. cog-exec timing** — In cog-exec, **RDLONG** is 9–16 cycles. In hub-exec, it's 9–26 cycles. Inner loops that hammer hub should run from cog memory when possible.

What We've Learned

Look at all the ground we've covered:

- ✓ How to read and write hub memory from cog code
- ✓ Why pointer registers (PTRA, PTRB) exist and when to use them
- ✓ How the FIFO pipeline accelerates sequential hub reads
- ✓ Block transfers with **SETQ** for moving big chunks at once
- ✓ The pattern for clearing or filling buffer-sized regions

You now have the entire cog↔hub vocabulary at your disposal.

Coming Up Next

In Chapter 5, we'll unleash the P2's mathematical muscle:

- Hardware multiply and divide that don't require library calls
- 64-bit arithmetic without sweating
- A preview of the CORDIC coprocessor (it gets a full chapter later)
- Fixed-point math for fast trigonometry

But first, take a break. Hub access is one of those topics where letting it sit a day helps it lock in. Try one of the experiments above tomorrow.

Have Fun! And remember — the FIFO is doing real work even when your code looks idle. Trust it.

Chapter 5: Mathematics Unleashed

Hardware multiply and divide - finally!

The Hook: Hardware Multiply

```
MUL      x, y                ' 16x16->32 bit unsigned multiply
' Result in x (lower 16 bits of each operand used)

' For full 32x32->64 bit multiply, use CORDIC:
QMUL     x, y                ' Start 32x32->64 multiply
' ... other work (55 clocks) ...
GETQX    low                 ' Get lower 32 bits
GETQY    high                ' Get upper 32 bits
```

Remember doing this with shifts and adds? Those days are over!

The Multiplication Revolution

```
' Unsigned multiply
      MUL      result, value    ' result = low 32 bits

' Signed multiply
      Muls     result, value    ' Signed version

' Scale by a power of two
      SHR      result, #1       ' Scale by 0.5 (divide by 2)
```

Division Without Tears

```
' Start division
      QDIV     dividend, divisor ' Start the operation

' Get results (takes 55 clocks)
      GETQX    quotient         ' Get quotient
      GETQY    remainder        ' Get remainder

' Fractional division
      QFRAC    numerator, denominator
      GETQX    fraction         ' 32-bit fraction
```

64-Bit Operations

```
' 64-bit add
    ADD    low1, low2 wc
    ADDX   high1, high2

' 64-bit multiply (uses CORDIC)
    QMUL   x, y           ' Start 32x32->64 multiply
    ' ... 55 clocks ...
    GETQX  low             ' Lower 32 bits
    GETQY  high           ' Upper 32 bits
```

Real-World Example: Fixed-Point Math

```
' 16.16 fixed point multiply (uses CORDIC for full precision)
fixed_mul
    QMUL   a, b           ' Start 32x32->64 unsigned multiply
    ' ... 55 clocks (do other work) ...
    GETQX  low            ' Lower 32 bits
    GETQY  high           ' Upper 32 bits
    ' Extract middle 32 bits for 16.16 result:
    SHL    high, #16      ' Upper 16 bits of result
    SHR    low, #16       ' Lower 16 bits of result
    MOV    a, low         ' Combine for 16.16 format
    OR     a, high
```

The Medicine Cabinet

Quick math reference:

- **MUL** D, S — unsigned $16 \times 16 \rightarrow 32$
- **MULS** D, S — signed $16 \times 16 \rightarrow 32$
- **QMUL** D, S — full $32 \times 32 \rightarrow 64$ (read via **GETQX**/**GETQY** after 55 clocks)
- **QDIV** D, S — full 32-bit divide (read via **GETQX** quotient / **GETQY** remainder after 55 clocks)
- **QFRAC** D, S — fractional divide (returns 32-bit fraction in **GETQX**)
- 64-bit add: **ADD** + **ADDX** chained with **WC**

For everyday integer work, **MUL**/**MULS** are 2 clocks and you're done. For precision (full 64-bit results, fixed-point math, signed division), **QMUL**/**QDIV** route through the CORDIC and pay 55 clocks — but they don't block the cog, so you can interleave other work.

Your Turn: Experiments

Stretch your math muscles:

1. **Compute the average:** Read 8 longs from hub, sum them with **ADD**, then divide by 8 using a **SHR** (or **QDIV** if you want exact). Compare both approaches.
2. **Fractional reciprocal:** Use **QFRAC** to compute $2^{32} / x$ for various x values. You've just built a hardware reciprocal table.
3. **Pipeline overlap:** Start a **QMUL**, do 30+ clocks of other work (compute something else, update a counter), then **GETQX/GETQY**. Measure total cycles vs. doing the multiply blocking-style.
4. **64-bit counter:** Build a 64-bit increment loop using **ADD + ADDX**. Add to it in a tight loop and watch the high word advance when the low one wraps.

Common Gotchas

The math instructions hide a few traps:

1. **MUL is unsigned, MULS is signed** — Using the wrong one on negative values gives spectacular nonsense. When in doubt, **MULS**.
2. **MUL gives only low 32 bits** — For the full 64-bit result, you must use **QMUL + GETQX/GETQY**. The high word is silently discarded by **MUL**.
3. **GETQX/GETQY block until ready** — They wait for the CORDIC. If you call them too early, your cog stalls. If you call them later than necessary, you've wasted cycles. The sweet spot is starting the CORDIC, doing exactly 55 clocks of other work, then reading.
4. **Don't queue a new CORDIC op while one is pending** — A second **QMUL/QDIV/QFRAC** before reading results overwrites the queue. Use **GETQX/GETQY** first.

What We've Learned

A whole arithmetic library, in your back pocket:

- ✓ Hardware multiply (**MUL**, **MULS**) in just 2 clocks
- ✓ Division and modulo via **QDIV + GETQX/GETQY**
- ✓ Fractional division via **QFRAC** for fast reciprocals
- ✓ 64-bit chains using **ADD/ADDX**, **SUB/SUBX** with **WC**
- ✓ Fixed-point patterns for 16.16 math

You no longer have to write multiply-by-shift-and-add. Those days really are over.

Coming Up Next

In Chapter 6 we'll meet the P2's most quietly powerful feature: **conditional execution**. Every instruction can be conditional — no branches needed, deterministic timing preserved. We'll cover:

- The C and Z flags and how to update them
- Building complex conditions with **WC**, **WZ**, **WCZ** combined with **IF_x** prefixes
- **SKIP** and **SKIPF** for skipping patterns of instructions

If conditional execution doesn't change how you think about flow control, nothing will.

Have Fun! And remember — the CORDIC is a coprocessor with infinite patience. Use it.

Chapter 6: Flags and Decisions

Making choices at machine speed

The Hook: Any Instruction Can Be Conditional

```

        CMP     x, y wcz      ' Compare x and y
if_a    MOV     result, x     ' If x > y, result = x
if_be   MOV     result, y     ' If x <= y, result = y
        ' Max function in 3 instructions!

```

The C and Z Flags

```

' Z Flag - Zero detection
        SUB     x, y wz      ' Z=1 if x equals y
if_z    JMP     #equal      ' Jump if equal

' C Flag - Carry/Borrow
        ADD     x, y wc      ' C=1 if overflow
if_c    JMP     #overflow   ' Handle overflow

```

Complex Conditions

```

' Combining flags
        CMP     x, y wcz
if_a    JMP     #greater     ' x > y (unsigned)
if_b    JMP     #less        ' x < y (unsigned)
if_z    JMP     #equal       ' x == y

' Signed comparisons
        CMPS    x, y wcz     ' Signed compare
if_gt   JMP     #greater     ' x > y (signed)
if_lt   JMP     #less        ' x < y (signed)

```

Skip Patterns - Conditional Blocks

```

SKIPF  pattern      ' Set skip pattern
ADD    x, #1        ' Maybe executed
SUB    y, #1        ' Maybe executed
MOV    z, #0        ' Maybe executed
' Pattern determines what runs!

```

The Medicine Cabinet

Conditional execution quick reference:

- **WC** — write the C flag with the instruction's carry-out
- **WZ** — write the Z flag based on result == 0
- **WCZ** — write both flags
- **IF_C, IF_NC** — execute next instruction only if C set / clear
- **IF_Z, IF_NZ** — same for Z
- **IF_A, IF_B, IF_AE, IF_BE** — unsigned comparisons after **CMP**
- **IF_GT, IF_LT, IF_GE, IF_LE** — signed comparisons after **CMPS**
- **IF_E, IF_NE** — equal / not equal (same as **IF_Z/IF_NZ** after compare)

Any instruction can carry a condition prefix and still take exactly 2 clocks. No branches, no pipeline flush, no surprise.

Your Turn: Experiments

Practice the conditional mindset:

1. **Branchless absolute value:** Compute `result = ABS(x)` using **ABS** — then prove to yourself the same thing works branchless using **CMPS** + a conditional **NEG**.
2. **Saturating add:** Add two unsigned values; if the result overflows (**WC** sets C), clamp to `$FFFF_FFFF`. No jumps.
3. **Min/max:** Build a min function using **CMP** + conditional **MOV**. Then build max. Three instructions each.
4. **Skip-pattern selector:** Use **SKIPF** with a runtime bit pattern to selectively execute 4 different instructions based on a state byte. This is the building block for hand-coded jump tables.

Common Gotchas

The conditional mindset has its own traps:

1. **Forgetting the effect** — `CMP x, y` with *no flag effect* doesn't update C or Z. You need **WC**, **WZ**, or **WCZ**. The default for **CMP** is no update.
2. **Signed vs. unsigned comparison** — **CMP** sets flags for unsigned semantics (**IF_A/IF_B**). **CMPS** sets them for signed (**IF_GT/IF_LT**). Mixing them produces baffling bugs near zero and at the high bit.
3. **One instruction follows the prefix** — `IF_C MOV x, y` conditionally runs *only* the **MOV**. The instruction after that runs unconditionally. To skip a block, use **SKIPF**.
4. **C flag inversion on subtract** — **SUB** sets C on *borrow*, not carry. So after `SUB a, b wc, C=1` means $a < b$ (unsigned). Many P1 programmers expect the opposite — re-check.
5. **Skip patterns are LSB-first** — In **SKIPF**, bit 0 controls the next instruction, bit 1 the one after, etc. The numbering can feel backward; sketch it out.

What We've Learned

You can now think like a P2 programmer:

- ✓ Every instruction can carry a condition prefix
- ✓ Flags are updated only when you ask (**WC**, **WZ**, **WCZ**)
- ✓ Unsigned and signed comparisons use different instructions (**CMP** vs. **CMPS**)
- ✓ Complex multi-way logic without a single jump using **IF_x** prefixes
- ✓ **SKIP** and **SKIPF** for skipping arbitrary patterns of instructions

You've learned to write branchless code that preserves deterministic timing. That's the deepest part of P2 thinking.

Coming Up Next

Chapter 7 reveals the CORDIC coprocessor in full — the math beast you glimpsed in Chapter 5. Trigonometry at the speed of logic gates:

- Rotate points in three instructions
- Polar↔Cartesian conversion
- Hardware square root and natural log
- Pipelined operations for graphics and signal processing

If you thought hardware multiply was nice, wait until you see what the CORDIC does.

Have Fun! And remember — every **JMP** you avoid is a pipeline flush you didn't pay for.

Chapter 7: CORDIC Magic

Trigonometry at the speed of logic gates

The Hook: Rotate a Point in 3 Lines

Let me show you something that, on most processors, would take a coffee break of instructions, a math library, and probably a tear or two:

```
' Rotate point (x,y) by angle - that's it!
  SETQ   y_coord      ' Set Y coordinate
  QROTATE x_coord, angle ' Start rotation by angle
  GETQX   new_x        ' Get rotated X (55 clocks later)
  GETQY   new_y        ' Get rotated Y
```

Read that again. *Four lines*, and a 2D rotation is done. No floating-point library. No lookup tables. No iterative approximation. The P2 has a dedicated trigonometric coprocessor sitting next to every cog, just waiting for you to wake it up. You're about to learn how.

Let me show you something even more impressive:

```
' Calculate sine and cosine simultaneously
  QROTATE ##$7FFF_FFFF, angle ' D=radius (max), S=angle
  GETQX   cosine              ' cos(angle); $7FFF_FFFF = 1.0
  GETQY   sine                ' sin(angle); $7FFF_FFFF = 1.0
  ' Both trig functions in 55 clocks total!
```

What Just Happened?

CORDIC stands for COordinate Rotation DIgital Computer. It's a method invented in 1959 for calculating trigonometric functions using only shifts and adds - no multiplies needed! Each P2 cog has its own dedicated CORDIC unit built into the hardware.

Think of CORDIC as your mathematical co-processor that can:

- Rotate points around the origin
- Convert between rectangular and polar coordinates
- Calculate sine, cosine, tangent
- Compute square roots and magnitudes
- Find arctangent (angle between points)
- Even do logarithms and exponentials!

All of this in exactly 55 clock cycles. Every time. No variation.

The CORDIC Pipeline - Your Mathematical Assembly Line

Here's the beautiful part: CORDIC operations are pipelined. While one calculation is running, you can start another. Let's see what that buys us:

```
' Generate sine wave samples rapid-fire
MOV     angle, #0
MOV     count, #256

generate
QROTATE ##$7FFF_FFFF, angle    ' D=radius, S=angle
ADD     angle, ##$0100_0000    ' Increment angle (no wait!)

    ' Do other work while CORDIC calculates
ADD     sample_ptr, #4
SUB     count, #1

GETQY   sample                 ' Get sine result
WRLONG  sample, sample_ptr     ' Store it

TJNZ    count, #generate      ' Test-jump-not-zero
    ' Generated 256 samples with perfect overlap!
```

The pipeline means you're not really waiting 55 clocks - you're getting useful work done while CORDIC churns away in the background. Uff! That's free math!

Core CORDIC Operations

QROTATE - The Rotation Engine

Here's a subtle detail: CORDIC operations work on 2D coordinates (X, Y), but the **QROTATE** instruction only takes one coordinate directly. The solution? **SETQ** loads the Y coordinate into the Q register, then **QROTATE** takes X from its first operand. It's a two-instruction dance that becomes second nature:

```
' Basic rotation: rotate point (x,y) by angle
SETQ    y                        ' First: load Y into Q register
QROTATE x, angle                 ' Then: X from operand, Y from Q
GETQX   new_x                   ' Result: X' = X*cos(θ) - Y*sin(θ)
GETQY   new_y                   ' Result: Y' = X*sin(θ) + Y*cos(θ)
```

The angle format is special: it's a 32-bit unsigned value where:

- \$0000_0000 = 0 degrees
- \$4000_0000 = 90 degrees
- \$8000_0000 = 180 degrees
- \$C000_0000 = 270 degrees
- \$FFFF_FFFF = just under 360 degrees

This makes angle math incredibly easy - just use regular addition and subtraction!

QVECTOR - From Rectangular to Polar

```
' Convert (x,y) to polar (radius, angle)
  QVECTOR x, y          ' X in D, Y in S (no SETQ for QVECTOR)
  GETQX  radius         ' sqrt(x2 + y2)
  GETQY  angle          ' atan2(y, x)
```

Perfect for:

- Finding distances between points
- Converting joystick input to angle/magnitude
- Radar and sonar applications

The Power of 32-Bit Precision

CORDIC uses 32-bit precision throughout:

- Angles: 32 bits (0.00000008 degree resolution!)
- Coordinates: 32 bits signed
- Results: Full 32-bit or 64-bit when needed

Real-World Example: Spinning a Sprite

Let's rotate a sprite around its center:

```
' Rotate sprite vertices around center
' (PTRA for vertex iteration - only PTRB support ++ post-increment.)
rotate_sprite
    MOV    ptra, ##sprite_data
    MOV    vertex_count, #4          ' 4 corners

next_vertex
    RDLONG x, ptra++                 ' Get X coordinate, advance pointer
    RDLONG y, ptra++                 ' Get Y coordinate, advance pointer
```

continues on next page →

```

' Center sprite at origin
SUB    x, center_x
SUB    y, center_y

' Rotate by current angle
SETQ   y
QROTATE x, rotation_angle

' While waiting, we can prepare
MOV    temp_x, center_x
MOV    temp_y, center_y

' Get rotated coordinates
GETQX  x
GETQY  y

' Translate back to position
ADD    x, temp_x
ADD    y, temp_y

' Store rotated vertex (back up PTR 8 to overwrite source longs)
SUB    ptr, #8
WRLONG x, ptr++
WRLONG y, ptr++

DJNZ   vertex_count, #next_vertex

' Increment rotation for animation
ADD    rotation_angle, ##$0100_0000 ' ~1.4 deg (1/256 rotation)

```

Your Turn: CORDIC Experiments

Your Turn

Your Turn: Create a circular motion pattern

Starting code:

continues on next page →

→ continued from previous page

```

        ORG      0

        MOV      angle, #0
        MOV      radius, ##100          ' 100 pixel radius

.loop   QROTATE radius, angle          ' D=X (radius), S=angle
        ' Add code to:
        ' 1. Get X,Y coordinates
        ' 2. Add screen center offset
        ' 3. Draw pixel at that position
        ' 4. Increment angle
        ' 5. Loop back to .loop

```

Goal: Make a dot trace a perfect circle on screen Hint: After **QROTATE**, use **GETQX**/getqy to get coordinates Success Check: Smooth circular motion, no gaps

Your Turn

Your Turn: Distance calculator

Starting code:

```

        ' Calculate distance between two points
        MOV      x1, #10
        MOV      y1, #20
        MOV      x2, #40
        MOV      y2, #60

        ' Calculate differences
        SUB      x2, x1          ' dx
        SUB      y2, y1          ' dy

        ' Your code here: use qvector to find distance

```

Goal: Calculate the distance between the two points Hint: **QVECTOR** dx, dy gives you radius (distance) in QX Success Check: Distance should be 50 units

The Medicine Cabinet

Feeling overwhelmed by all this trigonometry? Here's your simplified prescription:

Too Complex? Just remember these three patterns:

Pattern 1: Get sine/cosine

```
QROTATE ##$7FFF_FFFF, angle    ' D=radius, S=angle
GETQX  cos_value
GETQY  sin_value
```

Pattern 2: Rotate a point

```
SETQ    y
QROTATE x, angle
GETQX   new_x
GETQY   new_y
```

Pattern 3: Get distance

```
QVECTOR dx, dy
GETQX   distance
```

Master these three and you can do 90% of what you need!

Advanced CORDIC: The Pipeline Dance

Here's where CORDIC gets really powerful - overlapping operations:

```
' Process multiple points while calculating
process_points
    MOV    count, #16
    MOV    ptr, ##point_array

    ' Start first calculation
    RDLONG x, ptr++
    RDLONG y, ptr++
    SETQ   y
    QROTATE x, angle

process_loop
    ' Start next calculation immediately
    RDLONG x, ptr++ wz    ' Z flag tells us if done
    if_nz RDLONG y, ptr++
    if_nz SETQ   y
```

continues on next page →

```

if_nz QROTATE x, angle      ' New calculation starts

    ' Get previous result
    GETQX  prev_x
    GETQY  prev_y

    ' Store previous result
    WRLONG prev_x, ptrb++
    WRLONG prev_y, ptrb++

    DJNZ   count, #process_loop

    ' Don't forget last result!
    GETQX  prev_x
    GETQY  prev_y
    WRLONG prev_x, ptrb++
    WRLONG prev_y, ptrb++

```

See what happened? We started each new CORDIC operation immediately after the previous one, then retrieved results later. This pipeline approach means we're effectively getting one rotation every few instructions instead of waiting 55 clocks each time!

CORDIC for Graphics

Want to draw a spiral? CORDIC makes it trivial:

```

' Expanding spiral generator
spiral
    MOV    angle, #0
    MOV    radius, #1

draw_spiral
    QROTATE radius, angle      ' D=X (radius), S=angle
    GETQX  x
    GETQY  y

    ' Convert to screen coordinates
    SAR    x, #16              ' Scale down
    SAR    y, #16
    ADD    x, #320             ' Center X
    ADD    y, #240             ' Center Y

    ' Plot pixel (simplified)
    CALL   #plot_pixel

    ' Expand spiral
    ADD    angle, ##$0400_0000 ' Rotate 5.625 degrees (1/64 turn)
    ADD    radius, ##100        ' Expand slowly

```

```

        CMP        radius, ##30000 wcz
if_b JMP        #draw_spiral

```

CORDIC for Audio

Generate perfect sine waves for audio:

```

' Audio tone generator using CORDIC
tone_generator
    MOV        phase, #0
    MOV        frequency, ##$0100_0000 ' ~1.4 deg/sample (1/256 rot)

sample_loop
    QROTATE    ##$7FFF_FFFF, phase      ' D=radius, S=angle
    ADD        phase, frequency         ' Increment phase

    ' Do other audio processing while waiting
    RDLONG     volume, ##volume_addr

    GETQY      sample                   ' Get sine value
    SAR        sample, #16               ' Scale to 16-bit
    MULS       sample, volume            ' Apply volume

    ' Output to DAC
    WYPIN      sample, #AUDIO_PIN

    ' Wait for sample period (48kHz)
    WAITX      ##4166                   ' 200MHz / 48kHz

    JMP        #sample_loop

```

Common CORDIC Gotchas

Before you pull your hair out debugging, know these:

1. **One result at a time** - Each cog has its own CORDIC, but starting a new operation before retrieving your result overwrites it!
2. **55 clocks is exact** - Not 54, not 56. Always exactly 55 clocks from operation start to result ready.
3. **Don't forget SETQ** - For two-operand operations (QROTATE with X,Y), you must load Y into Q first.
4. **Results are scaled** - When rotating a vector of length \$7FFF_FFFF, the X/Y results come back scaled so that \$7FFF_FFFF represents 1.0 (full-scale signed).

5. **Angles wrap naturally** - Adding \$1_0000_0000 to an angle is the same as adding 0. Use this!

What About QLOG, QEXP?

Don't worry, we won't leave logarithms and exponentials behind. CORDIC handles those too:

```
' Base-2 logarithm
  QLOG    value
  GETQX   result      ' log2(value) in 5.27 fixed point

' Base-2 exponential (2^x)
  QEXP    value
  GETQX   result      ' 2^value
```

These are less commonly used but incredibly powerful for DSP and scientific calculations.

Interlude

Jack Volder's Gift to Computing

In 1959, Jack Volder was working on navigation computers for aircraft. He needed to calculate trigonometric functions, but the computers of the day couldn't handle the complex math quickly enough.

His insight? Any angle can be decomposed into a sequence of smaller, fixed angles. By choosing these angles cleverly (arctan of powers of 2), he could rotate vectors using only shifts and adds - no multiplication needed!

The B-58 bomber's navigation computer was the first to use CORDIC. Today, it's in your P2, calculating sines and cosines faster than those room-sized computers could add two numbers.

From military navigation to your LED projects - quite a journey for an algorithm!

What We've Learned

Let's celebrate your new CORDIC powers:

- ✓ Understood CORDIC's rotate and vector operations
- ✓ Generated sine and cosine values
- ✓ Calculated distances and angles
- ✓ Learned the pipeline technique for speed
- ✓ Created rotating graphics
- ✓ Built an audio tone generator

That's serious mathematical muscle!

Coming Up Next

Chapter 8 brings us back to Earth with “Basic I/O” - the fundamental pin operations that make the real world connection. We'll save smart pins for another manual and focus on the essentials: making pins go high and low, reading buttons, and basic timing.

But first, take a moment to appreciate what you just learned. CORDIC is unique to the Propeller 2 - most microcontrollers would need extensive software libraries to do what you just did in three instructions!

Have Fun! And remember - with CORDIC, you're not just calculating trigonometry, you're doing it at hardware speed. That's magical!

Chapter 8: Basic I/O

Making the real world connection

The Hook: One Pin, Three Instructions, Infinite Possibilities

Watch this:

```
' Complete button-and-LED program
.loop  TESTP   #BUTTON_PIN wc  ' Read button into C flag
      if_c  DRVH   #LED_PIN      ' If pressed, LED on
      if_nc DRVL   #LED_PIN      ' If not pressed, LED off
      JMP    #.loop              ' Repeat forever
```

Four lines. Complete input/output program. No configuration registers, no data direction setup, no port manipulation. Just pure, simple I/O.

But wait, let me show you the same thing with even more elegance:

```
' Even simpler - button controls LED directly
.loop  TESTP   #BUTTON_PIN wc  ' Read button
      DRVC     #LED_PIN        ' Drive LED from C flag!
      JMP      #.loop
```

Three lines! The **DRVC** instruction drives the pin to match the C flag. Input becomes output. Simple becomes simpler.

Understanding P2 Pins

Let's unpack what makes those three lines so short. Every P2 pin is bidirectional and incredibly capable. Unlike older microcontrollers where you set data direction registers (remember TRIS bits? DDRA? Yeah, we don't miss those either), P2 pins change direction on the fly based on the instruction you use.

Here's the mental model:

- **Output instructions** (DRVH/DRVL/DRVNOT) automatically drive the pin (direction becomes output)
- **Float instructions** (FLTL/FLTH) make the pin high-impedance (direction becomes input)
- **Reading a pin** (TESTP) works regardless of its direction
- No setup required!

Digital Output: Making Things Happen

The Fundamental Four

```
DRVH    #56      ' Drive pin 56 HIGH (3.3V)
DRVL    #56      ' Drive pin 56 LOW (0V)
DRVNOT   #56      ' Toggle pin 56
FLTTL   #56      ' Float pin 56 (high-Z)
```

That's it. Four instructions, 90% of your output needs covered. We'll meet a few specialized cousins below, but if you only remember these four, you'll get a lot done.

Conditional Driving

Here's where P2 gets clever:

```
DRVC     #56      ' Drive pin to match C flag
DRVNC    #56      ' Drive pin to NOT C flag
DRVZ     #56      ' Drive pin to match Z flag
DRVNZ    #56      ' Drive pin to NOT Z flag
```

And the really clever one:

```
DRVNOT   #56 wcz    ' Toggle pin AND read old state to C
' C now contains what the pin WAS before toggling
```

Random and Pattern Outputs

```
DRVRND   #56      ' Drive pin to a random level (hardware PRNG)
OUTL     #56      ' Set OUT bit low (dir unchanged)
OUTH     #56      ' Set OUT bit high (dir unchanged)
```

Digital Input: Reading the World

Basic Pin Reading

```

        TESTP    #BUTTON_PIN wc ' Read pin into C flag
if_c JMP        #pressed        ' Branch if high
if_nc JMP        #not_pressed    ' Branch if low

```

Or read into Z flag for zero/non-zero testing:

```

        TESTP    #SENSOR_PIN wz ' Read pin into Z flag
if_z JMP        #sensor_low      ' Jump if pin low (Z=1 when pin=0)
if_nz JMP       #sensor_high     ' Jump if pin is high

```

Reading Multiple Pins

```

' Read 8 pins at once (pins 0-7)
MOV     mask, #$FF      ' Pins 0-7
TESTB   ina, #0 wc      ' Test pin 0
RCL     result, #1      ' Rotate C into result
TESTB   ina, #1 wc      ' Test pin 1
RCL     result, #1
' ... continue for all 8 pins

```

Pin Timing: When Things Happen

Waiting for Pin Changes

```

' Wait for pin to go high
wait_high
        TESTP    #SIGNAL_PIN wc
if_nc JMP        #wait_high

' Wait for pin to go low
wait_low
        TESTP    #SIGNAL_PIN wc
if_c JMP         #wait_low

```

But there's a better way - hardware-assisted waiting with Smart Events:

```

WAITSE1          ' Wait for event 1
WAITSE2          ' Wait for event 2
' Configure events to watch pins - super efficient!

```

Real-World Example: Button Debouncing

Mechanical buttons bounce. Here's how to handle it:

```

' Debounced button reader
read_button
    MOV     debounce, #0

check_button
    TESTP   #BUTTON_PIN wc
    if_c ADD     debounce, #1    ' Count high readings
    if_nc MOV    debounce, #0    ' Reset on any low

    CMP     debounce, #10 wcz   ' Need 10 consecutive highs
    if_ae JMP    #button_confirmed

    WAITX    ##200_000          ' Wait 1ms at 200MHz
    JMP     #check_button

button_confirmed
    ' Button definitely pressed
    DRVH     #LED_PIN

```

Bit-Banged Serial (The Basics)

Yes, you'll usually reach for smart pins for serial — but it's worth seeing how to do it the hard way once, just so you appreciate what smart pins are doing for you. Here's how:

```

' Bit-bang serial transmit at 115200 baud
tx_byte
    OR      data, ##$100        ' Add stop bit
    SHL     data, #1            ' Add start bit (0)
    MOV     bits, #10           ' 1 start + 8 data + 1 stop

    GETCT   time                ' Get current time

tx_loop
    SHR     data, #1 wc         ' Get next bit into C
    DRVC    #TX_PIN            ' Output bit

    ADDCT1  time, bit_time      ' Next bit time
    WAITCT1                ' Wait for it

```

continues on next page →

```

        DJNZ    bits, #tx_loop
        RET

bit_time LONG    100_000_000 / 115200 ' Clock cycles per bit

```

Your Turn: I/O Experiments

Your Turn

Your Turn: Create a light chaser

Starting code:

```

        ORG     0

        MOV     pattern, #1      ' Start with one LED

.loop   MOV     pins, pattern    ' Your code here
        ' Make pattern rotate through pins 56-63
        ' Add delay between changes
        ' Wrap around at the end, then jmp #.loop

```

Goal: Create a rotating light pattern on LEDs Hint: Use **SHL** and check for overflow Success Check: Single lit LED rotating through all positions

Your Turn

Your Turn: Reaction timer

Starting code:

```

        ORG     0

        ' Turn on LED after random delay
        GETRND  delay
        AND     delay, ##$3FFF_FFFF ' Limit range
        WAITX   delay
        DRVH    #LED_PIN

        GETCT   start_time
        ' Your code: wait for button press
        ' Calculate reaction time

```

Goal: Measure reaction time between LED and button press Hint: Use **GETCT** after button detection Success Check: Time measured in clock cycles

The Medicine Cabinet

Feeling overwhelmed by all these pin operations? Here's the simplified prescription:

Just need something working? Remember these patterns:

Output pattern:

```
DRVH    #PIN    ' Make it high
DRVL    #PIN    ' Make it low
DRVNOT  #PIN    ' Toggle it
```

Input pattern:

```
TESTP   #PIN wc ' Read it
if_c    JMP    #high ' It's high
if_nc   JMP    #low  ' It's low
```

Timed pattern:

```
.loop   DRVNOT  #LED
        WAITX   ##50_000_000
        JMP     #.loop
```

That's 80% of all I/O right there!

Advanced Pin Control

Pin Groups

You can control multiple pins at once:

```
DRVH    #LED_BASE addpins 3 ' Drive 4 pins high (base+3)
DRVL    #LED_BASE addpins 7 ' Drive 8 pins low
```

Direct Pin Manipulation

For when you need absolute control:

```
MOV     outa, pattern    ' Set output register directly
MOV     dira, ##$FF      ' Set direction reg (rare in P2!)
```

But honestly? You'll rarely need these. The individual pin instructions are cleaner and clearer.

Common I/O Gotchas

Before you pull your hair out wondering why a pin “won’t work,” save yourself debugging time and skim these:

1. **Pin numbers are 0-63** - Not port.bit notation like other MCUs
2. **No pullup/pulldown by default** - Use external resistors or configure smart pin modes (advanced topic)
3. **Pins float on reset** - All pins start as inputs (floating)
4. **Reading output pins** - You CAN read a pin you’re driving (reads the actual pin state)
5. **3.3V logic levels** - P2 is 3.3V, not 5V tolerant!

Timing Is Everything

Now here’s something we’ll keep coming back to: P2 I/O is deterministic. When you execute:

```
DRVH    #56
DRVL    #57
```

Pin 56 goes high and pin 57 goes low at EXACTLY the same clock cycle. No skew, no uncertainty. Uff! Try doing that on an interrupt-driven MCU. This determinism is what makes P2 perfect for precise timing applications.

Real-World Example: Servo Control

Even without smart pins, controlling a servo is easy:

```
' Standard servo control (1-2ms pulse every 20ms)
servo_control
    MOV    position, ##300_000    ' 1.5ms = center at 200MHz

servo_loop
    DRVH    #SERVO_PIN
    WAITX   position              ' 1-2ms high pulse
    DRVL    #SERVO_PIN
    WAITX   ##4_000_000           ' Rest of 20ms at 200MHz

    ' Adjust position as needed
    RDLONG  position, ##position_addr
    FGE     position, ##200_000    ' Limit to 1ms min
```

continues on next page →


```
FLE    position, ##400_000    ' Limit to 2ms max  
  
JMP    #servo_loop
```

The Power of Simple

Here's something beautiful about P2's I/O philosophy: it's transparent. Unlike microcontrollers with complex GPIO configurations, port multiplexing, and alternate functions, P2 pins just... work.

Want an output? Drive it. Want an input? Read it. Want it to float? Float it.

No setup, no configuration, no confusion.

What We've Learned

Look at your new I/O skills:

- ✓ Understood P2's automatic pin direction
- ✓ Mastered the four fundamental output instructions
- ✓ Learned pin reading and conditional testing
- ✓ Created debounced inputs
- ✓ Built bit-banged serial
- ✓ Discovered deterministic timing

A Note About Smart Pins

You might wonder - if basic I/O is this simple, why do we need smart pins?

Well, while you CAN bit-bang serial at 115200 baud, or generate PWM, or measure frequencies using the techniques in this chapter, smart pins do all of this in hardware, freeing your cog for more important work.

For smart pin details: See the dedicated "P2 Smart Pins Manual" which covers all 32 modes, from simple PWM to complex protocol generation. Smart pins deserve their own complete treatment!

Coming Up Next

Chapter 9 takes us into "Streaming Data" - the P2's incredible FIFO system that can move megabytes of data without breaking a sweat. We'll see how to stream video, audio, and massive data blocks at maximum speed.

Have Fun! Remember, every embedded system ultimately comes down to pins going high and low. You've just mastered the fundamentals that everything else builds upon!

Chapter 9: Streaming Data

Moving mountains of data without breaking a sweat

The Hook: 4KB in 4 Instructions

Watch this data transfer magic:

```
' Copy 1000 longs (4KB) at maximum speed
  SETQ    ##1000-1      ' Setup for 1000 longs
  RDLONG  buffer, source ' Read them all!
  SETQ    ##1000-1      ' Setup for 1000 longs
  WRLONG  buffer, dest   ' Write them all!
  ' 4KB moved in microseconds!
```

Four instructions. Four kilobytes. Faster than DMA on most processors. And we're just getting started...

Block Transfers: The Power Move

The **SETQ** instruction is your gateway to block transfers:

```
' Basic block read
  SETQ    #16-1          ' Transfer 16 longs (minus 1!)
  RDLONG  buffer, hubaddr ' Reads 16 consecutive longs

' Basic block write
  SETQ    #16-1          ' Transfer 16 longs
  WRLONG  buffer, hubaddr ' Writes 16 consecutive longs
```

Here's the trick: **SETQ** tells the next hub instruction how many longs to transfer. The “-1” is because it's a count from 0 (yes, we'll trip over that off-by-one at least once — everyone does).

The FIFO: Your Streaming Pipeline

The FIFO (First In, First Out) is P2's streaming engine. Think of it as a conveyor belt between hub memory and your cog:

```
' Start FIFO reading
  RDFAST  #0, ##data_start ' Start fast read at data_start
```

continues on next page →

→ continued from previous page

```

' Now read data at maximum speed
stream_loop
    RFLONG    value                ' Read from FIFO (no waiting!)
    ' Process value here
    ADD      accumulator, value
    DJNZ     count, #stream_loop

' The FIFO keeps feeding data automatically

```

The beauty? The FIFO reads ahead automatically. While you're processing one value, it's already fetching the next. No hub timing slots to worry about!

Writing Through the FIFO

Reading was nice — writing is symmetric. We'll use **WRFAST** instead of **RDFAST**, and **WFLONG** instead of **RFLONG**, and that's it:

```

' Start FIFO writing
    WRFAST   #0, ##dest_buffer

' Stream data out
write_loop
    ' Generate or process data
    MOV      value, calculation
    WFLONG   value                ' Write to FIFO
    DJNZ     count, #write_loop

' Data streams to hub automatically

```

Real-World Example: Screen Buffer Clear

Let's clear a 320x240x4 byte screen buffer (~307KB - fits in hub!):

```

' Ultra-fast screen clear
clear_screen
    MOV      color, ##$00_00_00_00    ' Black (4 bytes per pixel)
    MOV      pixels, ##320*240        ' 76,800 pixels

    WRFAST   #0, ##screen_buffer      ' Start FIFO write

clear_loop
    WFLONG   color                    ' Write 4-byte pixel
    DJNZ     pixels, #clear_loop

' 307KB cleared at maximum hub speed!
' Note: Hub RAM is 512KB - plan buffer sizes accordingly

```

Streaming with the Streamer

The streamer is different from the FIFO - it's a dedicated DMA engine that can move data between hub memory and pins:

```
' Configure streamer for video output
  SETXFRQ ##PIXEL_FREQ      ' Set the pixel (NCO) output rate

' Start streaming video data to pins
  XINIT  ##STREAM_CMD, #0    ' Start streamer
  ' Data flows from hub to pins automatically!
```

FIFO and Cog Execution

Here's something amazing - you can execute code from hub through the FIFO:

```
' Execute large program from hub
  ORGH   $1000              ' Code in hub memory

hub_code
  ' This code is in hub but executes like it's in cog
  ADD    x, y
  SUB    a, b
  ' Can be megabytes of code!
```

When you call or jump to hub code, the FIFO automatically feeds instructions to the cog. It's like having unlimited code space!

The Medicine Cabinet

Feeling overwhelmed by all this streaming? Here's your prescription:

Just need to move data? Use these simple patterns:

Block read pattern:

```
SETQ    #SIZE-1
RDLONG  buffer, source
```

Block write pattern:

```
SETQ    #SIZE-1
WRLONG  buffer, dest
```

FIFO read pattern:

```
        RDFAST  #0, ##source
.loop    RFLONG  value
        ' Process value
        DJNZ    count, #.loop
```

That's 90% of streaming right there!

Advanced Streaming Techniques

Circular Buffers with FIFO

```
' Circular buffer reading
    RDFAST  ##BUF_BLOCKS, ##buffer ' D[13:0]=block count; auto-wraps

circular_loop
    RFLONG  value                  ' Read from FIFO
    ' Process value
    ' FIFO automatically wraps at buffer end!
    JMP     #circular_loop
```

Processing Pipeline with FIFO

```
' Read data with FIFO, process, write via PTRB
  RDFAST  #0, ##source      ' Set up FIFO for reading
  MOV     dest_ptr, ##dest

pipeline
  RFLONG  input              ' Get next input from FIFO

  ' Scale using 16x16 multiply (result in input)
  MUL     input, #SCALE_FACTOR

  WRLONG  input, dest_ptr    ' Write result via PTRB
  ADD     dest_ptr, #4
  DJNZ    count, #pipeline
```

Note: FIFO can only read OR write at a time, not both. Use PTRB/PTRA for the other direction.

Your Turn: Streaming Experiments

Your Turn

Your Turn: Fast memory fill

Starting code:

```
ORG      0

MOV      pattern, ##$DEADBEEF
MOV      dest, ##$1000
MOV      count, #256

' Your code here: Fill 256 longs with pattern
' Use SETQ and WRLONG
```

Goal: Fill memory with pattern using block transfer Hint: You'll need **SETQ** #255 (not #256) Success Check: Memory filled in one operation

Your Turn

Your Turn: Data filter pipeline

Starting code:

continues on next page →

```

    ORG      0

    RFAST   #0, ##input_data      ' FIFO for reading
    MOV     ptr, ##output_data    ' PTR for writing
    MOV     count, #100

filter_loop
    RFLONG  value                  ' Read from FIFO
    ' Your code: Simple filter
    ' Maybe average with previous value?
    WRLONG  result, ptr++          ' Write via PTR
    DJNZ    count, #filter_loop

```

Goal: Process streaming data through simple filter Hint: Keep previous value in a register

Success Check: Output is filtered version of input

Common Streaming Gotchas

Before you pull your hair out wondering why your transfer is one long short, or why your FIFO won't cooperate, skim these:

1. **SETQ uses count-1** - For 16 longs, use SETQ #15, not SETQ #16
2. **FIFO is shared per cog** - Can't use FIFO for both code execution and data streaming simultaneously
3. **Write synchronization** - **WRFAST** writes complete in the background. To force a flush, issue the next **RFAST**/WRFAST with D[31]=0 (it waits for the prior **WRFAST** to finish) rather than relying on a fixed delay
4. **Hub alignment** - Block transfers work best with long-aligned addresses
5. **FIFO depth** - The FIFO holds (cogs+11) = 19 longs. It refills automatically, so you rarely outrun it.

Performance Numbers

Let's talk speed:

- **Block transfer:** Up to 1 long per clock (at 200MHz = 800MB/s!)
- **FIFO streaming:** Up to 1 long per clock sustained
- **Random hub access:** 9-16 clocks per access
- **Streamer to pins:** Up to sysclock/1 rate

Uff! That’s seriously fast. Most microcontrollers need dedicated DMA controllers, peripheral co-processors, and a stack of config registers to achieve what P2 does with two instructions. You’re not paying for that DMA controller — it’s already in the silicon.

Real-World Example: Audio Buffer

Stream audio samples through processing:

```
' Audio processing pipeline
audio_process
    RDFAST    #0, ##input_buffer    ' FIFO for reading input
    MOV      ptra, ##output_buffer  ' PTR for writing output
    MOV      samples, ##BUFFER_SIZE

process_loop
    RFLONG    left_sample            ' Read left from FIFO
    RFLONG    right_sample           ' Read right from FIFO

    ' Apply simple low-pass filter
    ADD      left_filtered, left_sample
    SHR      left_filtered, #1      ' Average with previous

    ADD      right_filtered, right_sample
    SHR      right_filtered, #1

    ' Apply volume
    Muls      left_filtered, volume
    Muls      right_filtered, volume

    ' Output processed samples via PTR
    WRLONG    left_filtered, ptr++
    WRLONG    right_filtered, ptr++

    DJNZ      samples, #process_loop
```

What We’ve Learned

Your streaming skills now include:

- ✓ Block transfers with **SETQ**
- ✓ FIFO reading and writing
- ✓ Streaming pipeline concepts
- ✓ Circular buffer techniques
- ✓ Parallel processing while streaming
- ✓ Real-world applications

Coming Up Next

Chapter 10 explores “hub execution” - how to break free from the 496-instruction limit and run massive programs directly from hub memory. It’s like having your cake and eating it too!

Have Fun! Remember, streaming is about throughput, not just speed. It’s the difference between carrying one brick at a time and using a wheelbarrow!

Chapter 10: Hub Execution

Breaking free from the 496-instruction limit

The Hook: Unlimited Code Space

Remember fretting about fitting your code into 496 cog instructions? Watch this:

```
        ORGH    $400                ' Place code in hub memory

        ' This can be thousands of instructions!
main    MOV     x, #0
        MOV     y, #0
        CALL    #huge_function      ' Can be massive
        CALL    #another_big_one
        CALL    #yet_another
        ' Keep going... no limit!

huge_function
        ' 1000 instructions? No problem!
        ' 10000 instructions? Still fine!
        RET
```

Your code now lives in hub memory's 512KB instead of cog memory's 2KB. That's 256 times more space!

Cog vs Hub Execution: The Trade-offs

Let's be honest about the differences:

Cog Execution (traditional):

- ✓ Fast: exactly 2 clocks per instruction
- ✓ Deterministic: perfect for real-time
- ✗ Limited: only 496 instructions
- ✓ Self-contained: runs independently

Hub Execution (the new way):

- ✓ Fast sequential: 2 clocks per instruction (same as cog-exec, thanks to the 19-stage FIFO prefetch)
- ✗ Slower on branches: minimum 13 clocks per branch (the FIFO refill cost; +1 if target isn't long-aligned)
- ✓ Unlimited: 512KB of code space!

- ✓ Flexible: can call cog routines

The beauty? You can mix both in the same program! Sequential code in hub runs at full speed — only branches show the hub-execution penalty.

How Hub Execution Works

Let's peek behind the curtain. When the processor encounters a jump or call to a hub address ($\geq \$400$), it automatically switches to hub execution mode. The FIFO starts streaming instructions from hub memory:

```

        ORG      0                ' Start in cog

cog_code
    ' This runs from cog RAM
    CALL    #hub_function    ' Call into hub
    ' Back in cog mode here

        ORGH     $1000           ' Switch to hub addresses

hub_function
    ' This runs from hub RAM via FIFO
    ' Can be huge!
    RET                    ' Returns to cog code

```

The magic happens automatically. No mode switching instructions needed!

Real-World Example: Menu System

Here's something that would never fit in cog RAM:

```

        ORGH     $2000

menu_system
    CALL    #draw_menu_frame
    CALL    #display_options
    CALL    #get_user_input
    CALL    #process_selection

    CMP     selection, #1 wcz
if_e CALL    #option_1_handler
    CMP     selection, #2 wcz
if_e CALL    #option_2_handler
    CMP     selection, #3 wcz
if_e CALL    #option_3_handler
    ' ... many more options

```

continues on next page →

↪ continued from previous page

```

        JMP      #menu_system

draw_menu_frame
    ' 200 instructions for fancy graphics
    RET

display_options
    ' 300 instructions for text rendering
    RET

option_1_handler
    ' 500 instructions for configuration
    RET

    ' Thousands of instructions total - no problem!

```

The Hub Execution FIFO

You met the FIFO in the previous chapter as a data-streaming engine. Same hardware, different job here: it reads ahead, keeping a buffer of upcoming *instructions* ready for the cog. Think of it as a moving sidewalk for your code:

```

    ' The FIFO maintains performance by reading ahead
hub_loop
    ADD     x, y           ' FIFO has next instructions ready
    SUB     a, b           ' No waiting for hub access
    MUL     c, d           ' Instructions stream smoothly
    ' FIFO automatically refills as needed

```

This read-ahead behavior is the whole reason sequential hub code matches cog-exec speed — the FIFO is doing the waiting for you, in parallel with the cog running instructions it already prefetched.

Mixing Cog and Hub Code

Here's the real power - combining both modes:

```

        ORG      0

    ' Critical timing code in cog
critical_loop
    TESTP   #TRIGGER_PIN wz    ' Test trigger pin
    if_nz   JMP      #critical_loop    ' Loop until triggered
    DRVH    #CRITICAL_PIN      ' Immediate response!
    CALL    #hub_process    ' Do complex processing
    JMP     #critical_loop

```

continues on next page →

→ continued from previous page

```

        ORGH    $4000

' Complex processing in hub
hub_process
    ' Hundreds of instructions for data analysis
    ' Not time-critical, so hub execution is fine
    RET

```

Time-critical code stays in cog RAM for deterministic timing. Complex code lives in hub RAM for space.

Your Turn: Hub Execution Experiments

Your Turn

Your Turn: Build a simple calculator

Starting code:

```

        ORG     0
        JMP     #calculator    ' Jump to hub code

        ORGH    $1000
calculator
    ' Your code here:
    ' 1. Display menu
    ' 2. Get operation choice
    ' 3. Get two numbers
    ' 4. Call appropriate function
    ' 5. Display result

add_function
    ' Addition code
    RET

subtract_function
    ' Subtraction code
    RET

' Add more functions - no size limit!

```

Goal: Create a multi-function calculator Hint: Each function can be as complex as needed

Success Check: Multiple operations working

The Medicine Cabinet

Overwhelmed by execution modes? Here's the simple version:

Keep it simple:

1. **Small, time-critical code** → Put in cog (org 0)
2. **Large, complex code** → Put in hub (orgh \$400+)
3. **Don't overthink it** → The processor handles the switch

Basic pattern:

```

        ORG      0
        JMP      #main      ' Jump to hub

        ORGH     $400
main     ' Your big program here

```

That's it. Let the processor worry about the details!

Advanced Hub Execution

Long Jumps and Calls

Hub addresses need 20 bits, so jumping far requires special handling:

```

' Jump to distant hub code
        JMP      #\far_away    ' \ forces a 32-bit absolute address

        ORGH     $40000        ' Far away in hub
far_away
        ' Code here

```

Hub Data Access from Hub Code

When executing from hub, you can still access hub data:

```

        ORGH     $1000

hub_code
        RDLONG   value, ##hub_data  ' Read hub data
        ADD      value, #1
        WRLONG   value, ##hub_data  ' Write back

        ORGH     $8000

```

continues on next page →

```

hub_data
    LONG    $12345678

```

Performance Optimization

To maximize hub execution speed:

```

' Align branch targets to 8-byte boundaries
    ALIGNL                ' Align to long boundary
loop_start
    ' Loop code here
    DJNZ    count, #loop_start

' Keep critical loops small
' Consider moving inner loops to cog RAM

```

Common Hub Execution Gotchas

Before you cram everything into hub and call it a day, know these:

1. **Speed variation** - Don't use hub execution for precise timing
2. **FIFO conflicts** - Can't stream data while executing from hub
3. **Address confusion** - Remember: <\$200 is cog, \$200-\$3FF is LUT, ≥\$400 is hub
4. **Stack depth** - Still limited to 8-level hardware stack
5. **Relative jumps** - Work differently in hub mode

Real-World Example: Command Parser

```

        ORGH    $2000

command_parser
    CALL    #get_command_line
    CALL    #tokenize

    ' Compare against commands
    MOV     ptrb, #command_string
    MOV     ptrb, ##cmd_help
    CALL    #string_compare
    if_z    JMP    #help_command

    MOV     ptrb, ##cmd_run
    CALL    #string_compare
    if_z    JMP    #run_command

```

continues on next page →


```
' Many more commands...

help_command
    ' 500 instructions of help text display
    RET

run_command
    ' 1000 instructions of program execution
    RET

string_compare
    ' 50 instructions of string comparison
    RET

' Thousands of instructions total
' Would need multiple cogs without hub execution!
```

When to Use Hub Execution

Perfect for:

- User interfaces and menus
- Command processors
- Complex algorithms
- String manipulation
- Protocol handlers
- Error handling and recovery

Avoid for:

- Interrupt handlers (if you use them)
- Precise timing loops
- Bit-banged protocols
- Real-time control loops

What We've Learned

You've mastered hub execution:

- ✓ Understanding cog vs hub trade-offs
- ✓ Automatic mode switching
- ✓ Mixing cog and hub code
- ✓ FIFO streaming of instructions
- ✓ When to use each mode

- ✓ Real-world applications

Coming Up Next

Chapter 11 tackles the controversial topic: “Why No Interrupts?” We’ll explore why the Propeller philosophy says you don’t need them, and why that’s actually a good thing!

Have Fun! Hub execution is like having a sports car that can also carry furniture - you get both speed and capacity when you need them!

Chapter 11: Why No Interrupts?

The most controversial P2 feature explained

The Hook: Interrupts Without Interrupts

Here's a traditional interrupt-driven button handler:

```
' Traditional approach (not P2!) – pseudocode for the interrupt-driven style
ISR(BUTTON_INTERRUPT)
    ' Interrupt service routine
    buttonPressed = true
    ' Return to interrupted code
END_ISR
```

And here's the P2 way:

```
' Dedicated cog watching button
button_watcher
    TESTP    #BUTTON_PIN wc
    if_c WRLONG  ##1, ##button_flag
    JMP      #button_watcher

' Main cog doing important work
main_code
    ' Never interrupted!
    ' Checks button_flag when convenient
```

No interrupt latency. No context switching. No priority inversion. No critical sections. Just clean, deterministic, parallel processing.

The Interrupt Problem

Let me tell you a story. You're concentrating on a complex problem when someone taps your shoulder. You stop, handle their request, then try to remember where you were. Now imagine this happening randomly, unpredictably, dozens of times per second.

That's interrupts.

Traditional processors need interrupts because they only have one processor. Something important happens? Stop everything and handle it! But this causes:

- **Latency:** Time to save context and jump to handler
- **Jitter:** Variable response time depending on what was interrupted

- **Priority inversion:** Low-priority task blocks high-priority
- **Race conditions:** Shared data access problems
- **Debugging nightmares:** Non-reproducible timing bugs

The Propeller Solution

Eight processors. No sharing required.

```
' cog 0: Main application
main_app
    ' Complex calculations
    ' Never interrupted
    RDLONG    command, ##mailbox wz
    if_nz CALL    #process_command
    JMP      #main_app

' cog 1: Serial port handler
serial_handler
    ' Continuously monitors serial
    TESTP    #RX_PIN wc
    if_c CALL    #receive_byte
    JMP      #serial_handler

' cog 2: Motor control
motor_control
    ' Precise timing loops
    ' Never disrupted
    WAITX    motor_time
    DRVNOT    #STEP_PIN
    JMP      #motor_control

' cog 3: Sensor monitor
sensor_monitor
    ' Watches multiple sensors
    ' Responds instantly
    ' ... and so on
```

Each cog does one thing perfectly. No interruptions. No conflicts. Just pure, focused execution.

Real-World Example: Perfect Servo Control

With interrupts, servo pulses jitter. With dedicated cogs, they're perfect:

```
' cog dedicated to servo control
servo_cog
    GETCT    pulse_time
```

continues on next page →

↪ continued from previous page

```

servo_loop
    ' Generate 8 servo pulses simultaneously
    MOV     servo_mask, #0xFF      ' 8 servos
    OR      outa, servo_mask      ' All high

    MOV     index, #0

check_servos
    RDLONG  width, ptr[index]      ' Get pulse width
    ADDCT1  pulse_time, width      ' Set compare time

    WAITCT1                                ' Wait for exact time
    BITL    outa, index            ' Turn off this servo

    INCMOD  index, #7 wc           ' C set when index wraps 7->0
    if_nc   JMP     #check_servos  ' Loop until all 8 servos done

    ' Wait for 20ms frame
    WAITX   ##4_000_000
    JMP     #servo_loop

    ' Result: 8 servos with ZERO jitter!

```

Try that with interrupts. I'll wait. Actually, I won't - it's impossible to achieve this precision with interrupts.

“But P2 HAS Interrupts!”

Yes, it does. And you probably shouldn't use them.

Well, let me be more nuanced. P2 has interrupts for those rare cases where you absolutely need them:

```

    ' Setting up an interrupt (not recommended!)
    SETSE1  #%001<<6 + PANIC_BUTTON  ' SE1 triggers when pin goes high
    SETINT1 #EVENT_SE1                ' Enable INT1 on SE1 event

int1_handler
    ' Interrupt code here
    RETI1

```

When might you use them?

- Porting legacy code that requires interrupts
- Ultra-low-power designs where cogs must sleep
- Theoretical minimum latency response (but dedicated cog is usually faster!)

Uff! Even writing interrupt code feels wrong on a Propeller!

The Medicine Cabinet

Still thinking you need interrupts? Here's your medicine:

Think you need an interrupt for...

Fast response?

```
' Dedicated cog responds in ~4 clocks
watcher
    TESTP    #INPUT_PIN wz          ' Test pin state
    if_nz JMP    #watcher          ' Loop until pin high
    DRVH     #RESPONSE_PIN          ' Instant response!
```

Multiple events?

```
' One cog watches everything
monitor
    TEST     sensors, #SENSOR1 wz
    if_nz CALL #handle_sensor1
    TEST     sensors, #SENSOR2 wz
    if_nz CALL #handle_sensor2
    ' Check all sensors every loop
```

Periodic tasks?

```
' Perfect timing without interrupts
    GETCT    next_time
.loop   ADDCT1 next_time, ##PERIOD
    WAITCT1                      ' Exact timing
    CALL     #periodic_task
    JMP      #.loop
```

See? No interrupts needed!

The Event System: Better Than Interrupts

P2 has something better than interrupts - events:

```
' Configure event to watch pin
    SETSE1   #%001<<6 + BUTTON_PIN  ' Rising edge event

' Main code runs normally
main_loop
    ' Do work...
    POLLSE1  wc                      ' Check if event occurred
```

continues on next page →

```
if_c CALL    #handle_button  ' Handle when convenient
      ' Continue work...
      JMP     #main_loop
```

Events are like interrupts that wait politely for you to check them. No rudeness!

Interrupt Horror Stories

Let me share why we avoid interrupts:

Story 1: The Jittery Display

Approach	Problem	Result
With Interrupts	Display updates interrupted by serial	Visible glitches, tearing, inconsistent timing
With Cogs	Display Cog runs uninterrupted	Perfect, smooth, glitch-free display

Story 2: The Missed Pulse

Approach	Problem	Result
With Interrupts	Motor step interrupted by sensor read	Missed step, motor stalls, position lost
With Cogs	Motor Cog never misses a beat	Perfect positioning, no lost steps

Story 3: The Debugging Nightmare

Approach	Problem	Result
With Interrupts	Bug only appears under specific timing	Days of debugging, hair loss, coffee overdose
With Cogs	Deterministic timing, reproducible behavior	Bug found in minutes, sanity preserved

Your Turn: Cog vs Interrupt Challenge

Your Turn

Your Turn: Build a reaction timer without interrupts

Starting code:

```

        ORG      0

' cog 0: Main game logic
    SETQ      ##button_flag      ' PTR for new cog
    COGINIT   #1, @button_watcher

game_loop
    ' Random delay
    GETRND    delay
    WAITX     delay

    DRVH      #LED_PIN
    WRLONG    #0, ##button_flag
    GETCT     start_time

    ' Wait for button (no interrupt!)
wait_press
    RDLONG    pressed, ##button_flag wz
    if_z      JMP      #wait_press

    GETCT     end_time
    SUB       end_time, start_time
    ' Display reaction time

' cog 1: Button watcher
    ORGH      $400
button_watcher
    ' Your code here

```

Goal: Implement button watcher cog Hint: Continuously monitor and set flag Success

Check: Perfect timing without interrupts

The Philosophy Deep Dive

The Propeller philosophy is about **determinism over responsiveness**.

Traditional processors optimize for average-case performance:

- Interrupts handle rare events
- Most code runs uninterrupted

- When events happen, everything stops

Propeller optimizes for worst-case determinism:

- Every cog runs predictably
- No surprises, ever
- Timing is guaranteed

It's the difference between a talented soloist who might miss a note and an orchestra where everyone plays their part perfectly.

When Interrupts Actually Make Sense

I'll admit it - there are rare cases where interrupts are appropriate:

1. **Power-critical applications** where cogs must sleep
2. **Legacy code ports** that fundamentally require interrupts
3. **Single-cog designs** (but why waste the P2's power?)

But in 15 years of Propeller programming, I've needed interrupts exactly... never.

Common “But What About...” Questions

Q: “But what about interrupt priority?” A: Cogs don't have priority. They're all equal. Design your system accordingly.

Q: “How do I handle critical events?” A: Dedicate a cog to critical events. It will respond faster than any interrupt.

Q: “Isn't dedicating a whole cog wasteful?” A: You have eight! And a focused cog is simpler than interrupt-riddled code.

Q: “What about power consumption?” A: Use WAITSE/WAITCT for low-power waiting. Cog sleeps until event.

What We've Learned

You now understand the Propeller way:

- ✓ Why interrupts cause problems
- ✓ How cogs eliminate interrupt need
- ✓ event system as polite alternative
- ✓ Real-world benefits of no interrupts

- ✓ Rare cases where interrupts might be used
- ✓ The philosophy of determinism

Coming Up Next

Chapter 12 shows you “Optimization Mastery” - how to make your PASM2 code blazingly fast. We’ll explore the pipeline, instruction pairing, and timing tricks that squeeze every drop of performance from the P2.

Have Fun! And remember - in a world of interruptions, be a cog: focused, deterministic, and uninterruptible!

Chapter 12: Optimization Mastery

Making the fast faster

The Hook: Double Your Speed with One Change

Let me show you a loop that looks fine — until you realize you're paying for the same thing twice. Watch:

```
' Before optimization: 13 clocks
.loop  RDLONG  value, ptra      ' 9-16 (cog-exec) / 9-26 (hub-exec)
      ADD     value, #1        ' 2 clocks
      WRLONG  value, ptra      ' 3-10 (cog-exec) / 3-20 (hub-exec)
      ADD     ptra, #4         ' 2 clocks
      DJNZ    count, #.loop    ' 2/4 (cog-exec) / 2/13-20 (hub-exec)

' After optimization using PTR expressions:
.loop  RDLONG  value, ptra      ' Read from current address
      ADD     value, #1        ' Process
      WRLONG  value, ptra++    ' Write and increment in one!
      DJNZ    count, #.loop    ' Saved the ADD instruction
```

Almost twice as fast! The secret? Understanding how P2 really works.

Understanding the Pipeline

Before we hunt for clocks to save, let's understand where they go in the first place. P2 uses a **5-stage pipeline** - up to five instructions are in flight at once, each at a different stage. You don't need the stage-by-stage breakdown; what matters is the payoff: when the pipeline stays full, an instruction effectively completes every **2 clocks**.

This is why, while one instruction executes, the next is already being fetched:

```
ADD     x, y      ' Executing while next inst fetches
SUB     a, b      ' Fetching while previous executes
' Perfect overlap = maximum throughput
```

Instruction Timing Basics

Not all instructions are created equal:

```

' 2-clock instructions (most ALU operations)
    ADD    x, y          ' 2 clocks
    MOV    a, b          ' 2 clocks
    AND    c, d          ' 2 clocks

' Variable timing (hub access)
    RDLONG value, hubaddr ' 9-16 clocks (hub slot wait)
    WRLONG value, hubaddr ' 3-10 clocks (variable)

' Long operations (CORDIC)
    QROTATE x, angle      ' 2 clocks to start
    GETQX   result        ' 2 clocks (but wait 55 for result)

' Special cases
    MUL     x, y          ' 2 clocks
    QDIV    x, y          ' 2 clocks to start
    GETQX   result        ' 2 clocks (but wait 55 for result)

```

REP: The Speed Loop

REP creates hardware-accelerated loops with zero overhead:

```

' Traditional loop: overhead per iteration
.loop  ADD    sum, value    ' 2 clocks
        ADD    ptr, #4      ' 2 clocks
        DJNZ   count, #.loop ' 2 or 4 clocks (4 if branch taken)

' REP loop: 0 clocks overhead!
        REP    #2, count    ' Repeat next 2 instructions
        ADD    sum, value    ' 2 clocks
        ADD    ptr, #4      ' 2 clocks = 4 total!

```

That's 33% faster just by using **REP**!

Hub-exec note: **REP** works in hub-exec too, but each iteration executes a hidden jump to loop back — and that hidden jump pays the 13+ clock hub-branch cost per iteration. So a 2-instruction **REP** loop that takes 4 clocks in cog-exec balloons to ~17+ clocks per iteration in hub-exec. For time-critical inner loops, keep **REP** in cog or LUT memory. Hub-exec **REP** works correctly; it just isn't zero-overhead there.

SKIP: Conditional Execution on Steroids

SKIP and **SKIPF** let you conditionally execute patterns of instructions:

```

' Traditional: multiple jumps
        CMP     x, #5 wcz
if_a    JMP     #greater
if_b    JMP     #less
        JMP     #equal

' With SKIP: cancel optional steps per a precomputed pattern
'   bit N (LSB first) skips the Nth instruction after SKIP
        SKIP    config_mask      ' e.g. %010 runs steps 0 and 2, skips step 1
        CALL    #setup_uart      ' Step 0
        CALL    #setup_spi       ' Step 1
        CALL    #setup_timer     ' Step 2
        ' One pattern picks which steps run – no jumps, no stalls!

```

Hub Access Optimization

Hub timing is critical for performance:

```

' Unaligned hub access: variable timing
        RDLONG  v1, ##$1001      ' Not long-aligned, slower

' Aligned hub access: predictable timing
        RDLONG  v1, ##$1000      ' Long-aligned, faster

' Sequential access: maximum speed
        RDLONG  v1, ptr++        ' Hardware manages sequence
        RDLONG  v2, ptr++        ' Optimal hub slot usage
        RDLONG  v3, ptr++        ' Maximum throughput

```

The FIFO Fast Path

For ultimate speed, use the FIFO:

```

' Traditional hub reading: ~6 clocks average per long
.loop   RDLONG  value, ptr++
        ADD     sum, value
        DJNZ    count, #.loop

' FIFO reading: 2 clocks per long!
        RDFAST  #0, ptr          ' Start FIFO
.loop   RFLONG  value            ' 2 clocks, always!
        ADD     sum, value        ' 2 clocks
        DJNZ    count, #.loop    ' 2 clocks
        ' 3x faster for sequential reads!

```

Parallel Operations

CORDIC operations can overlap with other work:

```
' CORDIC overlaps with other instructions
  QMUL    x, y          ' Start 32x32->64 multiply (CORDIC)
  ' 55 clocks to do other work!
  ADD     a, b          ' These execute during CORDIC
  SUB     c, d
  MOV     index, #0
  RDLONG  data, ptr++
  ' ... more work
  GETQX   low_result    ' Get CORDIC result (lower 32 bits)
  GETQY   high_result   ' Get CORDIC result (upper 32 bits)

' QROTATE overlap
  QROTATE x_coord, angle ' Start rotation (D=X, S=angle)
  ' 55 clocks of other work!
  GETQX   new_x         ' Get rotated X
  GETQY   new_y         ' Get rotated Y
```

Note: **MUL**/**MULS** are 2-clock ALU instructions that complete immediately (16x16->32). Use **QMUL** for 32x32->64 with CORDIC overlap.

Real-World Example: Fast Memory Copy

Let's optimize a memory copy routine:

```
' Version 1: Basic (slow)
copy_basic
  RDLONG  temp, source
  WRLONG  temp, dest
  ADD     source, #4
  ADD     dest, #4
  DJNZ    count, #copy_basic
  ' ~13 clocks per long

' Version 2: Better pointers
copy_better
  RDLONG  temp, ptr++
  WRLONG  temp, ptrb++
  DJNZ    count, #copy_better
  ' ~8 clocks per long

' Version 3: Block transfer (ultimate)
copy_ultimate
  SUB     count, #1      ' SETQ needs count-1 (0 = 1 long)
  SETQ    count          ' Setup block transfer
  RDLONG  buffer, source ' Read all at once
```

continues on next page →

→ continued from previous page

```

SETQ    count          ' (count already decremented)
WRLONG  buffer, dest    ' Write all at once
' ~1 clock per long for large blocks (hub maximum)!

```

The Medicine Cabinet

Optimization overwhelming you? Start with these simple improvements:

Three easy wins:

1. **Use PTRA/PTRB** instead of manual pointer math

```

' Slow
      RDLONG x, addr
      ADD    addr, #4

' Fast
      RDLONG x, ptra++

```

2. **Align your data** to long boundaries

```

      ALIGNL          ' Force long alignment
data  LONG    $12345678

```

3. **Use REP** for tight loops (note: hub-address expressions like ptra++ only work with hub-access instructions, so we read first, then add)

```

      REP    @.end, count
      RDLONG val, ptra++    ' Fetch next long from hub
      ADD    sum, val       ' Accumulate
.end

```

Just these three changes often double performance!

Your Turn: Optimization Challenges

Your Turn

Your Turn: Optimize a checksum calculator

Starting code:

continues on next page →

```

' Slow version
checksum_slow
    MOV     sum, #0
    MOV     addr, ##buffer
    MOV     count, #256

.loop    RDBYTE temp, addr
        ADD     sum, temp
        ADD     addr, #1
        DJNZ    count, #.loop

```

Goal: Make it at least 4x faster Hint: Read longs instead of bytes, use FIFO Success Check:
Same checksum, much faster

Advanced Techniques

Instruction Pairing

Some instruction pairs execute specially:

```

' ## syntax handles AUGS automatically
    MOV     x, ##$12345678 ' Assembler generates AUGS + MOV
    ' Same result, cleaner code!

' ALTD + instruction = indirect addressing
    ALTD     index, #array
    MOV     0-0, value      ' Stores to array[index]

```

Pipeline-Aware Coding

Avoid pipeline stalls:

```

' Bad: result needed immediately
    ADD     x, y
    CMP     x, #10 wcz      ' Stall waiting for x

' Good: interleave operations
    ADD     x, y
    MOV     a, b             ' Do something else
    CMP     x, #10 wcz      ' Now x is ready

```


Unrolling Loops

Sometimes removing the loop is faster. (Remember: `ptr++` only works with hub-access instructions like **RDLONG**, so each iteration reads then adds.)

```
' Looped version
    REP     @.end, #4
    RDLONG  val, ptr++
    ADD     sum, val
.end

' Unrolled version (faster for small counts)
    RDLONG  val, ptr++
    ADD     sum, val
    RDLONG  val, ptr++
    ADD     sum, val
    RDLONG  val, ptr++
    ADD     sum, val
    RDLONG  val, ptr++
    ADD     sum, val
```

Common Optimization Gotchas

Before you rewrite everything in **REP** and **SKIP**, a few sanity checks:

1. **Premature optimization** - Get it working first, then optimize
2. **Over-optimizing** - Sometimes clarity is worth 2 clocks
3. **Ignoring the big picture** - Optimize the bottleneck, not everything
4. **Breaking functionality** - Fast but wrong is useless
5. **Forgetting about power** - Faster isn't always better for battery life

Profiling and Measurement

Always measure your optimizations:

```
' Time your code
    GETCT   start_time

    ' Code to measure
    CALL    #function_to_test

    GETCT   end_time
    SUB     end_time, start_time
    ' end_time now contains exact clock cycles
```

Pitfall — CT wraps: **GETCT** reads a 32-bit free-running counter that wraps every ~21.5 seconds at 200 MHz ($2^{32} \div 200 \text{ MHz}$). For short measurements like the one above, the **SUB** trick masks the wrap correctly thanks to two's-complement arithmetic. But for a *scheduler* or *timer* running over minutes, hours, or days, you need one of two strategies:

1. **Capture the full 64-bit count.** **GETCT D WC** returns the upper 32 bits (with WC set); a plain **GETCT D** returns the lower 32. Read upper-then-lower-then-upper-again and retry if the upper word changed, then build a 64-bit timestamp.
2. **Work with deltas, not absolute time.** Compare $(\text{now} - \text{start})$ rather than $\text{now} > \text{deadline}$. The subtraction wraps correctly even when the counter does.

The 21.5-second wrap is *fast* — long enough that you won't see it in a basic example, short enough that real applications hit it constantly.

What We've Learned

You're now an optimization expert:

- ✓ Understanding the P2 pipeline
- ✓ Instruction timing knowledge
- ✓ **REP** and **SKIP** for zero-overhead loops
- ✓ FIFO for maximum throughput
- ✓ Parallel operation techniques
- ✓ Real-world optimization strategies

Coming Up Next

Chapters 13-15 dig into LUT memory, smart pins, and event-driven programming - with references to dedicated manuals for deep dives. Think of them as appetizers showing what's possible!

Have Fun! Remember, the best optimization is often a better algorithm. But when you need every last cycle, you now know how to get them!

Chapter 13: LUT Memory - Your Private Lookup Table

512 longs of fast, deterministic storage in every cog

The Hook: A Lookup Table in 3 Cycles

Need fast data lookup without hub timing? Every cog has its own private 512-long Lookup RAM (LUT):

```
' Sine table lookup - 3 clocks, every time
get_sine
    AND    angle, #$FF      ' Mask to table index
    RDLUT  value, angle     ' Read from LUT in 3 clocks!
    RET
```

No hub timing to worry about. No waiting for the egg beater. Just 3 clock cycles, guaranteed. The LUT is like having a personal data assistant that never takes a coffee break.

Why Another Memory?

You might be thinking, “Wait, I already have cog RAM and hub RAM - why do I need a third memory?” Excellent question!

Memory	Size per Cog	Access Time	Special Features
Cog RAM	512 longs	2 clocks	Instructions live here
Hub RAM	512 KB shared	9-16 clocks (hub slot wait)	Shared by all cogs
LUT RAM	512 longs	3 clocks	Private, deterministic; shareable with a neighbor cog (see below)

The LUT fills a sweet spot: faster than hub memory, doesn’t compete with your instruction space, and has a trick up its sleeve - neighboring cogs can share LUTs!

Reading and Writing the LUT

Basic LUT Access

```
' Write to LUT
WRLUT    ##$12345678, #100 ' Write 32-bit constant to LUT[100]
WRLUT    value, index      ' Write variable to LUT[index]

' Read from LUT
RDLUT    result, #100      ' Read LUT[100] into result
RDLUT    data, index       ' Read LUT[index] into data
```

Notice the operand order: **WRLUT** writes its first operand to the address in the second, while **RDLUT** reads from its second operand into the first. A bit backwards from what you might expect, but you'll get used to it.

Building a Lookup Table

Here's how to load a sine table into LUT:

```
' Copy 256-entry sine table from hub to LUT
load_sine_table
    MOV    index, #0
    LOC    ptr, #\sine_data_hub ' Hub address of table

.loop   RDLONG value, ptr++      ' Read from hub
        WRLUT value, index      ' Write to LUT
        ADD    index, #1
        CMP    index, #256 wz
    if_nz JMP    #.loop
    RET

' Now lookups are fast!
get_sine
    RDLUT    sine_value, angle ' 3 clocks!
    RET
```

Sidetrack

Bulk LUT Loading with SETQ2

For loading entire tables, **SETQ2** + **RDLONG** transfers hub data into LUT memory. The trick: when **SETQ2** is active, the **RDLONG** destination field (still 9-bit, 0–\$1FF) is interpreted as a LUT offset that maps physically to \$200–\$3FF:

continues on next page →

↪ continued from previous page

```

SETQ2  #256-1          ' 256 longs to LUT
RDLONG $000, hub_table_ptr ' Dest field 0 = physical LUT $200

```

The destination operand is the LUT offset, not the absolute address — so you write \$000 (LUT base), not \$200 (which would overflow the 9-bit field). Remember the -1 in **SETQ2** (same rule as **SETQ** for hub block transfers).

LUT Sharing Between Cogs

Here's something clever: adjacent cog pairs can share LUT data! When you enable LUT sharing with **SETLUTS**, writes your neighbor makes to their LUT are automatically *copied* to your LUT too.

```

' --- cog 1 (consumer) - MUST enable sharing FIRST ---
  SETLUTS #1          ' Enable LUT write copying FROM cog 0
  ' Now when cog 0 writes to its LUT, data is COPIED to our LUT

' --- cog 0 (producer) - writes AFTER consumer enables sharing ---
  WRLUT  message, #10  ' Write MY LUT[10] (copies to cog 1)
  WRLUT  #1, #0        ' Set ready flag (copies to cog 1)

' --- cog 1 (consumer) - reads its OWN LUT (which contains copies) ---
.wait  RDLUT  flag, #0  ' Read MY LUT[0] (contains copy from cog 0)
      CMP    flag, #1 wz
      if_nz  JMP    #.wait
      RDLUT  message, #10  ' Read MY LUT[10] (copied from cog 0)

```

The key instruction is:

- **SETLUTS**: Enable write copying - when neighbor writes with **WRLUT**, data is copied to YOUR LUT
- **RDLUT**: Read your own LUT (which now contains copied data)

Important: The consumer cog must enable **SETLUTS** *before* the producer writes, otherwise the writes won't be copied!

This gives you a 512-long shared buffer between cog pairs without touching hub memory. Perfect for high-bandwidth data passing!

Sidetrack**Which Cogs Are Neighbors?**

The LUT sharing pairs are fixed:

- Cog 0 ↔ cog 1
- Cog 2 ↔ cog 3
- Cog 4 ↔ cog 5
- Cog 6 ↔ cog 7

When you enable sharing, your odd/even companion cog's LUT *writes* are copied into your own LUT — so you read the copies from your own LUT. Sharing works only within the fixed pair; non-adjacent cogs cannot share.

Practical Examples

Enough theory — let's see what people actually use the LUT for in real code:

Fast Data Transformation

```
' Gamma correction table in LUT
' Input: 8-bit value in 'pixel'
' Output: Gamma-corrected value
gamma_correct
    AND    pixel, #$FF      ' Mask to 8 bits
    RDLUT  pixel, pixel     ' Transform via table
    RET

' Initialize gamma table (power law curve)
' Would be pre-calculated and loaded from hub
```

Circular Buffer in LUT

```
' Fast circular buffer using LUT
' 256-entry buffer at LUT addresses 0-255

buf_write_ptr    LONG    0
buf_read_ptr     LONG    0

put_byte
    WRLUT  data, buf_write_ptr
    ADD    buf_write_ptr, #1
    AND    buf_write_ptr, #$FF  ' Wrap at 256
    RET
```

continues on next page →

→ continued from previous page

```

get_byte
    RDLUT    data, buf_read_ptr
    ADD      buf_read_ptr, #1
    AND      buf_read_ptr, #$FF    ' Wrap at 256
    RET

```

Fast Stack in LUT

```

' Stack implementation in LUT
' Grows downward from $1FF
stack_ptr    LONG    $1FF

stack_push
    WRLUT    value, stack_ptr
    SUB      stack_ptr, #1
    RET

stack_pop
    ADD      stack_ptr, #1
    RDLUT    value, stack_ptr
    RET

```

LUT with the Streamer

Here's where LUT gets really interesting. The streamer can read directly from LUT to generate waveforms without any cog intervention:

```

' Fill LUT with waveform data
' Then let Streamer output it to DAC

load_waveform
    MOV      index, #0

.fill    MOV      value, index
        SHL      value, #24    ' Scale for DAC
        WRLUT    value, index
        ADD      index, #1
        CMP      index, ##512 wz
        if_nz    JMP      #.fill

' Now configure Streamer to read from LUT
' Streamer handles the rest - no cog cycles needed!

```

The streamer configuration for LUT reading is covered in detail in the Video and Audio manuals - but the key point is that your LUT becomes a 512-sample waveform buffer that plays automatically.

Common Gotchas

✗ **WRONG: Confusing LUT addresses**

```
' WRONG - This reads cog RAM, not LUT!
  MOV    value, $200      ' $200 is cog RAM address
```

✓ **RIGHT: Use RDLUT for LUT access**

```
' RIGHT - RDLUT addresses the LUT space
  RDLUT  value, #0        ' LUT address 0
```

✗ **WRONG: Reading LUT before neighbor writes**

```
' WRONG - No data to read yet!
  SETLUTS #1              ' Enable sharing
  RDLUT  data, #10        ' Empty - neighbor hasn't written!
```

✓ **RIGHT: Wait for neighbor's write signal**

```
' RIGHT - Wait for data to be copied
  SETLUTS #1              ' Enable sharing BEFORE neighbor writes
.wait  RDLUT  ready, #0    ' Check flag in MY LUT
      TJZ    ready, #.wait  ' Wait until neighbor writes
      RDLUT  data, #10      ' Now MY LUT has copied data
```

Medicine Cabinet

The Medicine Cabinet

LUT Memory Quick Reference

Instruction	Operation	Cycles
RDLUT D, S	Read LUTS into D	3
WRLUT D, S	Write D to LUTS	2
SETLUTS D	Enable LUT write copying (D[0]=1)	2

Memory Map:

- LUT addresses: 0-511 (512 longs = 2KB)
- Neighbor pairs: 0↔1, 2↔3, 4↔5, 6↔7

Best Uses:

- Lookup tables (sine, gamma, encoding)
- Fast circular buffers
- Cog-pair data sharing
- Streamer waveform source

Your Turn

Your Turn

Exercise 1: Build an 8-bit Encoder

Create a LUT-based ASCII to 7-segment display encoder. Load a 128-entry table where each entry maps an ASCII code to the 7-segment pattern for that character.

```
' Your code here:
' 1. Load segment patterns into LUT
' 2. Write encode_char routine
' Hint: rdlut segment_pattern, ascii_char
```

Your Turn

Exercise 2: High-Speed Cog Communication

Use LUT sharing to create a message passing system between cog 2 and cog 3:

- Cog 2 writes 8-long messages
- Cog 3 reads them without hub access
- Use a simple ready/ack protocol

continues on next page →

```
' Hint: Use LUT[0] as ready flag, LUT[1-8] as message buffer
```

What We've Learned

The LUT in your toolbox:

- ✓ 512 longs of fast, private memory in every cog
- ✓ 3-clock deterministic access via **RDLUT** / **WRLUT**
- ✓ Bulk loading via **SETQ2** + **RDLONG**
- ✓ cog-pair LUT sharing for high-bandwidth data passing
- ✓ streamer source for waveform generation

Coming Up Next

Chapter 14 hands you the keys to 64 autonomous I/O processors — smart pins. We'll cover the universal configuration pattern and the modes you'll reach for most often, then point you at the dedicated Smart Pins Manual for the deep dive.

Have Fun! And remember — a 3-clock private lookup table is a luxury most chips don't give you. Use it!

Chapter 14: Smart Pins Orientation

64 autonomous I/O processors waiting to do your bidding

The Hook: A UART in 4 Lines

Remember that tedious bit-bang serial from Chapter 8? Watch this:

```
' Configure pin as UART transmitter - done!
  DIRL    #TX_PIN                ' Reset pin first!
  WRPIN    ##P_ASYNC_TX, #TX_PIN  ' Configure as async TX
  WXPIN    ##BAUD_115200, #TX_PIN ' Set baud rate
  DIRH    #TX_PIN                ' Enable - runs on its own
```

That's it. The pin is now a fully autonomous UART transmitter. It handles start bits, stop bits, timing - everything. You just feed it bytes with **WYPIN** and it sends them. The pin has become a state machine.

And here's the mind-bending part: *every single one of the 64 pins can do this*. Or PWM. Or ADC. Or quadrature decoding. Or 28 other modes.

What Are Smart Pins, Really?

Each of the P2's 64 I/O pins contains its own little processor - a state machine that can operate completely independently of the cogs. This means:

- A pin configured as UART keeps sending/receiving without cog intervention
- A PWM output keeps running its duty cycle automatically
- An ADC samples continuously in the background
- A quadrature decoder tracks position even while your cog does other things

The cog only needs to configure the pin and occasionally read/write data. The pin does the rest.

The Universal Smart Pin Pattern

Every smart pin follows the same configuration pattern. This is **the most important thing to remember**:

```
' === THE SMART PIN RECIPE ===

' Step 1: RESET the pin (CRITICAL!)
```

continues on next page →

↪ continued from previous page

```

        DIRL    pin                ' Always start by resetting

' Step 2: CONFIGURE the mode
        WRPIN   mode, pin          ' What should this pin do?

' Step 3: SET the X parameter
        WXPIN   x_value, pin       ' Mode-specific parameter X

' Step 4: ENABLE the pin
        DIRH    pin                ' Start the magic!

' Step 5: WRITE Y / data -- AFTER enable!
        WYPIN   y_value, pin       ' Mode-specific parameter Y

```

Why is WYPIN last, *after* DIRH? For the trigger and serial modes the Y value is held at zero during reset, so a WYPIN written before DIRH is simply lost. Writing Y after enabling is the one ordering that works for *every* mode – so make it your habit. (For pure value modes the order doesn't matter, but one rule is easier to remember than two.)

Sidetrack

Why DIRL First?

The **DIRL** at the start isn't optional politeness - it's *required*. Smart pins must be reset before configuration to ensure they're in a known state. **SKIP** this and you'll get unpredictable behavior as old settings conflict with new ones.

Think of it like power-cycling a misbehaving device. Always start fresh.

The Core Instructions

Configuration Instructions

Instruction	Purpose
WRPIN mode, pin	Set the operating mode
WXPIN value, pin	Set X parameter (mode-specific)
WYPIN value, pin	Set Y parameter (mode-specific)
DIRH pin	Enable the smart pin
DIRL pin	Disable/reset the smart pin

Data Instructions

Instruction	Purpose
WYPIN data, pin	Write data to smart pin (same instruction!)
RDPIN data, pin	Read result, clear “ready” flag
RQPIN data, pin	Read result, keep “ready” flag
AKPIN pin	Acknowledge (clear “ready” flag only)

Status Instructions

Instruction	Purpose
TESTP pin WC	Check if IN flag is set (data ready)
TESTPN pin WC	Check if IN flag is clear

Understanding the IN Flag

Every smart pin has an IN flag that signals “something happened.” What that something is depends on the mode:

- **UART TX:** IN high = ready for another byte
- **UART RX:** IN high = byte received
- **ADC:** IN high = new sample ready
- **PWM:** IN high = period complete
- **Counter:** IN high = threshold reached

You check this flag with **TESTP** and clear it by reading with **RDPIN** (or explicitly with **AKPIN**).

```
' Wait for smart pin to be ready
wait_ready
    TESTP    #PIN wc        ' Check IN flag
    if_nc JMP    #wait_ready ' Loop if not ready
    RDPIN    data, #PIN     ' Read and clear flag
```

Sidetrack**Event-Driven Alternative**

Instead of polling with **TESTP**, you can use the event system:

```
SETSE1  #%001<<6 + PIN    ' Event when IN rises
WAITSE1                                ' Sleep until ready - no polling!
RDPIN   result, #PIN        ' Read the result
```

This is more efficient because your cog sleeps instead of spinning. See Chapter 15 for the full event story.

Common Smart Pin Modes

Here are the modes you'll use most often:

Asynchronous Serial (UART)

```
' Transmit mode
DIRL    #TX_PIN
WRPIN   ##P_ASYNC_TX | P_OE, #TX_PIN
WXPIN   ##(CLK_FREQ/BAUD)<<16 | 7, #TX_PIN ' Baud + 8 bits
DIRH    #TX_PIN

' Send a byte
send     TESTP   #TX_PIN wc      ' Wait for ready
if_nc   JMP      #send
WYPIN    BYTE, #TX_PIN          ' Send it
```

```
' Receive mode
DIRL    #RX_PIN
WRPIN   ##P_ASYNC_RX, #RX_PIN
WXPIN   ##(CLK_FREQ/BAUD)<<16 | 7, #RX_PIN
DIRH    #RX_PIN

' Get a byte
recv     TESTP   #RX_PIN wc      ' Check for received byte
if_nc   JMP      #recv
RDPIN    BYTE, #RX_PIN          ' Get it
SHR      BYTE, #24              ' Shift to low byte
```

PWM Output

```
' PWM mode - period + duty cycle
  DIRL    #PWM_PIN
  WRPIN   ##P_PWM_SAWTOOTH | P_OE, #PWM_PIN
  WXPIN   ##period, #PWM_PIN      ' Period in clocks
  DIRH    #PWM_PIN
  WYPIN   ##duty, #PWM_PIN        ' High time; Y after enable

' Change duty cycle on the fly
  WYPIN   ##new_duty, #PWM_PIN    ' Just update Y parameter
```

ADC Input

```
' ADC mode - continuous sampling
  DIRL    #ADC_PIN
  WRPIN   ##P_ADC | P_ADC_GPIO, #ADC_PIN
  WXPIN   ##13, #ADC_PIN          ' 14-bit mode (period = 2^13 clocks)
  DIRH    #ADC_PIN

' Read ADC value
read_adc
  RDPIN   adc_value, #ADC_PIN      ' Get N-bit ADC count (LSB-aligned)
```

Quadrature Encoder

```
' Quadrature decoder - A on pin, B on pin+1
  DIRL    #ENC_PIN
  WRPIN   ##P_QUADRATURE, #ENC_PIN
  DIRH    #ENC_PIN

' Read position
  RDPIN   position, #ENC_PIN       ' Get accumulated count
```


✓ **RIGHT: DIRH comes last**

```
' RIGHT - Configure completely, then enable
    DIRL    #PIN
    WRPIN   ##P_ASYNC_TX, #PIN
    WXPIN   ##BAUD, #PIN
    DIRH    #PIN                ' Enable last!
```

Medicine Cabinet

The Medicine Cabinet

Smart Pin Quick Reference

The Recipe:

1. **DIRL** pin — Reset the pin first
2. **WRPIN** mode, pin — Set the operating mode
3. **WXPIN** x, pin — Set X parameter
4. **DIRH** pin — Enable the smart pin
5. **WYPIN** y, pin — Write Y / data (after enable)

Common Modes:

- **UART TX:** P_ASYNC_TX — Serial transmit
- **UART RX:** P_ASYNC_RX — Serial receive
- **PWM:** P_PWM_SAWTOOTH — Sawtooth wave output
- **PWM:** P_PWM_TRIANGLE — Triangle wave output
- **ADC:** P_ADC — Analog input
- **Quadrature:** P_QUADRATURE — Encoder
- **NCO:** P_NCO_FREQ — Frequency output

Data Flow:

- **WYPIN** = Write data TO smart pin
- **RDPIN** = Read data FROM smart pin (clears IN)
- **TESTP** = Check if IN flag set

Golden Rule: **DIRL** before **WRPIN** · **WXPIN** before **DIRH** · **WYPIN** (data) after **DIRH**

Your Turn

Your Turn

Exercise 1: PWM LED Dimmer

Create a PWM output that dims an LED:

1. Configure a pin for PWM sawtooth mode
2. Set a 1 kHz period (at 160 MHz: period = 160,000)
3. Vary duty cycle from 0% to 100%

```
' Your code here:
' Hint: Change duty with WYPIN new_duty, #LED_PIN
```

Your Turn

Exercise 2: Simple Serial Echo

Set up UART at 115200 baud:

1. Configure RX on pin 63
2. Configure TX on pin 62
3. Echo every received byte back

```
' Your code here:
' At 160 MHz: baud_divisor = 160_000_000 / 115200 = 1389
' WXPIN format: (divisor << 16) | (bits - 1)
```

Going Deeper: This chapter covered the smart pin essentials - the configuration pattern and common modes. For complete coverage of all 32 modes, timing diagrams, and advanced techniques, see the dedicated “P2 Smart Pins Manual.”

What We’ve Learned

Smart pin essentials:

- ✓ Every pin contains its own state machine (32 modes available)
- ✓ The universal recipe: **DIRL** → **WRPIN** → **WXPIN** → **DIRH** → **WYPIN**
- ✓ The IN flag signals “something happened” (mode-specific)
- ✓ UART, PWM, ADC, quadrature — same configuration pattern
- ✓ smart pins free the cog for other work

Coming Up Next

Chapter 15 explores the event system — how to stop polling and start waiting, so your cog sleeps until something interesting happens. The companion to smart pins: when one tells the other a byte is ready, you want to be notified, not spinning.

Have Fun! And remember — every smart pin you configure is a coprocessor you don't have to babysit. That's leverage!

Chapter 15: Event-Driven Programming

Stop spinning, start waiting

The Hook: No More Polling Loops

Remember all those busy loops waiting for things to happen?

```
' OLD WAY: Spin waiting for serial data (burns CPU cycles!)
wait_rx TESTP    #RX_PIN wc      ' Check over and over
if_nc JMP        #wait_rx        ' Spin spin spin...
        RDPIN    data, #RX_PIN

' NEW WAY: Sleep until data arrives (zero CPU cycles!)
        SETSE1   #001<<6 + RX_PIN ' Wake on IN rise
        WAITSE1                      ' Sleep until event
        RDPIN    data, #RX_PIN
```

The event system lets your cog sleep while waiting. When the event happens, it wakes up instantly. No cycles wasted, and you respond the moment something happens.

Why Events Matter

Polling loops have two problems:

1. **They waste cycles** - The cog spins doing nothing useful
2. **They add latency** - You check periodically, so there's delay between "thing happened" and "you noticed"

The event system solves both. Your cog *sleeps* and *wakes the instant* something happens. It's like having a personal assistant tap your shoulder instead of constantly looking up to check.

The Four Selectable Events

Let's meet the cast. Every cog has four configurable event channels: SE1, SE2, SE3, and SE4. Each can be configured to trigger on different conditions:

Event	Configuration	Wait	Poll
SE1	SETSE1	WAITSE1	POLLSE1
SE2	SETSE2	WAITSE2	POLLSE2
SE3	SETSE3	WAITSE3	POLLSE3
SE4	SETSE4	WAITSE4	POLLSE4


Plus there are built-in timer events:

Timer	Wait	Poll
CT1	WAITCT1	POLLCT1
CT2	WAITCT2	POLLCT2
CT3	WAITCT3	POLLCT3

Configuring an Event

The **SETSE1** through **SETSE4** instructions take a 9-bit configuration value:

SETSE1/2/3/4 Configuration (9 bits):



Event Modes

Mode	Meaning
%000	Never (disabled)
%001	IN rises (smart pin ready)
%010	IN falls
%011	IN changes
%100	Pin is low (level)
%101	Pin is low (level)
%110	Pin is high (level)
%111	Pin is high (level)

EVENT_* Constants: When You Need Interrupts

While dedicated cogs are usually better than interrupts (see Chapter 11), sometimes you need them. The **SETINT1/2/3** instructions select which event triggers an interrupt using these constants:

Table 15.1.

Constant	Value	Description
EVENT_INT	%0000	Pin matches interrupt configuration
EVENT_CT1	%0001	CT equals CT1 (timer 1 target)
EVENT_CT2	%0010	CT equals CT2 (timer 2 target)
EVENT_CT3	%0011	CT equals CT3 (timer 3 target)
EVENT_SE1	%0100	Selectable event 1 triggered
EVENT_SE2	%0101	Selectable event 2 triggered
EVENT_SE3	%0110	Selectable event 3 triggered
EVENT_SE4	%0111	Selectable event 4 triggered
EVENT_PAT	%1000	SETPAT pattern detected
EVENT_FBW	%1001	Hub FIFO wrapped around
EVENT_XMT	%1010	Streamer needs data
EVENT_XFI	%1011	Streamer operation complete
EVENT_XR0	%1100	NCO frequency counter rolled

Continued on next page

Table 15.1. (Continued)

Constant	Value	Description
EVENT_XRL	%1101	Streamer read last LUT location (\$1FF)
EVENT_ATN	%1110	Another cog signaled attention
EVENT_QMT	%1111	CORDIC/PIX math complete

Using EVENT_* with SETINT:

```

' Enable INT1 when SE1 event occurs
  SETSE1  %#001<<6 + RX_PIN      ' SE1 = IN rise on RX_PIN
  SETINT1 #EVENT_SE1              ' INT1 fires when SE1 triggers

' Enable INT2 on timer match
  ADDCT2  target, ##200_000        ' Set timer 2 target
  SETINT2 #EVENT_CT2              ' INT2 fires when CT = CT2

' Enable INT3 when another cog signals
  SETINT3 #EVENT_ATN              ' INT3 fires on COGATN

```

Pro tip: You can also use these constants with WAITSE/POLLSE by first triggering them with the appropriate hardware condition, then waiting. But for most purposes, the SETSE mode bits (the table above this one) are what you'll configure directly.

Smart Pin Events

The most common use is waiting for a smart pin to have data:

```

' Wait for smart pin on pin 15 to be ready
  SETSE1  %#001<<6 + 15          ' IN rise on pin 15
  WAITSE1                                ' Sleep until ready
  RDPIN   data, #15                ' Get the data

```

Pin Edge Events

You can also wait for raw pin edges (without smart pin):

```

' Wait for rising edge on pin 5
  SETSE1  %#001<<6 + 5          ' Rising edge on pin 5
  WAITSE1                                ' Sleep until edge
  ' Edge detected!

' Wait for falling edge on pin 10

```

continues on next page →

```

SETSE2  ##010<<6 + 10    ' Falling edge on pin 10
WAITSE2
' Edge detected!

```

Timer Events

For precise timing, use the counter comparison events:

```

' Wait exactly 1 millisecond (at 200 MHz)
GETCT   target            ' Current time
ADD     target, ##200_000 ' +1ms at 200MHz
ADDCT1  target, #0        ' Set CT1 target
WAITCT1                               ' Sleep until CT >= CT1

' Alternative using WAITX (simpler but less precise)
WAITX   ##200_000        ' Wait ~1ms at 200MHz

```

The timer events are:

- **ADDCT1/ADDCT2/ADDCT3:** Set the comparison target
- **WAITCT1/WAITCT2/WAITCT3:** Wait until CT reaches target
- **POLLCT1/POLLCT2/POLLCT3:** Check (non-blocking) if target reached

Waiting vs Polling

Two ways to use events:

WAIT - Sleep Until Event

```

WAITSE1                               ' cog sleeps here
' Wakes instantly when event occurs

```

- Cog sleeps, uses no cycles
- Wakes immediately when event fires
- Can't do anything else while waiting

POLL - Check and Continue

```
POLLSE1 wc          ' Check event, clear if set
if_c JMP    #event_handler ' Handle if occurred
      ' Continue with other work...
```

- Cog keeps running
- Checks event flag, clears it
- Returns result in C flag
- Good for servicing multiple events

Multiple Events

With four SE channels, you can monitor multiple sources:

```
' Setup multiple events
SETSE1  %%001<<6 + RX_PIN      ' Serial data ready
SETSE2  %%001<<6 + BUTTON_PIN  ' Button pressed (rising edge)
SETSE3  %%001<<6 + ADC_PIN     ' ADC sample ready

event_loop
  POLLSE1 wc          ' Check serial
  if_c CALL    #handle_serial

  POLLSE2 wc          ' Check button
  if_c CALL    #handle_button

  POLLSE3 wc          ' Check ADC
  if_c CALL    #handle_adc

  JMP      #event_loop
```

Practical Examples

Time to put events to work. Three patterns you'll reach for again and again:

Timeout with Fallback

```
' Wait for serial data, but give up after 100ms
wait_with_timeout
  SETSE1  %%001<<6 + RX_PIN      ' Serial ready event

  GETCT   timeout
  ADD     timeout, ##16_000_000  ' 100ms at 160MHz
```

continues on next page →

```

        ADDCT1    timeout, #0

.wait    POLLSE1  wc          ' Check serial
    if_c    JMP      #.got_data
        POLLCT1  wc          ' Check timeout
    if_c    JMP      #.timed_out
        JMP      #.wait

.got_data
    RDPIN    data, #RX_PIN
    RET

.timed_out
    NEG      data, #1          ' Return error (-1)
    RET

```

Debounced Button Press

```

' Wait for clean button press with debounce
debounced_button
    SETSE1    ##001<<6 + BUTTON ' Rising edge
    WAITSE1   ' Wait for press

    WAITX     ##2_000_000        ' 10ms debounce at 200MHz

    TESTP     #BUTTON wc         ' Verify still pressed
if_nc    JMP      #debounced_button ' Bounce - try again
    RET      ' Clean press!

```

Precise Periodic Sampling

```

' Sample ADC at exactly 10 kHz
sample_loop
    GETCT     next_sample

.loop    ADDCT1  next_sample, ##16_000 ' 100us period
        WAITCT1 ' Wait for next slot

        RDPIN   sample, #ADC_PIN      ' Read sample
        WRLONG  sample, buffer_ptr    ' Store it
        ADD     buffer_ptr, #4

        JMP     #.loop

```

ATN - Inter-Cog Events

The ATN (attention) system lets cogs signal each other:

```
' cog 0: Signal another cog
  COGATN 0x0000_0010      ' Send ATN to cog 1

' cog 1: Wait for attention
  WAITATN                  ' Sleep until ATN received
  ' Another cog signaled us!
```

The **COGATN** instruction takes an 8-bit mask where each bit corresponds to a cog. Setting bit N sends attention to Cog N.

Common Gotchas

✗ **WRONG: Forgetting to clear event flag**

```
' WRONG - Event may fire before you're ready
  SETSE1 0x001<<6 + PIN
  ' ... do other stuff ...
  WAITSE1                  ' May return immediately!
```

✓ **RIGHT: Poll first to clear any pending event**

```
' RIGHT - Clear any stale event
  SETSE1 0x001<<6 + PIN
  POLLSE1                  ' Clear if already set
  ' ... do other stuff ...
  WAITSE1                  ' Now wait cleanly
```

✗ **WRONG: Using WAIT when you need to handle multiple sources**

```
' WRONG - Can only wait for one event at a time
  WAITSE1                  ' Stuck here until SE1
  ' SE2 might fire and be missed!
```

✓ **RIGHT: Use POLL loop for multiple events**

```
' RIGHT - Check all sources
.loop  POLLSE1 wc
    if_c  CALL    #handle_se1
        POLLSE2 wc
    if_c  CALL    #handle_se2
        JMP     #.loop
```

Medicine Cabinet

The Medicine Cabinet

Event System Quick Reference

Configure Events:

```
SETSE1/2/3/4  #%%MMM_PPPPPP    ' Mode and pin
```

Event Modes:

%%MMM	Trigger
%001	IN rises (smart pin ready)
%010	IN falls
%011	IN changes
%10x	Pin is low (level)
%11x	Pin is high (level)

Wait (blocking):

```
WAITSE1/2/3/4    ' Sleep until event
WAITCT1/2/3      ' Sleep until timer
WAITATN          ' Sleep until attention
```

Poll (non-blocking):

```
POLLSE1/2/3/4 WC ' Check event, clear flag, C=occurred
POLLCT1/2/3 WC   ' Check timer, C=reached
POLLATN WC       ' Check attention, C=received
```

continues on next page →

Timer Setup:

```
ADDCT1/2/3 target, #delta    ' Set comparison target
```

Inter-cog:

```
COGATN #mask    ' Signal cogs (bit per cog)
```

Your Turn

Your Turn**Exercise 1: Event-Driven Serial**

Rewrite a serial receive loop to use events instead of polling:

1. Configure SE1 for UART RX smart pin ready
2. Use WAITSE1 instead of **TESTP** loop
3. Measure the cycle count difference

```
' Your code here:
' Hint: setse1 #%%001<<6 + RX_PIN
```

Your Turn**Exercise 2: Dual Event Monitor**

Create a loop that monitors both a button (pin edge event) and a timer (periodic event):

1. SE1 = button press (rising edge)
2. CT1 = 1 second heartbeat
3. On button: toggle LED
4. On timer: print timestamp

```
' Your code here:
' Use POLL for both, handle whichever fires
```

What We've Learned

The event toolkit you now command:

- ✓ Four selectable event channels (SE1-SE4) per cog
- ✓ Three timer events (CT1-CT3) for precise scheduling
- ✓ **WAIT** (sleep) vs **POLL** (check-and-continue) tradeoffs
- ✓ The ATN inter-cog signaling system
- ✓ Common patterns: timeout-with-fallback, debounce, periodic sampling

Coming Up Next

Chapter 16 brings the whole journey together — orchestrating eight cogs in parallel harmony to build complete systems. It's where the P2 philosophy really shines.

Have Fun! And remember — every spin loop you replace with **WAITSE** is cog cycles you've handed back to your design. Be generous with events!

Chapter 16: Multi-Cog Orchestration

Bringing it all together in parallel harmony

The Hook: A Complete System in 8 Cogs

Watch this system architecture come alive:

```
' Main orchestrator (cog 0)
main_orchestrator
    ' Launch the orchestra (SETQ sets PTRA for new cog)
    SETQ    @sensor_params
    COGINIT #1, @sensor_cog
    SETQ    @motor_params
    COGINIT #2, @motor_cog
    SETQ    @comms_params
    COGINIT #3, @comms_cog
    SETQ    @display_params
    COGINIT #4, @display_cog
    SETQ    @safety_params
    COGINIT #5, @safety_cog
    SETQ    @logger_params
    COGINIT #6, @logger_cog
    SETQ    @debug_params
    COGINIT #7, @debug_cog

    ' Now coordinate them all
orchestrate
    RDLONG  sensor_data, ##SENSOR_MAILBOX wz
    if_nz CALL  #process_sensor_data

    RDLONG  command, ##COMMAND_MAILBOX wz
    if_nz CALL  #execute_command

    CALL    #update_system_state
    WRLONG  state, ##STATE_MAILBOX

    JMP     #orchestrate
```

Eight independent processors, each with a specific job, all working in perfect coordination. This is the true power of P2!

Communication Patterns

Eight processors running in parallel sounds wonderful — until you realize they need to talk to each other. Let's meet the three patterns you'll use 95% of the time.

The Mailbox Pattern

The simplest and most common — a single hub long that one cog writes and another reads:

```
' Producer cog
producer
    ' Generate data
    CALL    #calculate_result
    WRLONG  result, ##MAILBOX_ADDR

' Consumer cog
consumer
    RDLONG  data, ##MAILBOX_ADDR wz
    if_z JMP    #consumer          ' Wait for data
    WRLONG  #0, ##MAILBOX_ADDR     ' Clear mailbox
    CALL    #process_data
```

The Ring Buffer Pattern

For streaming data between cogs:

```
' Writer cog
writer_cog
    RDLONG  wr_ptr, ##WRITE_PTR
    WRLONG  data, wr_ptr
    ADD     wr_ptr, #4
    AND     wr_ptr, ##BUFFER_MASK ' Wrap around
    WRLONG  wr_ptr, ##WRITE_PTR

' Reader cog
reader_cog
    RDLONG  rd_ptr, ##READ_PTR
    RDLONG  wr_ptr, ##WRITE_PTR
    CMP     rd_ptr, wr_ptr wz
    if_z JMP    #reader_cog        ' Buffer empty

    RDLONG  data, rd_ptr
    ADD     rd_ptr, #4
    AND     rd_ptr, ##BUFFER_MASK
    WRLONG  rd_ptr, ##READ_PTR
```

The Command Queue Pattern

For sending commands between cogs:


```

' Command structure in hub
' +0: Command ID
' +4: Parameter 1
' +8: Parameter 2
' +12: Result/Status

' Commander cog
send_command
    WRLONG    cmd_id, ##CMD_BUFFER+0
    WRLONG    param1, ##CMD_BUFFER+4
    WRLONG    param2, ##CMD_BUFFER+8
    WRLONG    ##$FFFF, ##CMD_BUFFER+12 ' Mark as pending

wait_complete
    RDLONG    status, ##CMD_BUFFER+12
    CMP       status, ##$FFFF wz
    if_z JMP  #wait_complete

' Worker cog
process_commands
    RDLONG    status, ##CMD_BUFFER+12
    CMP       status, ##$FFFF wz
    if_nz JMP  #process_commands      ' No command

    RDLONG    cmd_id, ##CMD_BUFFER+0
    RDLONG    param1, ##CMD_BUFFER+4
    RDLONG    param2, ##CMD_BUFFER+8

    CALL      #execute_command
    WRLONG    result, ##CMD_BUFFER+12 ' Signal complete

```

Synchronization Techniques

Sometimes communication isn't enough — you need two or more cogs to *agree* on what happens when. That's where synchronization comes in.

Using Locks

When multiple cogs need atomic access to the same piece of data, P2 gives you 16 hardware locks. They're tiny and they're fast:

```

' Atomic increment using lock
atomic_increment
    LOCKTRY   #COUNTER_LOCK wc      ' C=1 if we got the lock
    if_nc JMP  #atomic_increment    ' Retry until we get it

    RDLONG    value, ##COUNTER
    ADD       value, #1

```

continues on next page →

```
WRLONG  value, ##COUNTER

LOCKREL #COUNTER_LOCK
```

Event Synchronization

Cogs waiting for specific events:

```
' cog 1: Signal event
    WRLONG  ##EVENT_FLAG, ##EVENT_ADDR

' cog 2: Wait for event
wait_event
    RDLONG  flag, ##EVENT_ADDR wz
    if_z JMP  #wait_event
    WRLONG  #0, ##EVENT_ADDR      ' Clear event
```

Real-World Example: Robot Controller

Let's build a complete robot control system:

```
' cog 0: Main Controller
main_controller
    CALL    #init_system

main_loop
    ' Read sensor hub
    RDLONG  distance, ##DISTANCE_SENSOR wz
    if_z JMP  #too_close

    ' Check for commands
    RDLONG  cmd, ##SERIAL_COMMAND wz
    if_nz CALL  #process_command

    ' Update motor speeds
    CALL    #calculate_motion
    WRLONG  left_speed, ##LEFT_MOTOR
    WRLONG  right_speed, ##RIGHT_MOTOR

    JMP     #main_loop

' cog 1: Ultrasonic Sensor
sensor_cog
    ' Trigger ultrasonic pulse
    DRVH    #TRIGGER_PIN
    WAITX   ##1000
    DRVL    #TRIGGER_PIN
```

```

        ' Measure echo time - wait for rising edge
.wait_hi
        TESTP    #ECHO_PIN wz
        if_nz JMP    #.wait_hi
        GETCT    start_time
.wait_lo                                ' Wait for falling edge
        TESTP    #ECHO_PIN wz
        if_z  JMP    #.wait_lo
        GETCT    end_time

        ' Calculate distance
        SUB      end_time, start_time
        ' Convert to distance...
        WRLONG   distance, ###DISTANCE_SENSOR

        WAITX    ##10_000_000          ' 50ms at 200MHz
        JMP      #sensor_cog

' cog 2: Left Motor Driver
left_motor_cog
        RDLONG   speed, ##LEFT_MOTOR wz
        if_z  JMP    #left_motor_cog      ' No speed set

        ' Generate motor control signals
        ' ... PWM generation code
        JMP      #left_motor_cog

' cog 3: Right Motor Driver
' (Similar to left motor)

' cog 4: Serial Communications
serial_cog
        ' Check for incoming commands
        TESTP    #RX_PIN wc
        if_nc JMP    #serial_cog

        CALL      #receive_byte
        ' Build command...
        WRLONG   command, ##SERIAL_COMMAND
        JMP      #serial_cog

' cog 5: LED Status Display
status_cog
        RDLONG   system_state, ##STATE_MAILBOX

        ' Display state on LEDs
        CMP      system_state, #STATE_RUNNING wz
        if_z  DRVH    #GREEN_LED
        if_nz DRVL    #GREEN_LED

        CMP      system_state, #STATE_ERROR wz
        if_z  DRVH    #RED_LED
        if_nz DRVL    #RED_LED

```

→ continued from previous page

```

        WAITX    ##10_000_000
        JMP      #status_cog

' cog 6: Safety Monitor
safety_cog
' Monitor critical systems
RDLONG  battery, ##BATTERY_VOLTAGE
CMP     battery, ##MIN_VOLTAGE wcz
if_b WRLONG ##STATE_SHUTDOWN, ##STATE_MAILBOX

' Check temperature
RDLONG  temp, ##TEMPERATURE
CMP     temp, ##MAX_TEMP wcz
if_a WRLONG ##STATE_OVERHEAT, ##STATE_MAILBOX

        JMP     #safety_cog

' cog 7: Debug Output
debug_cog
' Output system state for debugging
' ... debug code

```

Eight cogs, each doing one job perfectly, creating a responsive, reliable robot!

Your Turn: Multi-Cog Project

Your Turn

Exercise: Traffic Light Controller

Requirements:

- Cog 0: Main sequencer
- Cog 1: North-South lights
- Cog 2: East-West lights
- Cog 3: Pedestrian button watcher
- Cog 4: Timer/scheduler

Starting structure:

```

        ORG      0
' cog 0: Main sequencer
' Launch other cogs
' Coordinate light changes
' Handle pedestrian requests

' Your implementation here

```

continues on next page →

Goal: Working traffic light with pedestrian crossing Hint: Use mailboxes for cog communication
 Success Check: Lights change correctly, pedestrian button works

The Medicine Cabinet

Multi-cog systems overwhelming? Start simple:

Start with just 2 cogs:

```
' Main + Helper pattern
main    SETQ    @params          ' PTRA for new cog
        COGINIT #1, @helper
        ' Main work

helper  ' Support work
```

Use simple mailboxes:

```
' Fixed hub addresses for communication
MAILBOX_1 = $1000
MAILBOX_2 = $1004
```

Debug one cog at a time: TEST each cog in isolation before combining!

Design Principles for Multi-Cog Systems

The hardware gives you eight processors. Whether your *design* survives the journey is up to you. A few rules of thumb we've learned the hard way:

1. **Single Responsibility:** Each cog does ONE thing well
2. **Loose Coupling:** Cogs communicate through hub, not direct dependencies
3. **Clear Ownership:** Each piece of data has one writer
4. **Predictable Timing:** Real-time tasks get dedicated cogs
5. **Graceful Degradation:** System continues if one cog fails

Common Multi-Cog Gotchas

Before you pull your hair out wondering why the eight-cog dream turned into a debugging nightmare, skim these:

1. **Race conditions** - Use locks for shared write access
2. **Deadlocks** - Avoid circular dependencies
3. **Starvation** - Ensure all cogs get resources
4. **Communication overhead** - Don't over-communicate
5. **Debugging complexity** - Use LED indicators for each cog

What We've Learned

You've mastered multi-cog orchestration:

- ✓ Communication patterns (mailbox, ring buffer, queue)
- ✓ Synchronization techniques
- ✓ Real-world system architecture
- ✓ Design principles
- ✓ Common pitfalls and solutions

This is it - you now understand the full power of the Propeller 2!

Your Journey Continues

You've completed this manual, but your P2 journey has just begun:

1. **Build something amazing** - Put your knowledge to work
2. **Share with the community** - Your projects inspire others
3. **Explore other manuals** - smart pins, Video, I/O await
4. **Push boundaries** - P2 can do things we haven't imagined yet

Chapter Summary

Chapter Summary

Congratulations! You've mastered multi-cog orchestration!

You now understand:

- How to coordinate 8 parallel processors

continues on next page →

- Communication patterns between cogs
- Synchronization techniques
- Real-world system design

You did it! You're now fluent in PASM2 and ready to build incredible parallel systems!

Have Fun!

Remember what you've learned:

- Eight cogs working together are more powerful than any interrupt-driven system
- Parallel processing isn't harder, it's different
- The P2 way is about determinism and elegance
- Every complex system is just simple parts working together

Now go forth and create something amazing with your Propeller 2!

Epilogue: The Journey Forward

Well, here we are at the end... or should I say, at the beginning?

You've traveled from blinking your first LED to orchestrating eight parallel processors. You've mastered CORDIC mathematics, tamed the FIFO, and learned why interrupts are usually the wrong answer. That's quite a journey!

But here's the secret: everything you've learned is just the foundation. The P2 community continues to discover new techniques, new optimizations, new ways to use this remarkable chip. Every project pushes the boundaries a little further.

What Makes You Different Now

You're not just another embedded programmer anymore. You think in parallel. You see solutions that others miss. When someone says "that's impossible in real-time," you know better - you just dedicate a cog to it.

The Community Awaits

The Parallax forums are filled with fellow travelers on this journey. Share your projects. Ask questions. Help newcomers. The community that inspired this manual continues to grow because people like you contribute back.

One Last Story

I remember my first P2 project. I was trying to control 16 servos with perfect timing while reading sensors and communicating over serial. On my previous microcontroller, it was a nightmare of interrupts and jitter.

On the P2? Three cogs. Clean, simple, perfect timing. That's when I truly understood - this isn't just a different processor, it's a different philosophy of computing.

Your Challenge

Build something that wouldn't be possible without parallel processing. Something that would be a nightmare of interrupts on other processors. Then share it with the world.

Show them what eight cogs can do.

Show them the Propeller way.

"The best way to predict the future is to invent it."

— Alan Kay

And with your Propeller 2, you have everything you need to invent amazing futures.

Have Fun!

— The ghosts of deSilva, the P2 community, and one very enthusiastic AI

P.S. - Don't forget to blink an LED once in a while, just for old times' sake. It's still magical, even after all you've learned.

THE END

(But really, just the beginning...)

Further Reading

This teaching manual focuses on concepts, patterns, and building your understanding. For complete technical specifications, refer to these companion documents:

Propeller 2 Assembly Language (PASM2) Manual Complete PASM2 instruction details including syntax, timing, and flag effects for all 300+ instructions. Quick lookup reference for day-to-day development.

Parallax Propeller 2 Documentation (v35, Rev B/C silicon, 2021-05-18) Official silicon documentation from Parallax covering hardware specifications, electrical characteristics, and detailed register maps.

Appendix A: Platform Comparison

How P2 compares to other microcontrollers

If you're coming from another embedded platform, this appendix shows how the P2's approach differs and when those differences matter.

The Landscape

The embedded world is dominated by a handful of architectures:

Platform	Architecture	Typical Cores	Peripherals	Timing
STM32	ARM Cortex-M	1-2	Fixed location	Cache-dependent
ESP32	Xtensa/RISC-V	2	Fixed location	FreeRTOS scheduled
Arduino/AVR	AVR	1	Fixed location	Deterministic but slow
PIC32	MIPS	1	Fixed location	Interrupt-driven
P2 Propeller	Custom	8	Any pin	Deterministic

What Makes P2 Different

Eight Real Processors, Not Time-Slicing

On ARM, ESP32, or PIC, you typically have 1-2 cores that share time between tasks using interrupts or an RTOS. The P2 gives you eight complete, identical processors that run truly in parallel.

Traditional approach:

```
' Everyone fights for the same CPU
ISR(TIMER1_vect) { motor_control(); }   ' Might delay...
ISR(UART_RX_vect) { serial_handler(); } ' ...this
main() { while(1) { sensor_loop(); } } ' Hope we get time
```

P2 approach:

```
' Each task owns its own processor
COG0: COGINIT(1, @motor_control)  ' Coordinator launches workers
COG1: motor_control()              ' Dedicated - always on time
COG2: serial_handler()             ' Dedicated - never misses byte
COG3: sensor_loop()               ' Dedicated - consistent sample
COG4-7: ready for more
```

No interrupt priority juggling. No RTOS configuration. Each task owns its processor.

Smart Pins: Peripherals on Every Pin

Traditional MCUs have fixed peripheral assignments: UART1 is on PA9/PA10, SPI1 is on PB3/PB4/PB5, and if you need those pins for something else, you're stuck rerouting your PCB.

On P2, every pin contains a programmable state machine. Any pin can become a UART, SPI, PWM, ADC, quadrature decoder, or 27 other modes. The peripheral comes to your pin, not the other way around.

Deterministic Timing

ARM MCUs with cache have unpredictable timing. A memory read might take 1 cycle (cache hit) or 50+ cycles (cache miss). Even instruction timing varies—ARM instructions take 1-3+ cycles depending on the operation. This makes cycle-accurate timing extremely difficult.

P2 takes a different approach: nearly all instructions execute in exactly **2 clock cycles**. Want to know how long a code sequence takes? Count the instructions and multiply by 2. Hub memory uses round-robin access that gives every cog predictable, guaranteed access slots. Your timing loops work identically every time—no cache luck required.

Coming From ARM/STM32

You're used to configuring HAL structures, writing interrupt handlers, and managing DMA. Here's how P2 solves those problems:

Instead of...	On P2...	The Benefit
HAL_UART_Transmit()	Configure smart pin once, then WYPIN bytes	Pin handles all timing autonomously
HAL_TIM_PWM_Start()	Configure smart pin once, update with WYPIN	Pin runs independently—your cog is free
NVIC priority configuration	Nothing needed	All cogs equal, no priority inversion ever
HAL_DMA_Start()	Use built-in FIFO/Streamer	Simpler API, integrated into each cog
arm_sin_f32() library	QROTATE instruction	Hardware trig in exactly 55 clocks
FreeRTOS xTaskCreate()	COGINIT	True parallel execution, not scheduled

The result: Deterministic timing, zero interrupt conflicts, and I/O configuration that just works.

Coming From ESP32

You're used to WiFi/Bluetooth convenience and FreeRTOS abstractions. P2 takes a different approach:

ESP32 Way	P2 Way	The Benefit
Built-in WiFi/BT	Add WizNet or ESP module	You choose your connectivity—or skip it entirely
xTaskCreate()	COGINIT	Not scheduled—truly parallel, guaranteed timing
GPIO matrix routing	smart pins	32 modes per pin, far more capability
FreeRTOS timing	Deterministic hub	Cycle-accurate timing guaranteed
Arduino framework	Spin2/PASM2	Deeper control, deeper understanding

The result: 8 real cores running simultaneously, timing you can count on, I/O flexibility that eliminates peripheral conflicts.

Coming From Arduino/AVR

You'll find P2 familiar but dramatically more powerful:

Arduino Way	P2 Way	The Upgrade
<code>digitalWrite()</code>	DRVH/DRVL or smart pins	Similar syntax, vastly more capability
<code>delay()</code> blocks everything	WAITX or dedicated cog	Timing without blocking other tasks
One thing at a time	8 things truly parallel	Real concurrency, not fake multitasking
8-bit math limits	32-bit + hardware CORDIC	No more overflow worries, hardware trig
Libraries for everything	Growing ecosystem + OBEX	More control, deeper understanding

The result: Graduate from 8-bit limitations to 8 parallel 32-bit processors with hardware math and smart pins on every I/O.

When P2 Is the Right Choice

P2 excels when you need:

- **Multiple real-time tasks** running simultaneously without conflicts
- **Precise timing** that cache misses and interrupts can't disrupt
- **Video or audio generation** requiring cycle-accurate output
- **Flexible I/O** where any pin can become any peripheral
- **Hardware math** for motor control, signal processing, or robotics
- **Multiple motor/servo control** with dedicated cogs per channel
- **Protocol implementation** where smart pins handle timing autonomously

Platform Trade-offs

Every platform makes trade-offs. P2 optimizes for **determinism, parallelism, and flexibility** rather than:

If you need...	P2's answer
Built-in WiFi/Bluetooth	Add WizNet or ESP module—you choose connectivity
Massive library ecosystem	Growing OBEX + helpful community
Ultra-low-power sleep	External modules or different platform
Lowest unit cost at 100K+ volumes	P2 targets flexibility over commodity pricing

The honest reality: If your project is “connect to WiFi and display data,” an ESP32 does that with less effort. But if your ESP32 project is fighting timing jitter, missing deadlines, or running out of peripheral pins—that’s exactly what P2 solves.

Community Resources

While P2’s ecosystem is smaller than ARM or Arduino, it’s active and welcoming:

Parallax Forums - The heart of the P2 community. Chip Gracey (P2’s designer) participates actively, answering questions and discussing design decisions. You’ll find help from experienced developers who’ve solved problems you haven’t encountered yet.

P2 Object Exchange (OBEX) - A library of reusable Spin2 and PASM2 objects covering drivers, protocols, display interfaces, and more. Before writing something from scratch, check OBEX—someone may have already done the work.

Community Support - Unlike large platforms where your question disappears in a sea of posts, the P2 community is small enough that questions get noticed and answered. Many community members have decades of Propeller experience.

Coming from Arduino’s library-for-everything culture, you’ll write more code yourself—but you’ll understand it deeply, and help is always available when you get stuck.

The P2 Hardware Ecosystem

P2 isn’t just a chip - it’s a platform with expansion options:

Video & Audio:

- A/V Breakout Board: VGA, RCA, 80mW stereo, microphone input
- Digital Video Out: HDMI-type with differential signaling
- Built-in 8-bit DAC per pin (16-bit with dithering)

Connectivity:

- USB Host Board: Two USB-A ports
- USB Device Board: HID or CDC modes
- Serial Host/Device: RS-232 interfaces
- WizNet/ESP modules: Ethernet or WiFi

Development:

- P2 Eval Board: Complete development environment
- Edge Modules: 4MB or 32MB flash for embedding

- Breakout Boards: All 64 pins accessible

You add what you need - no paying for peripherals you won't use.

Summary

P2 represents a fundamentally different approach to embedded computing—one that eliminates entire categories of problems:

- **Eight processors** means your motor control never delays your serial handler
- **64 smart pins** means peripheral conflicts become impossible
- **Deterministic timing** means your code works the same way every time
- **Hardware CORDIC** means real-time math without floating-point libraries

Engineers who've fought interrupt priority inversions, missed timing deadlines, and PCB rework due to peripheral conflicts find P2 refreshing. You spend your time solving your actual problem, not fighting your MCU.

Welcome to the P2 community. You've got 8 processors, 64 smart pins, and a community that's been building amazing things since the original Propeller. Time to see what you can build.

Index

A

- ADC operations: Ch14
- **ADD** instruction: Ch3, Ch5
- ADDCT1/2/3: Ch15
- Address modes: Ch3
- **ALTD**/ALTS: Ch3
- Architecture: Ch2
- Assembly basics: Ch3

B

- Bit manipulation: Ch3
- Block transfers: Ch4, Ch9
- Booleans: Ch3

C

- C flag: Ch6
- **CALL**/RET: Ch3
- Clock timing: Ch2
- **CMP** instruction: Ch6
- Cog anatomy: Ch2
- Cog communication: Ch2, Ch16
- Conditional execution: Ch3, Ch6
- CORDIC: Ch7
- Counters: Ch2

D

- DAC operations: Ch14
- Debugging: Ch12
- Division: Ch5
- **DRVH**/DRVL: Ch1

E

- Egg beater: Ch2, Ch4
- Event system: Ch15

F

- FIFO: Ch4, Ch9
- Flags: Ch6
- Flow control: Ch3

G

- **GETQX**/GETQY: Ch5, Ch7

H

- Hardware multiply: Ch5
- Hub execution: Ch10
- Hub memory: Ch2, Ch4

I

- IN flag: Ch14, Ch15
- Immediate values: Ch3
- Instruction format: Ch3
- Interrupts: Ch11

J

- **JMP** instruction: Ch3

L

- LED control: Ch1
- Locks: Ch2, Ch16
- Logic operations: Ch3
- LUT memory: Ch13
- LUT sharing: Ch13

M

- Mailboxes: Ch2, Ch16
- Mathematics: Ch5
- Memory access: Ch4
- **MOV** instruction: Ch3
- **MUL**/MULS: Ch5
- Multi-cog: Ch16

O

- Optimization: Ch12

P

- Parallel processing: Ch2
- Pipeline: Ch7, Ch12
- Pins, Smart: Ch8, Ch14
- Platform comparison: Appendix A
- POLLSE1-4: Ch15
- PTRB/PTRA: Ch3, Ch4
- PWM: Ch8, Ch14

Q

- Q flag: Ch7
- **QDIV**: Ch5
- **QROTATE**: Ch7

R

- **RDBYTE**/RDWORD/RDLONG: Ch4
- **RDLUT**/WRLUT: Ch13
- **RDPIN**: Ch14, Ch15
- **REP** instruction: Ch3
- Rotation: Ch3, Ch7

S

- SETSE1-4: Ch15

- Shift operations: Ch3
- **SKIP** instruction: Ch3, Ch6
- Smart pins: Ch8, Ch14
- Streamer: Ch9, Ch13

T

- Timer: Ch2
- Timing: Ch2, Ch12
- Trigonometry: Ch7

U

- UART: Ch8, Ch14

V

W

- WAITSE1-4: Ch15
- WAITCT1-3: Ch15
- **WAITX**: Ch1
- **WRBYTE**/WRWORD/WRLONG: Ch4
- **WRPIN**/WXPIN/WYPIN: Ch14

Z

- Z flag: Ch6